

# **GreenCom Networks**

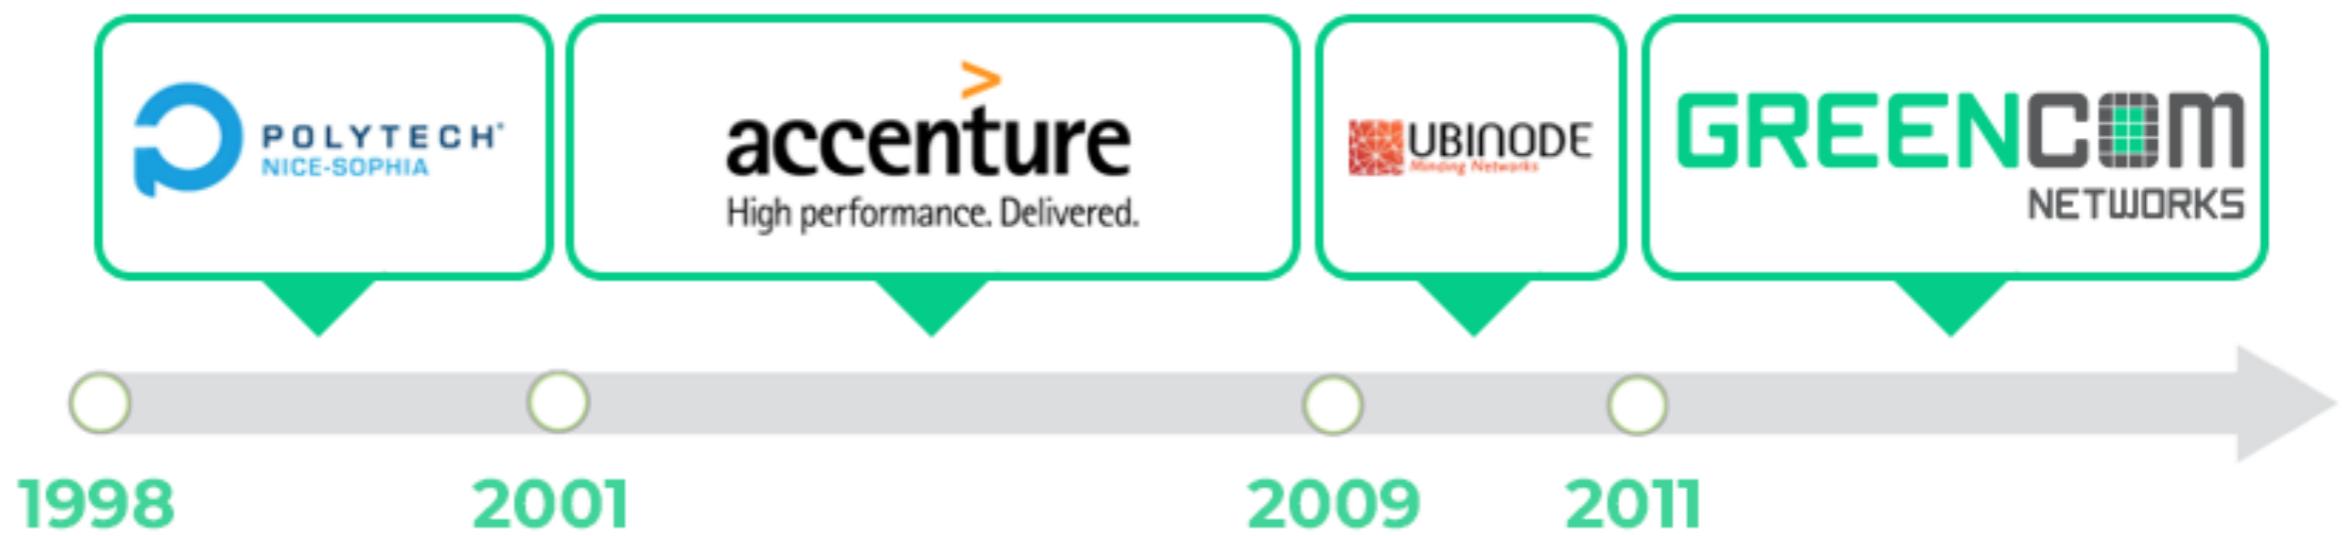
# **Platform Architecture**

January 2019

Sebastien Alegret  
CTO@ GreenCom Networks

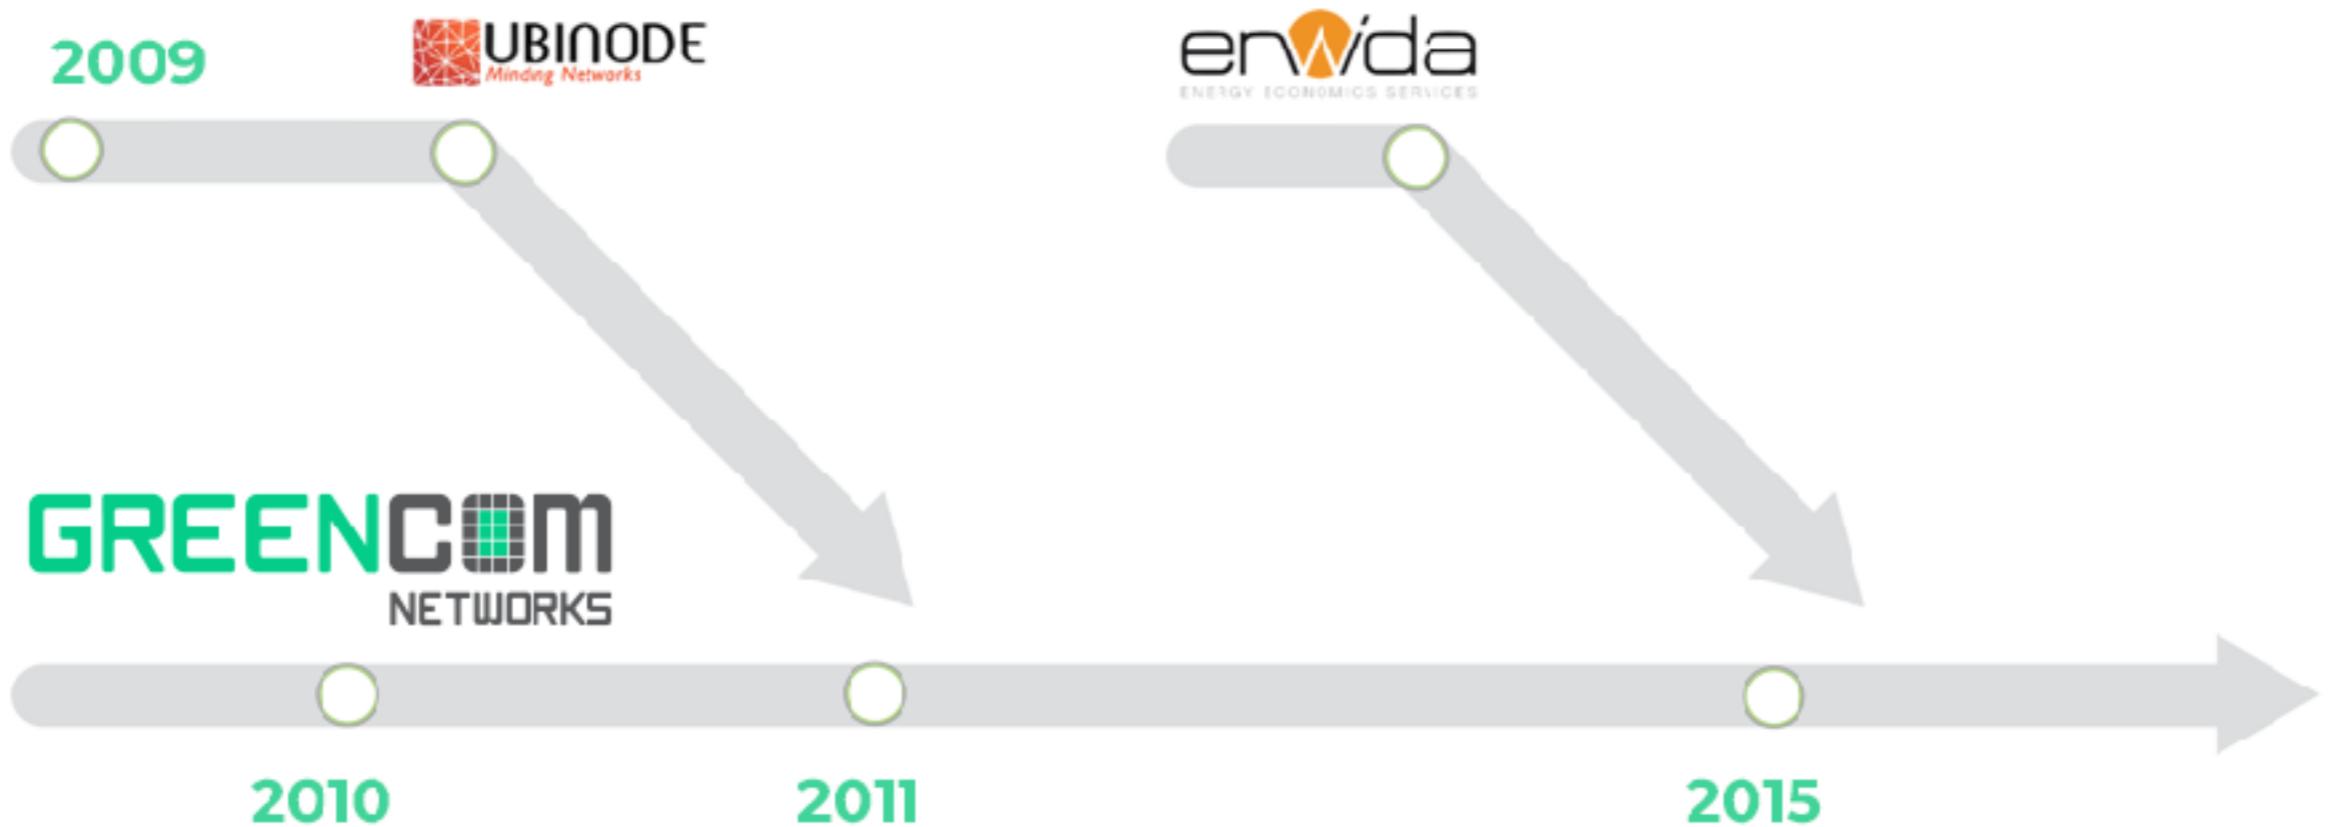
Mathias Chevalier  
Infrastructure/Devops Team

# A bit of (my) history

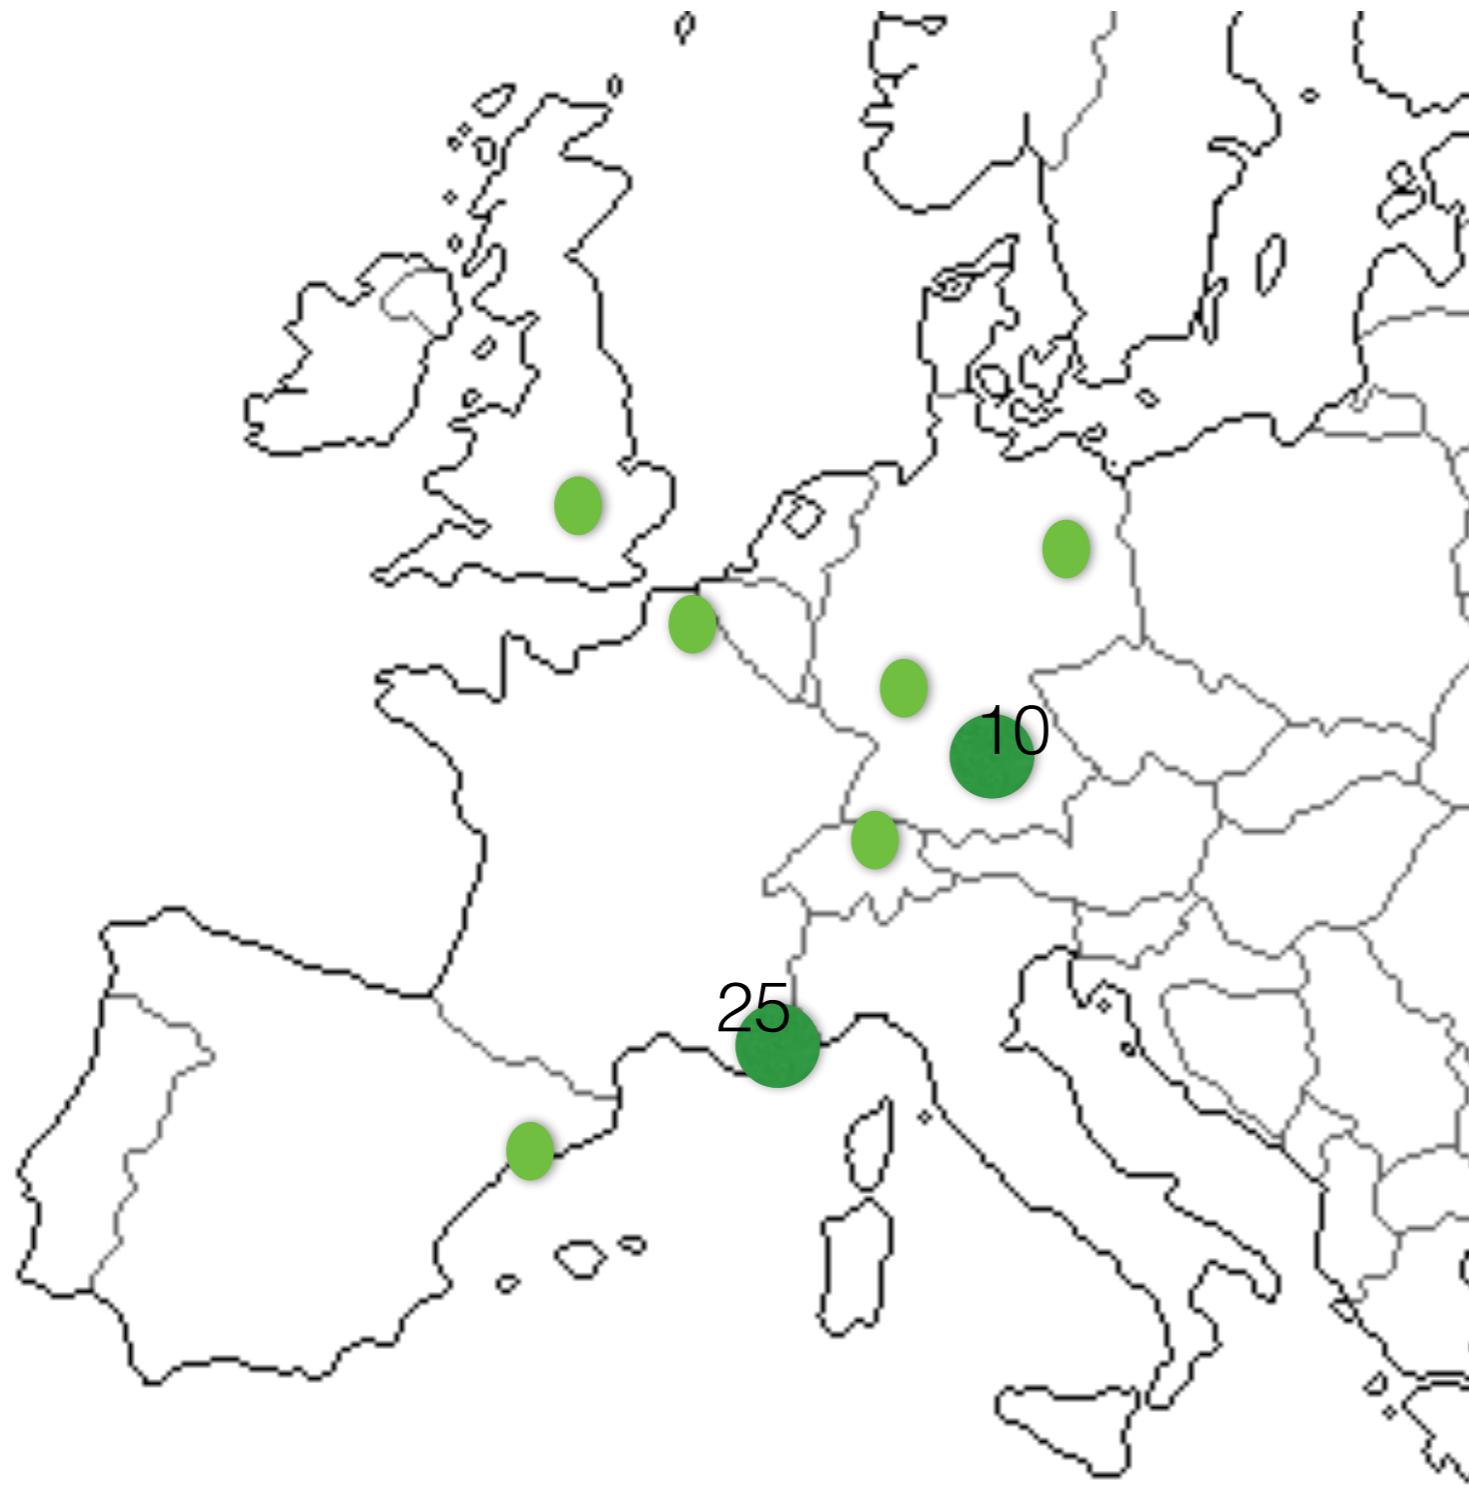


# **GreenCom Networks**

# A bit of GCN history



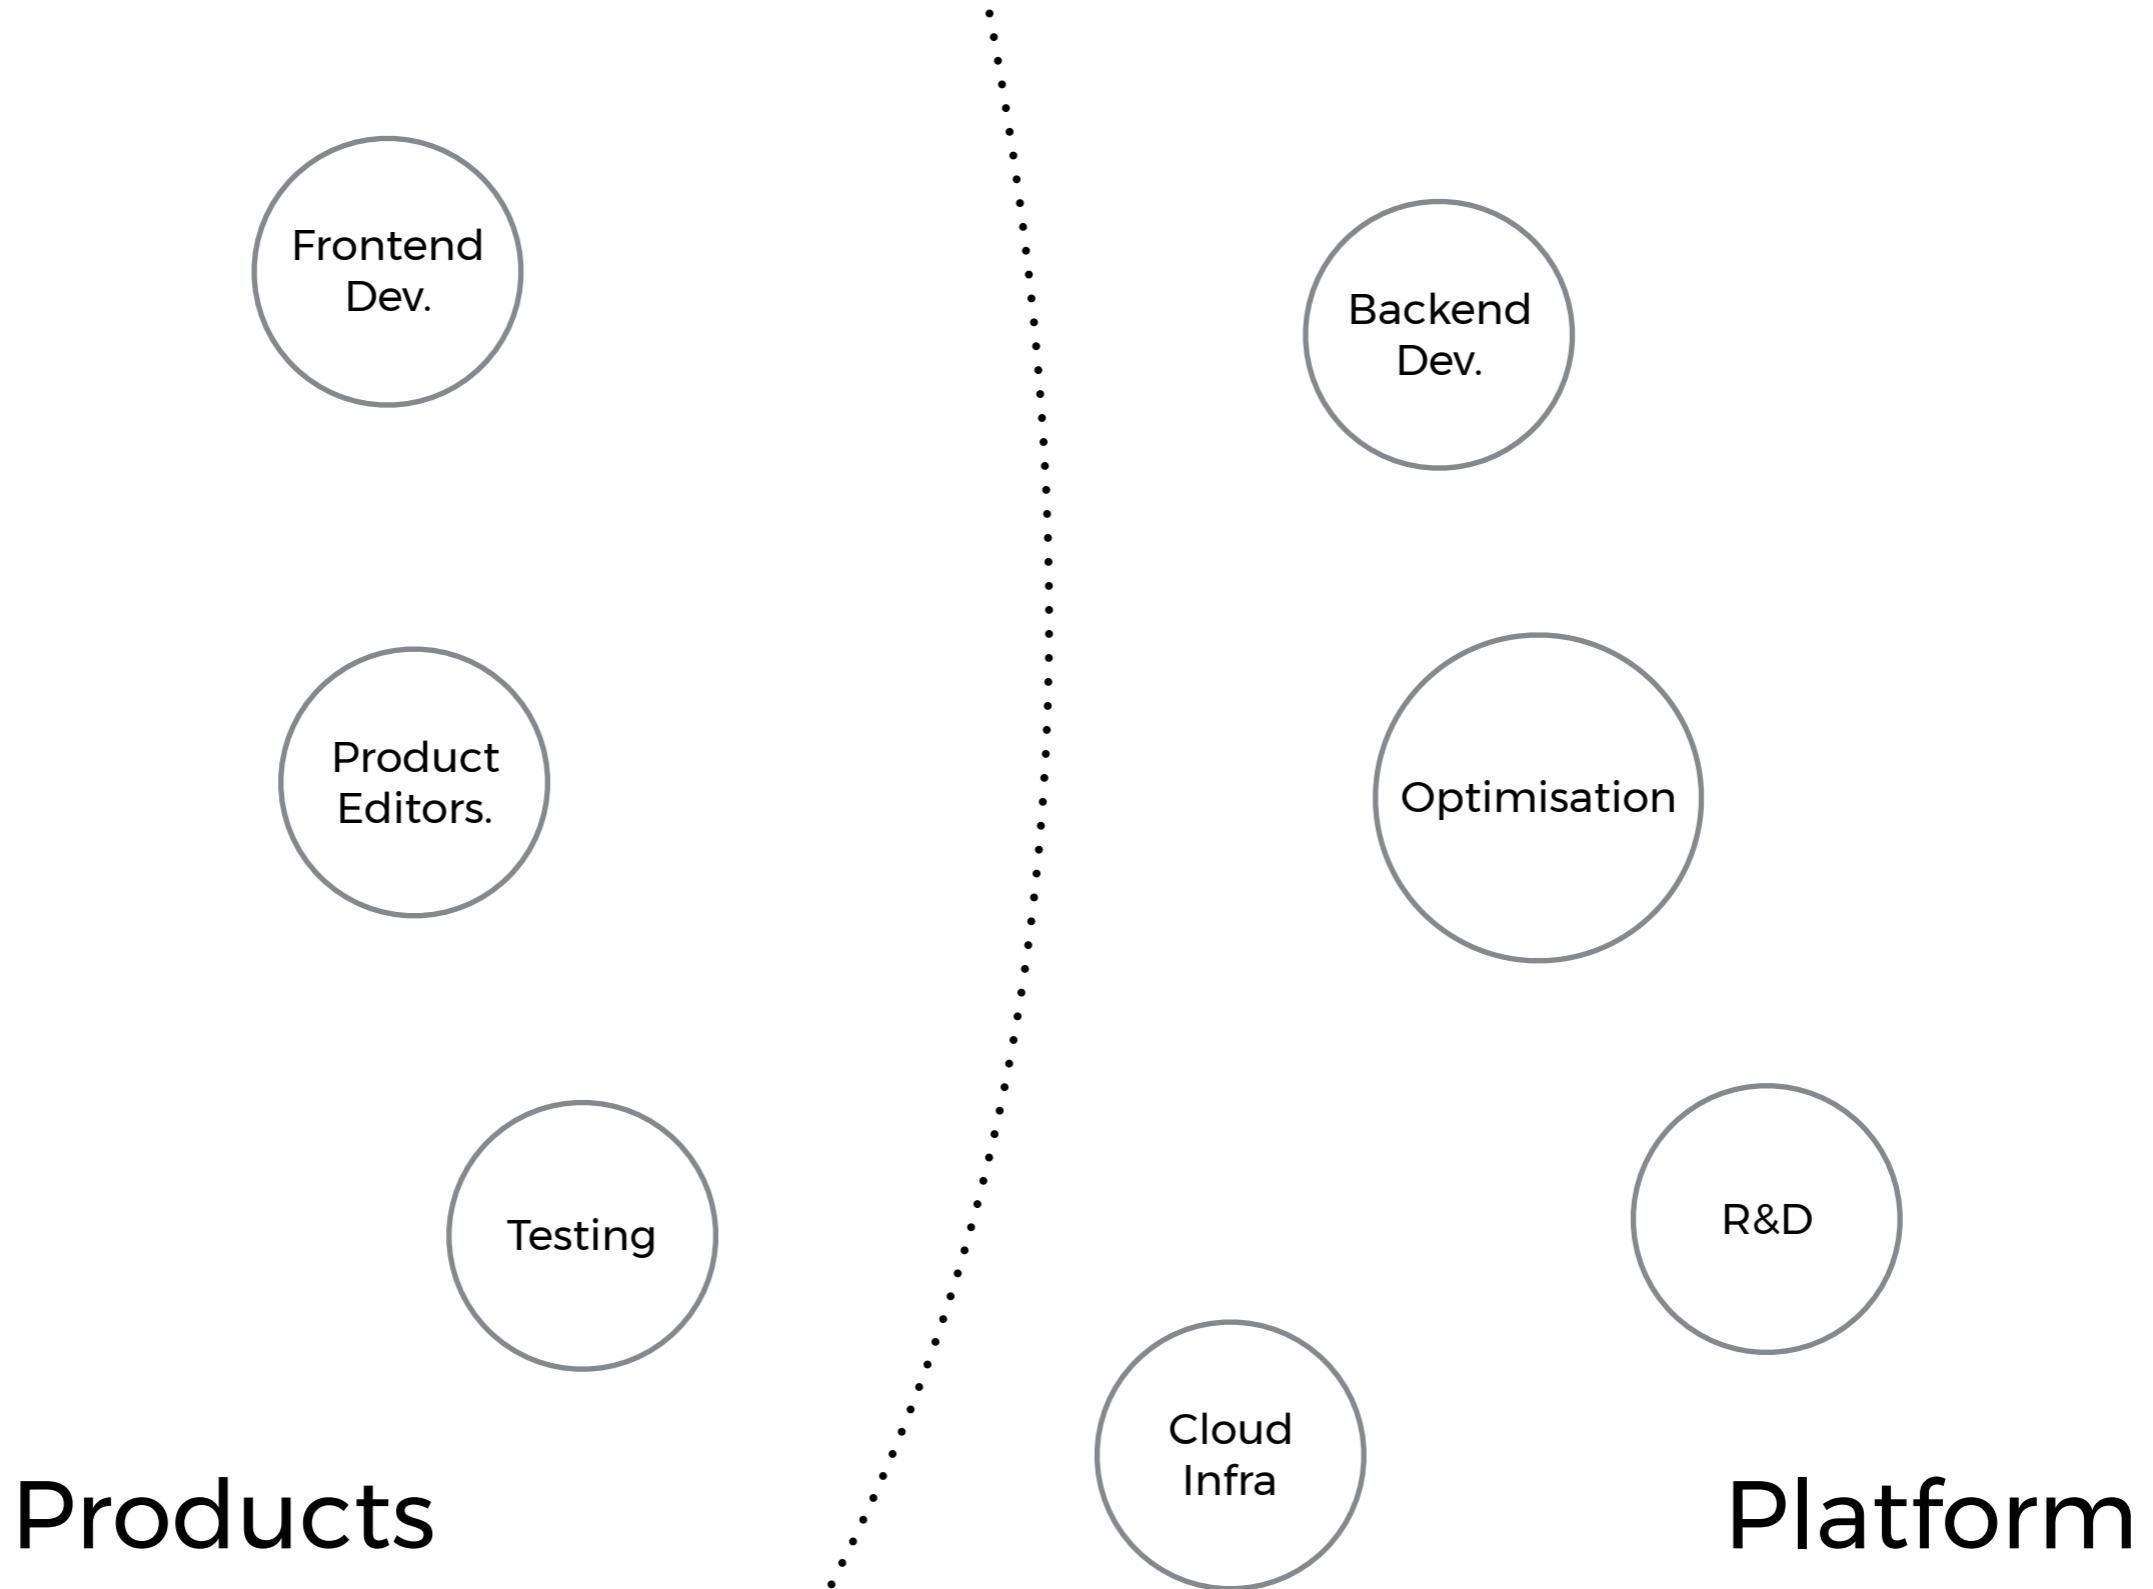
# Our team



HQ in Munich  
Offices in Munich / Sophia Antipolis  
People in  
- Berlin, Heidelberg (Germany)  
- London  
- Barcelona  
- Winterthur (Switzerland)  
- Lille  
- Somewhere in a jet



# Sophia Antipolis - Team Structure (originally)



# Sophia Antipolis - Team Structure (today)

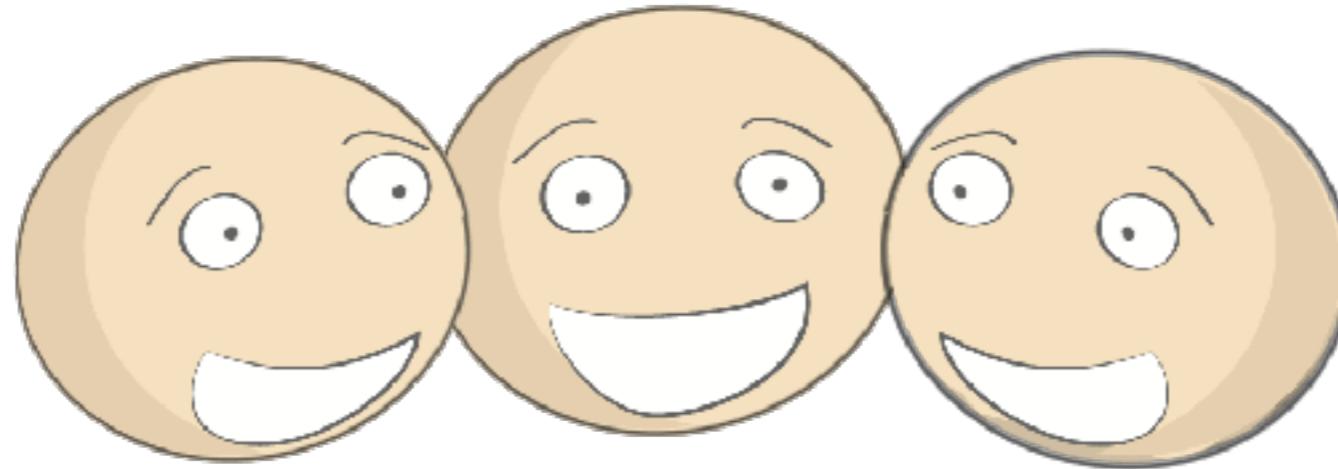
Things to do



People to do  
them

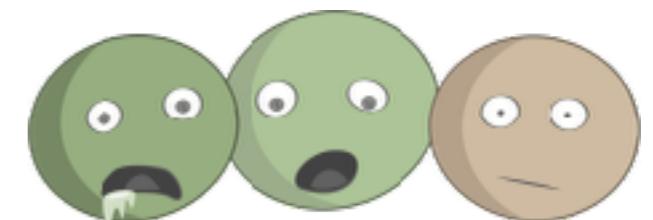
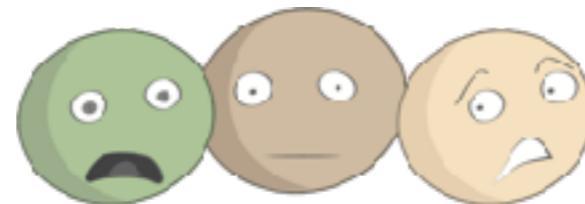
# Dev Methodology

**WE THINK THAT  
BLINDLY APPLYING METHODOLOGIES...**



**team of happy developers**

**... SLOWLY TURN THEM INTO ZOMBIES.**



**the walking devs**

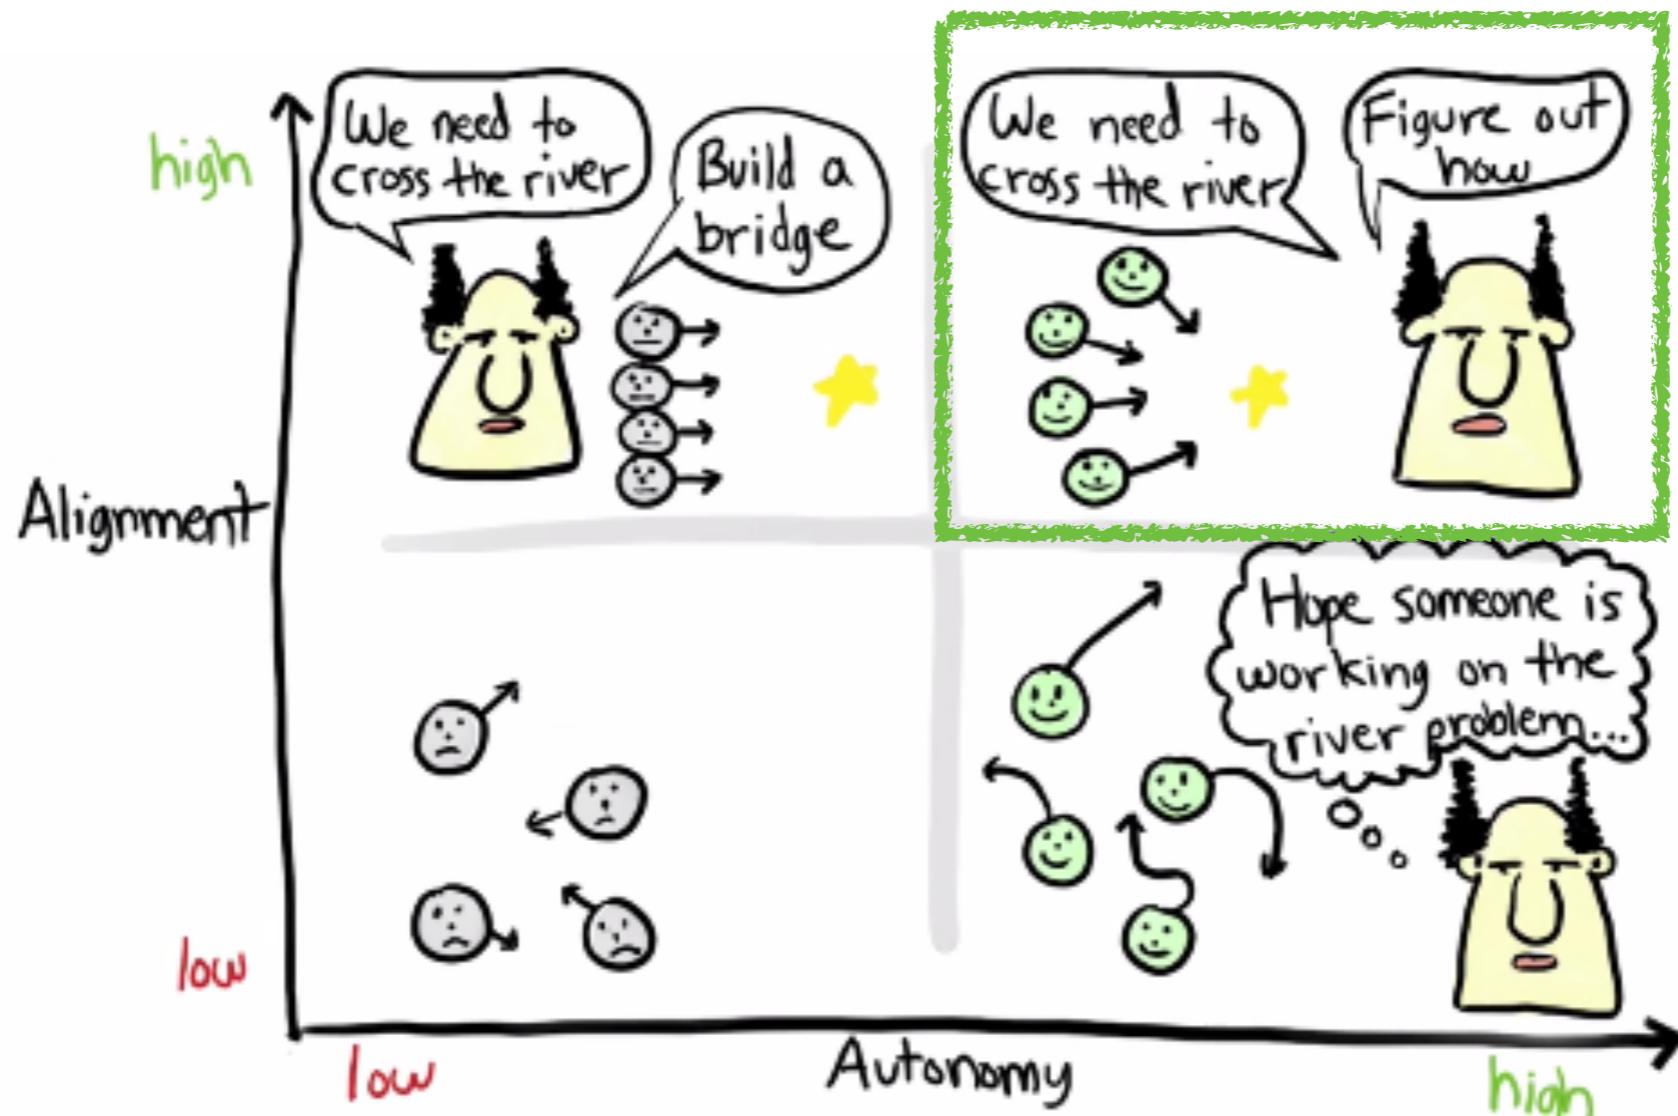


We've got **Things** to do and **People** to do them.

We do **Empower** the people (context, ownership).

We seek for **Autonomy** through **Alignment**.

# Autonomy through Alignment





# **Our world / Our business**

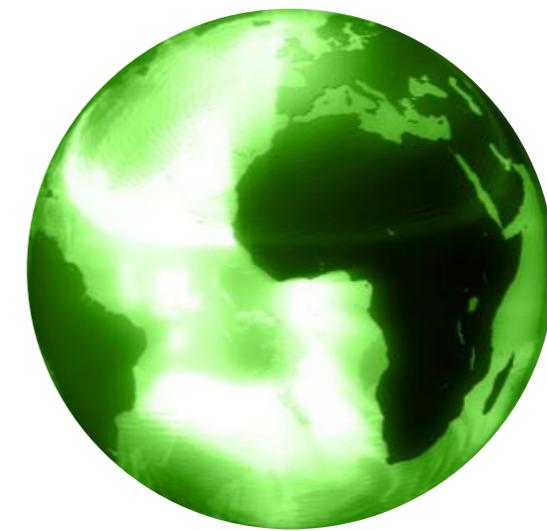
## **Our products**

# Paradigm Shift

## Old World



## New World



### Energy Supply

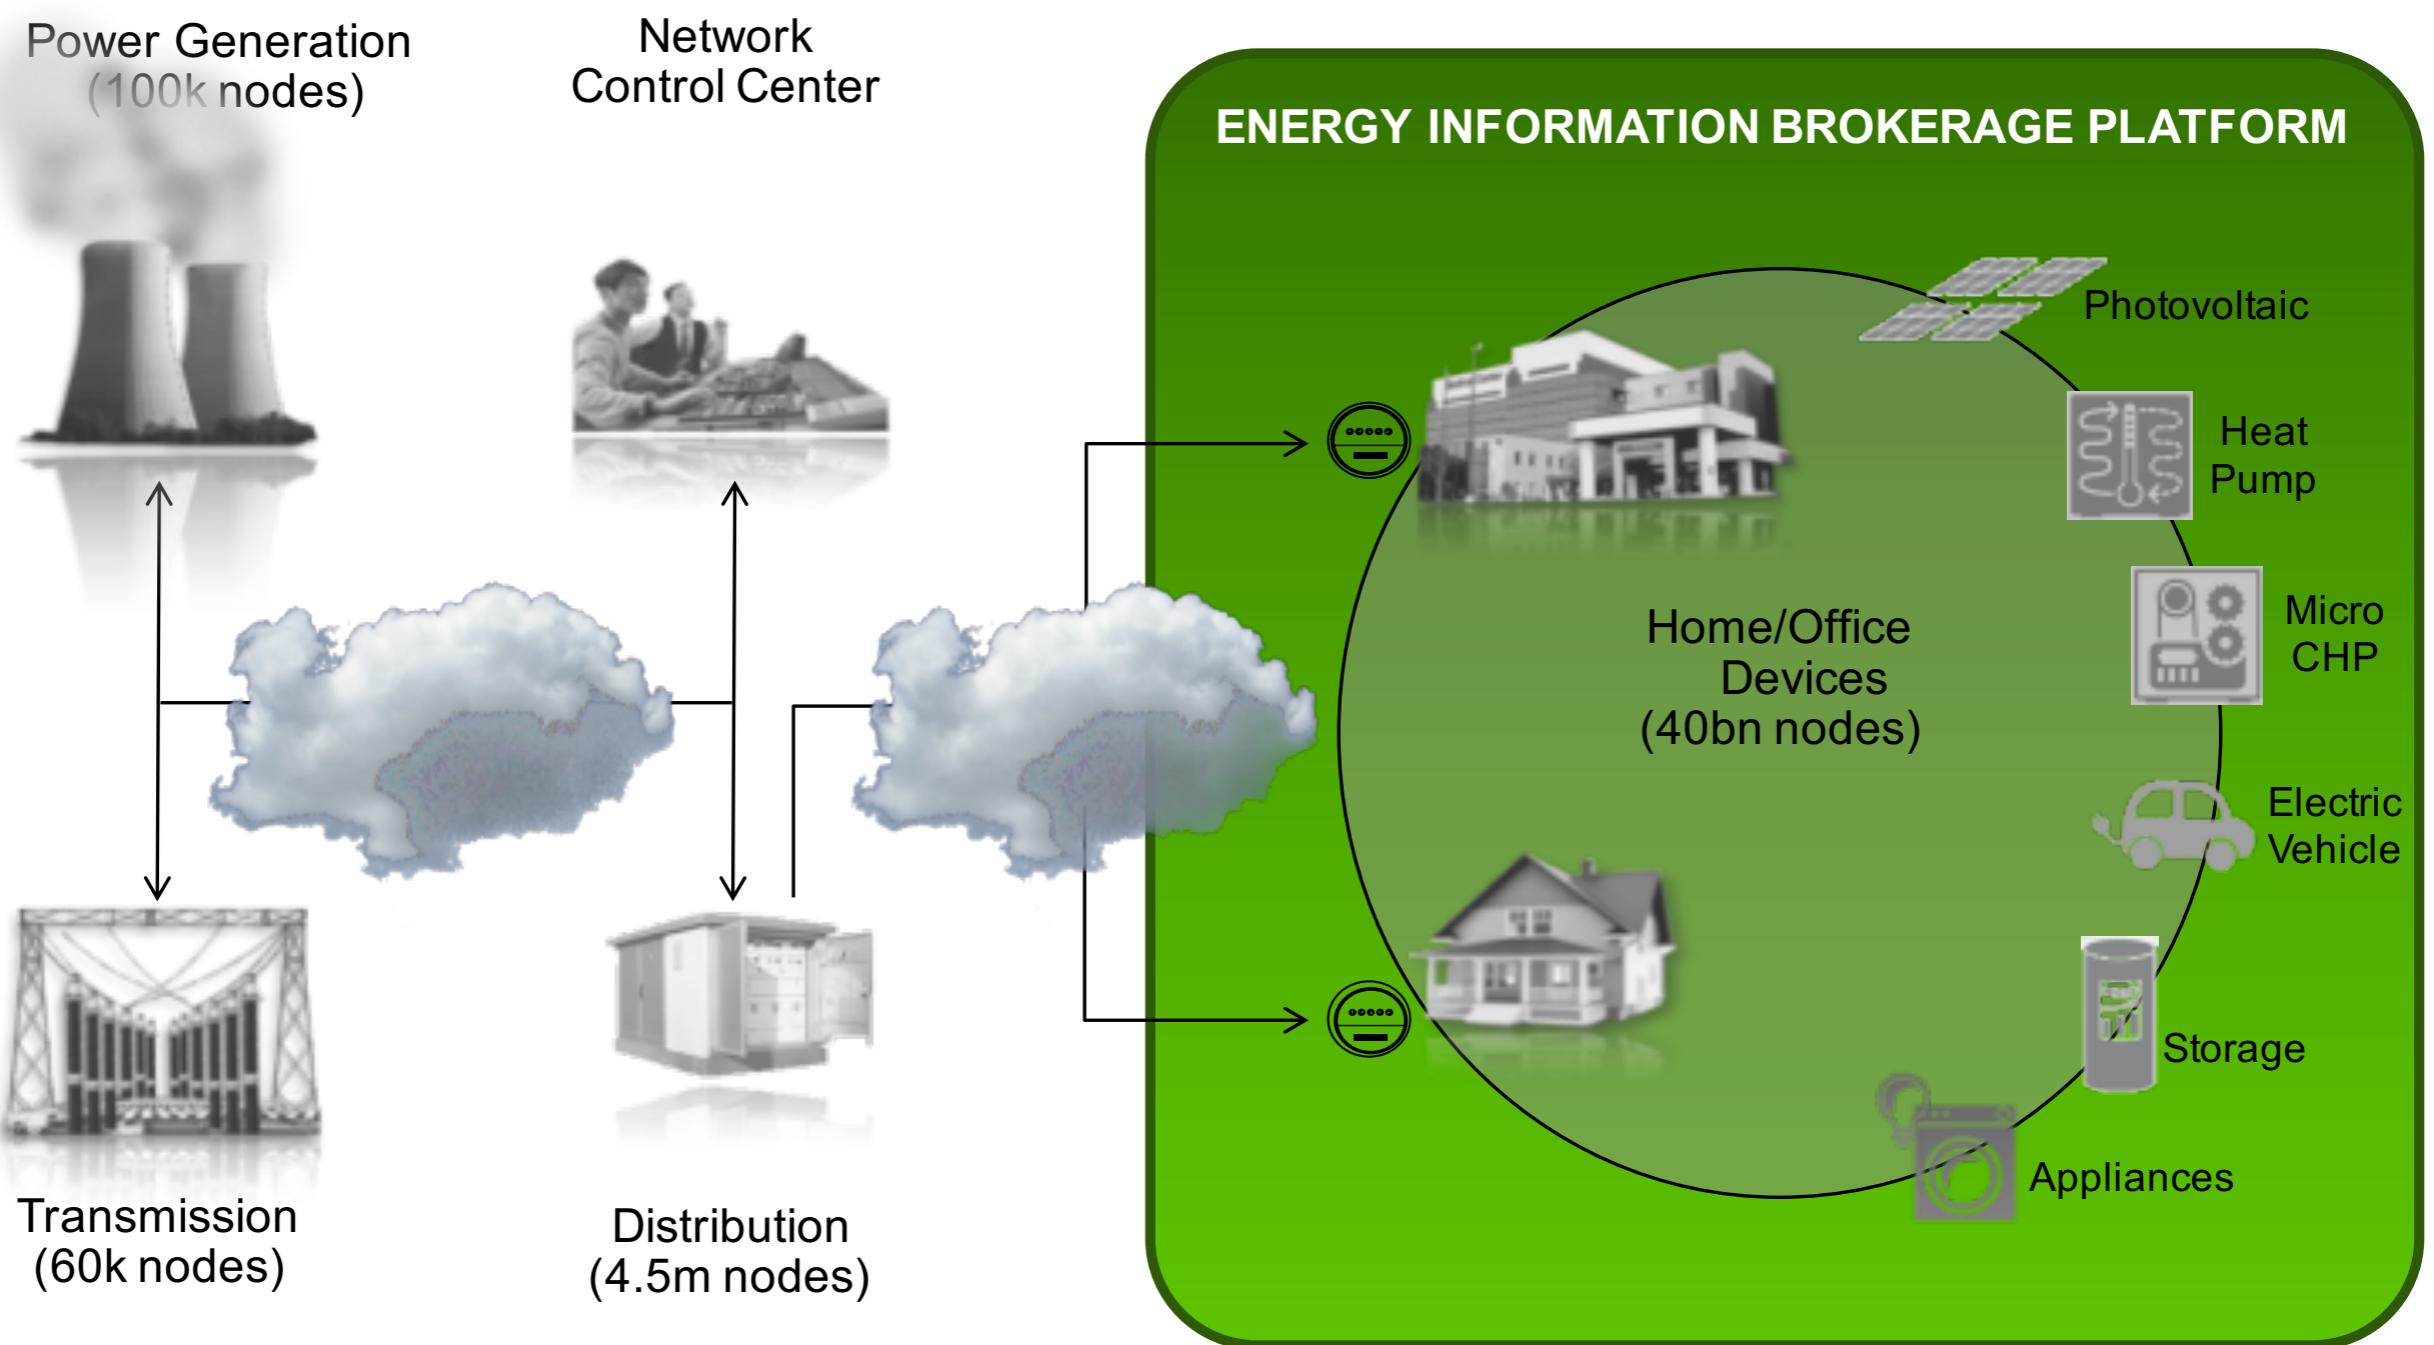
- Central conventional power generation
- Central system management
- Consumers or End Points

### Energy Management

- Central and distributed power generation based on conventional and renewable sources
- Central and distributed system management
- Customers and Prosumers

# Paradigm Shift

Cleantech and Communication Technology leads to a Paradigm Shift in the resources being used





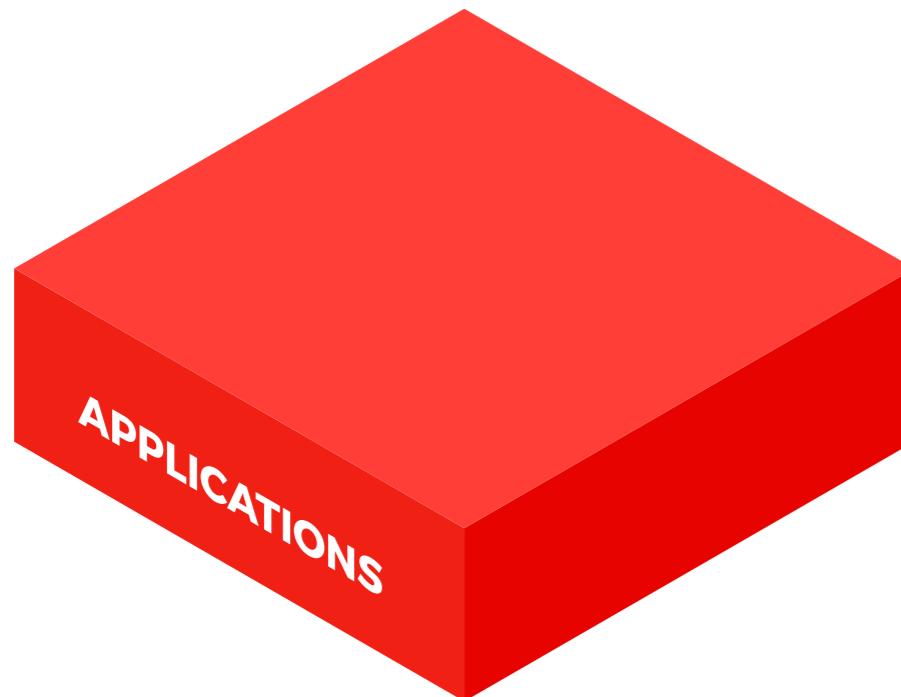
# **GreenCom Networks**

A software-as-a-service company that offers white-label solutions for the utility industry and energy service companies

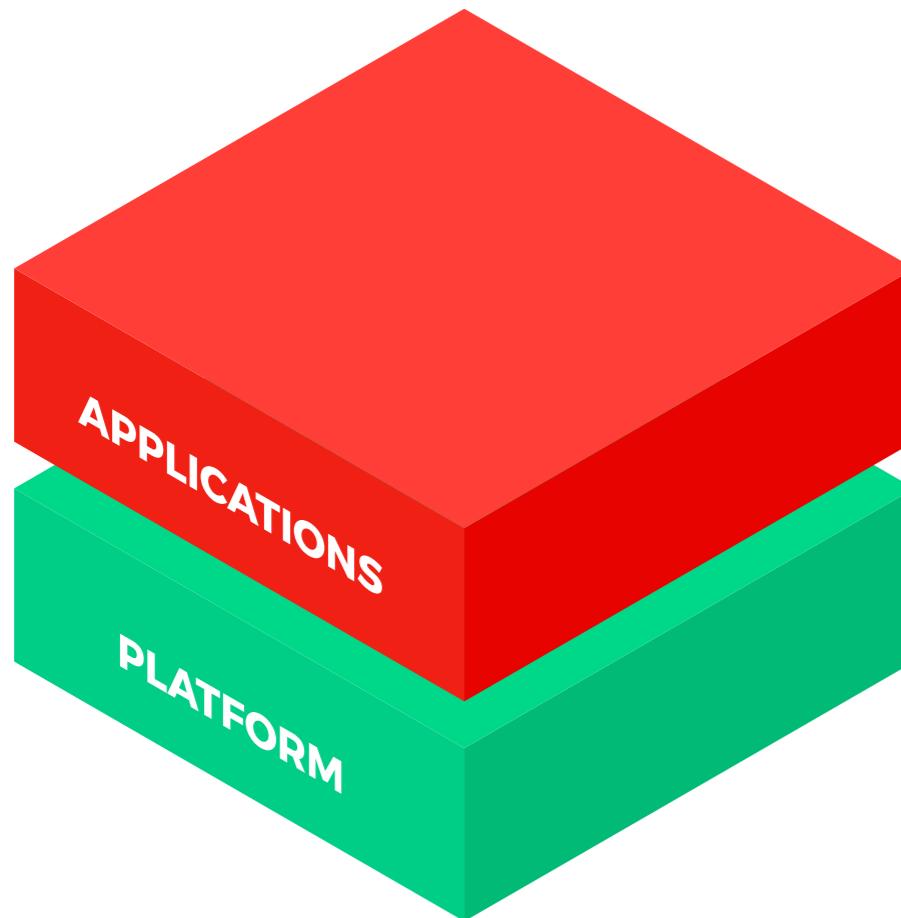


# **GreenCom Networks**

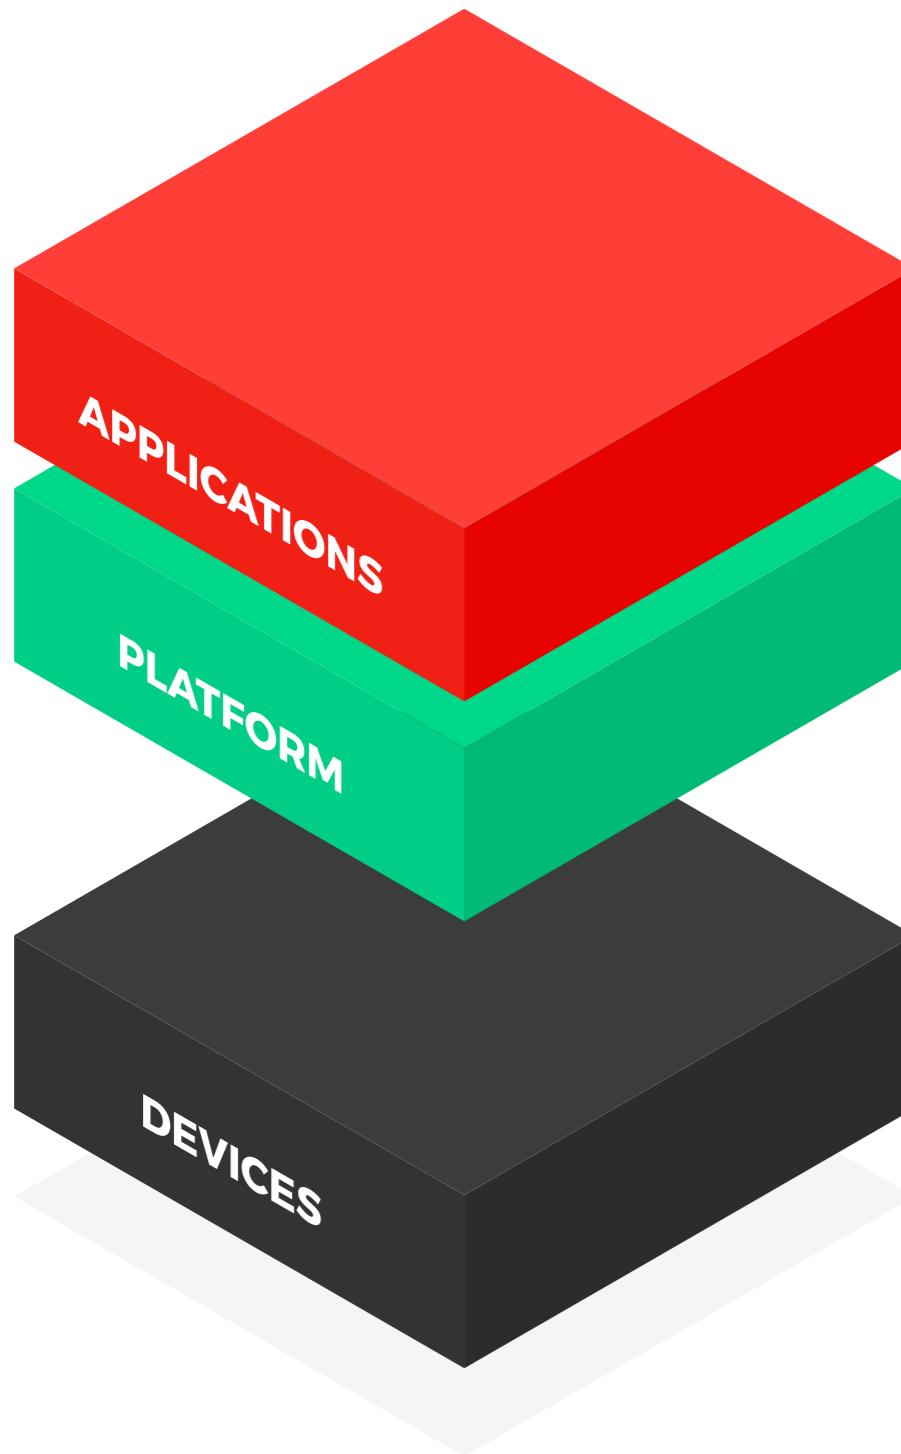
We offer energy related services  
based on our end-to-end IoT  
platform



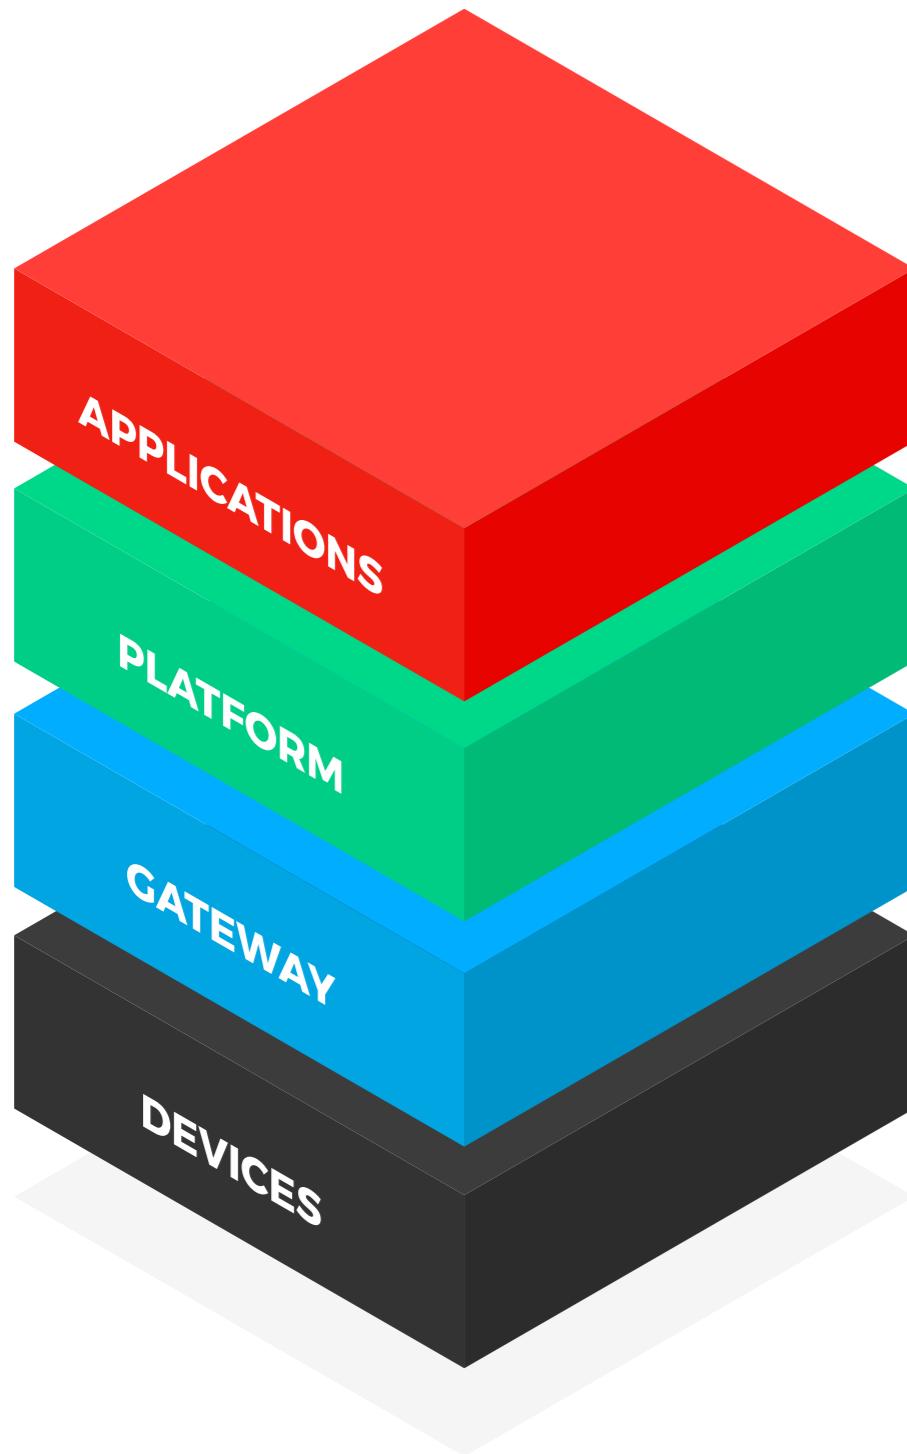
A typical IoT solution  
consists of a set of  
**applications**



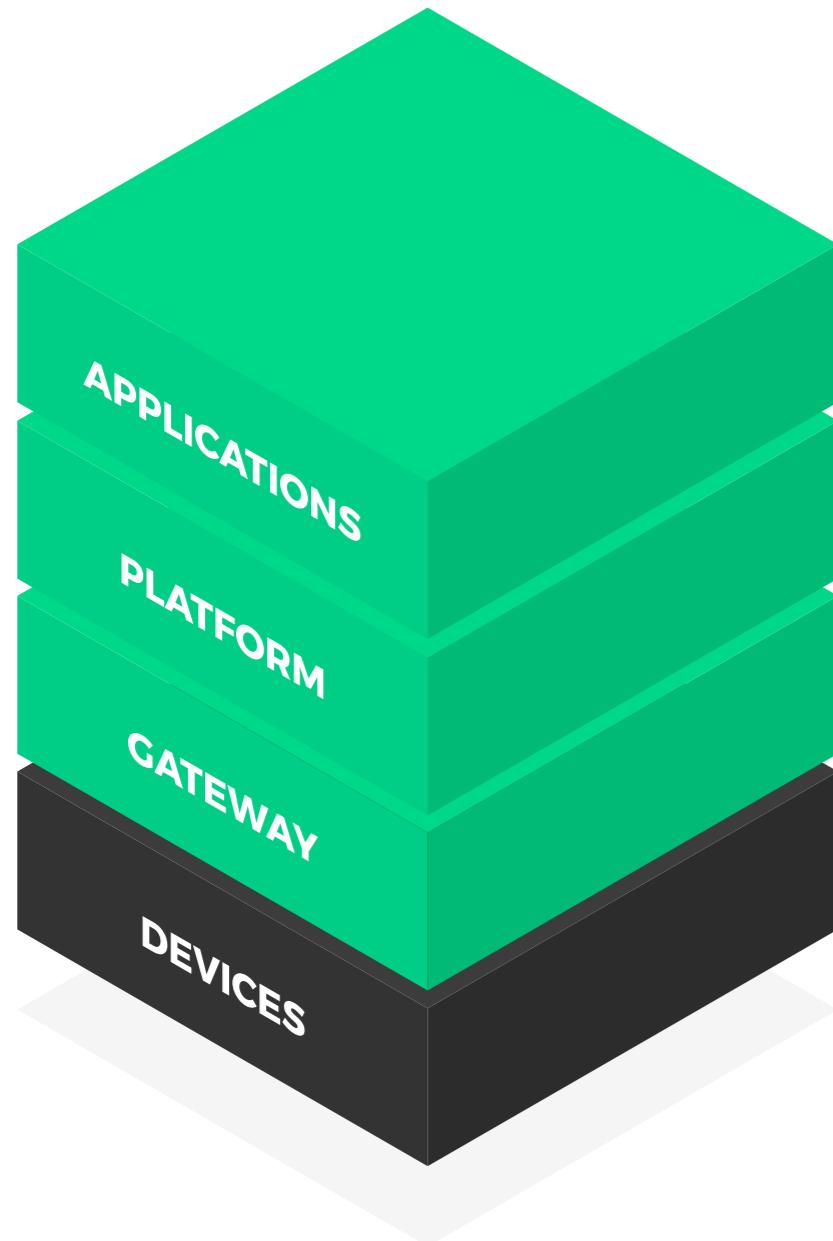
A typical IoT solution  
consists of a set of  
**applications** leveraging  
an IoT **platform**



A typical IoT solution consists of a set of **applications** leveraging an IoT **platform** to connect to a large amount of **devices**



A typical IoT solution consists of a set of **applications** leveraging an IoT **platform** to connect to a large amount of **devices** through some sort of **gateway**.



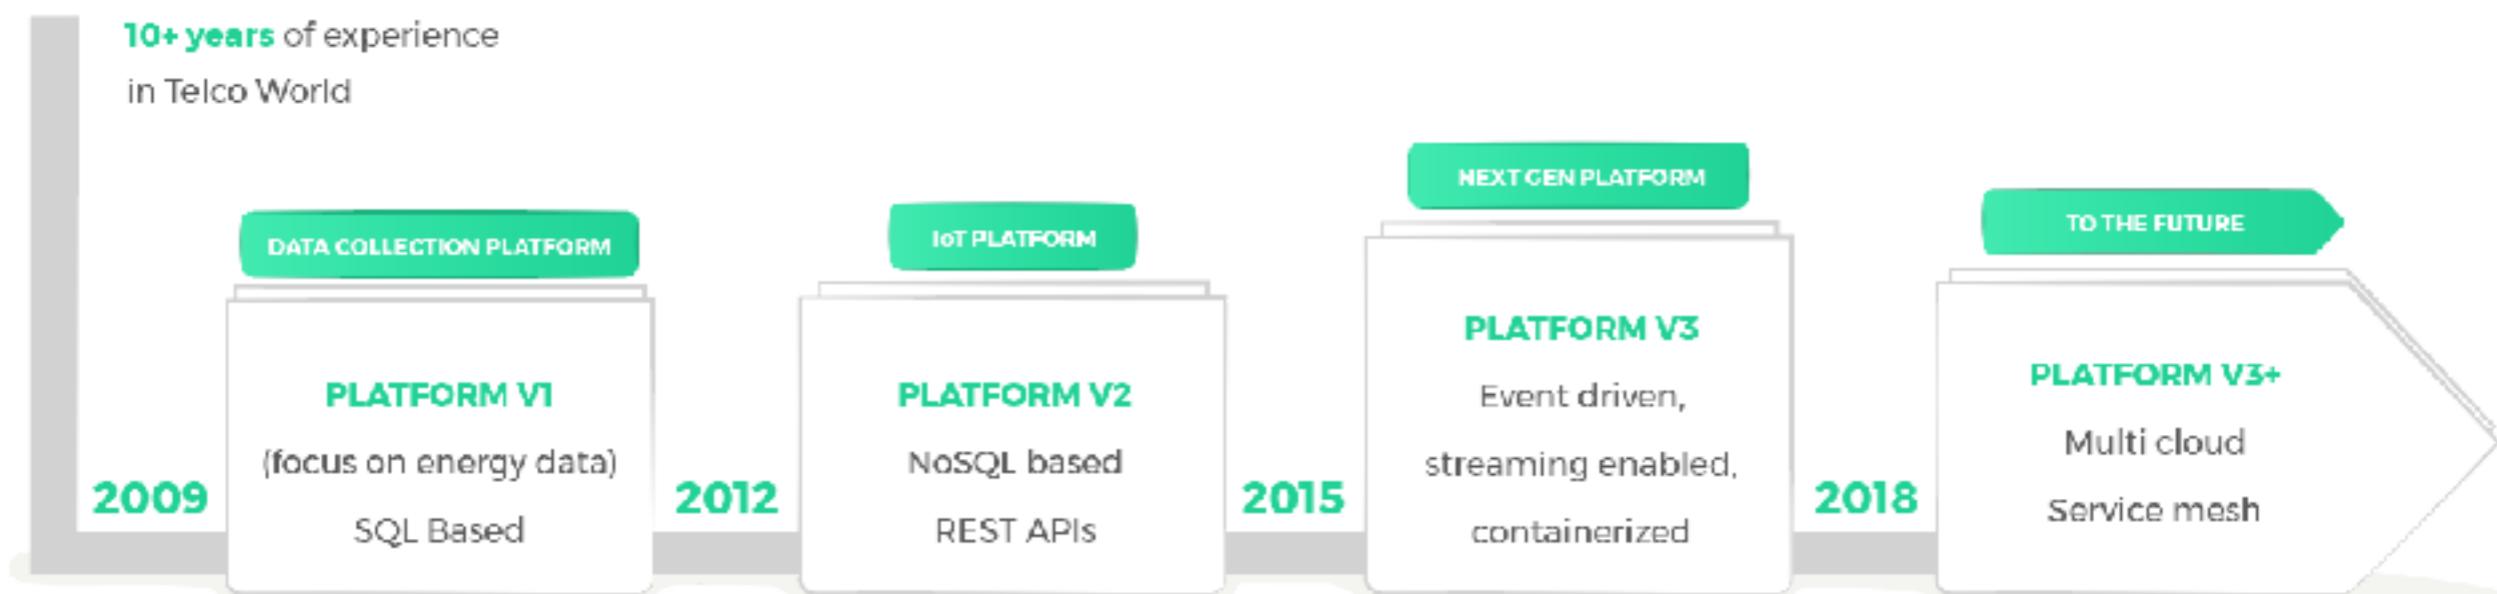
GreenCom Networks through its various offerings is acting in the **application, platform** and **gateway** spaces.

## Energy Information Brokerage Platform

- Energy related implementation of an IoE platform
- Built around a data collection engine and a high-throughput, low-latency message broker
- Offers ability to quickly create and deploy new energy related services
- secured
- multi tenant
- multi cloud provider
- scalable

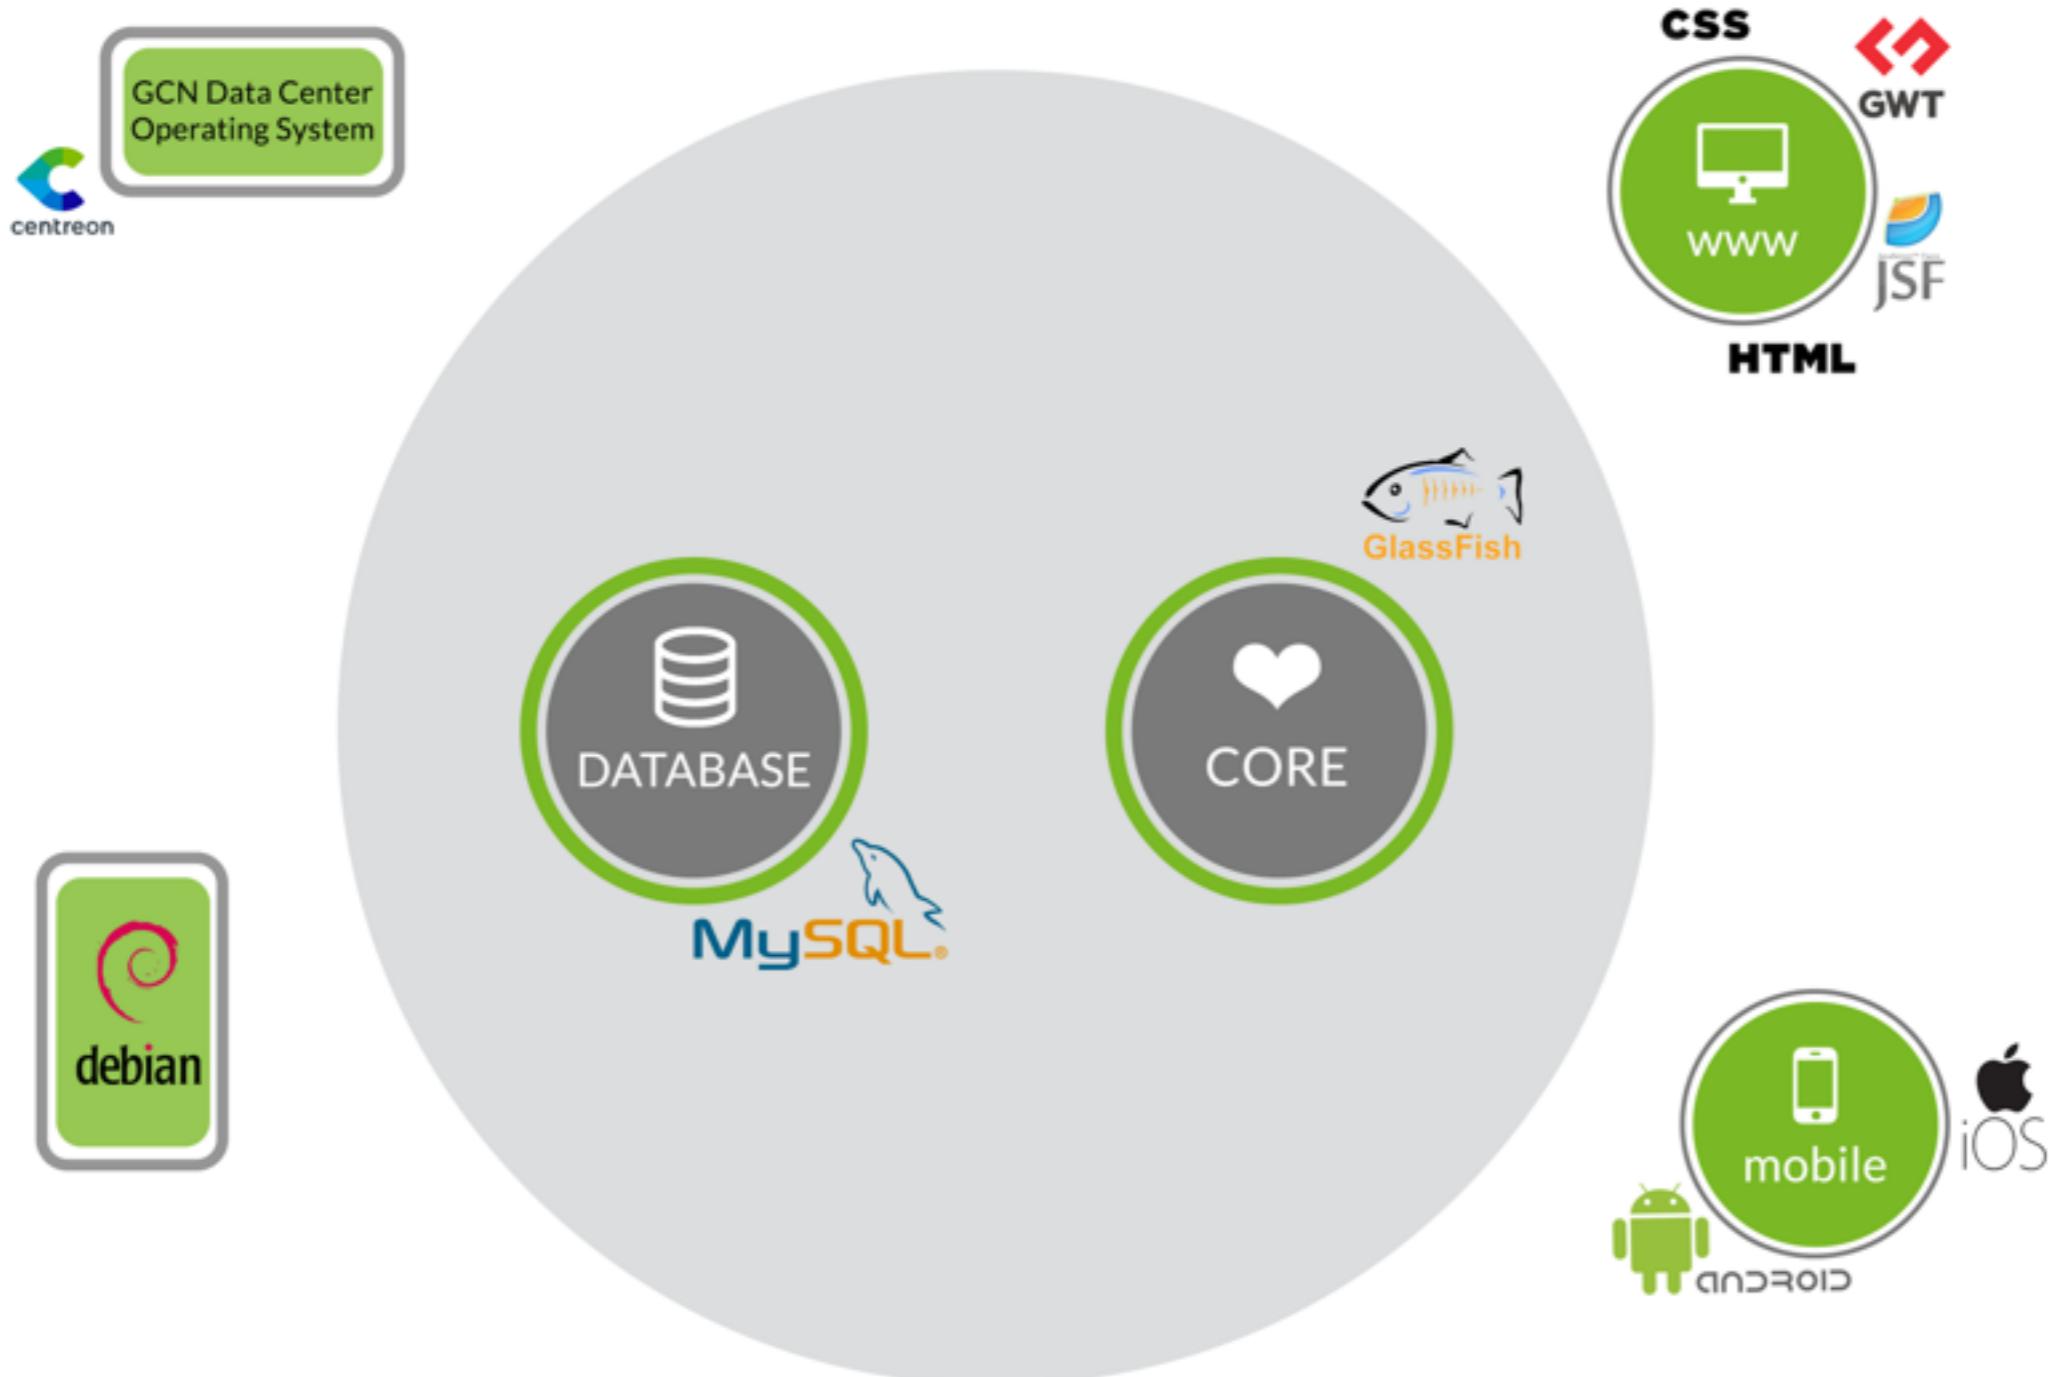
# **A bit of history**

# EIBP Roadmap (past to present)

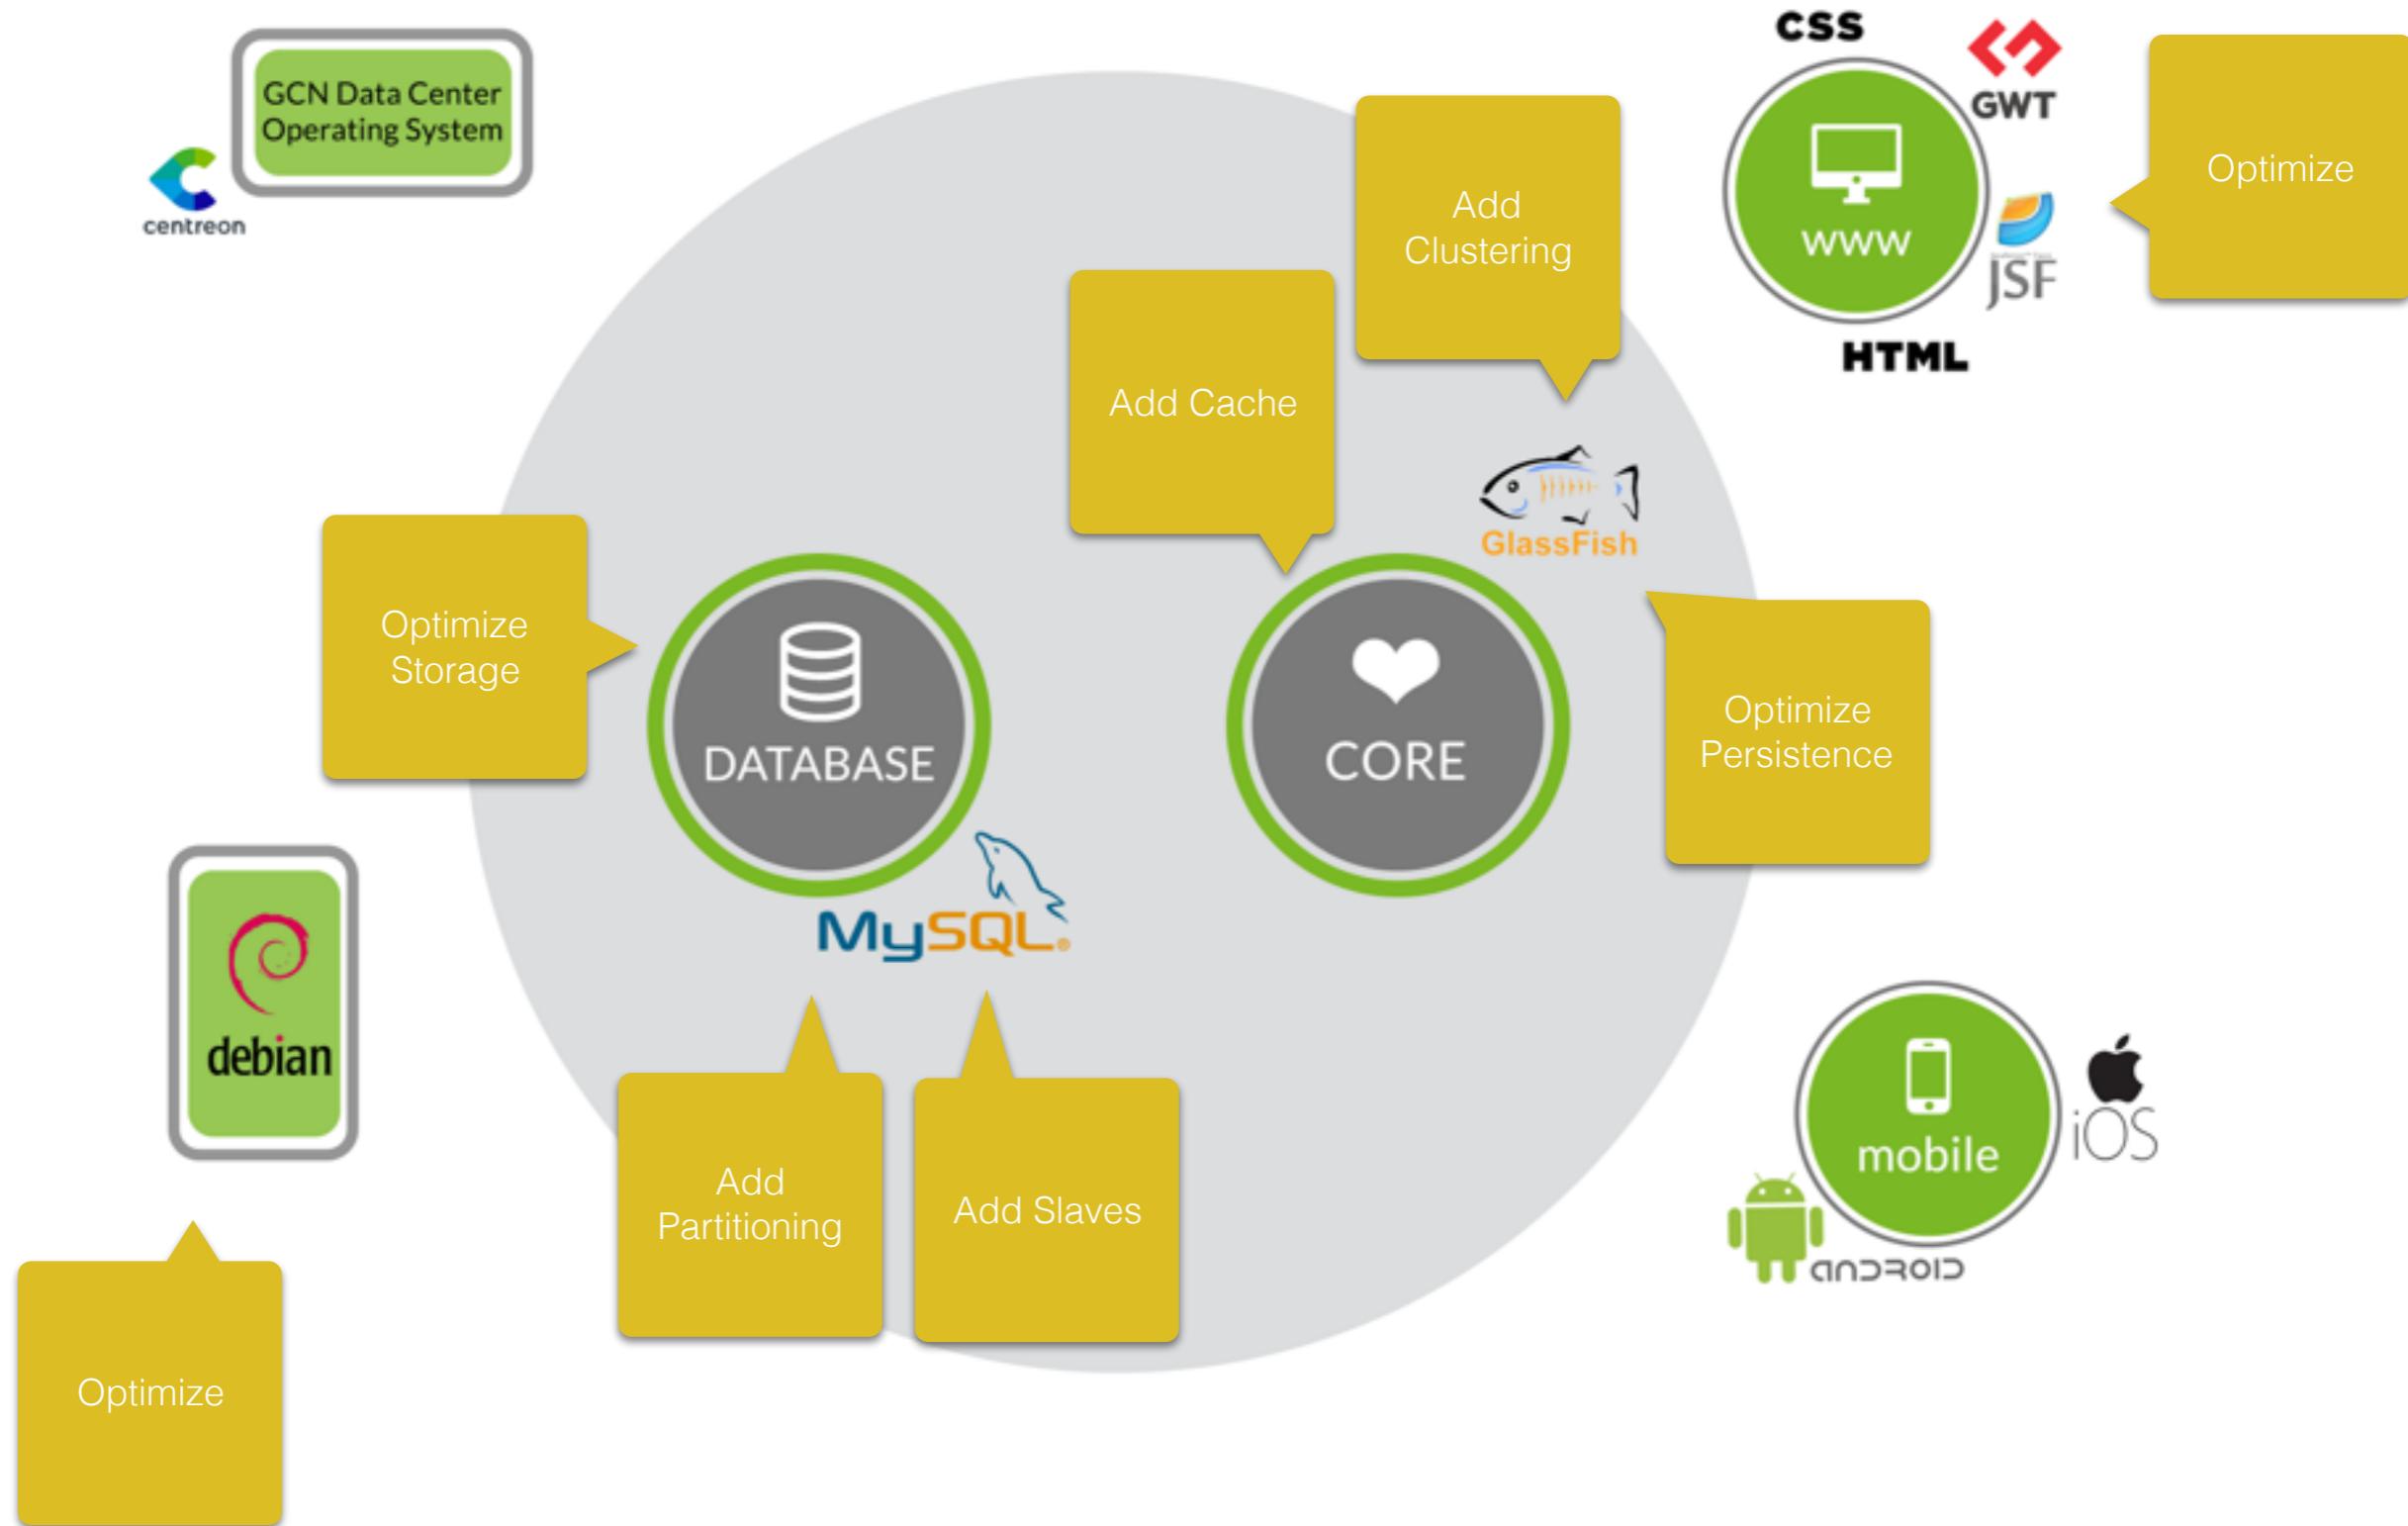


# EIBP Platform V1

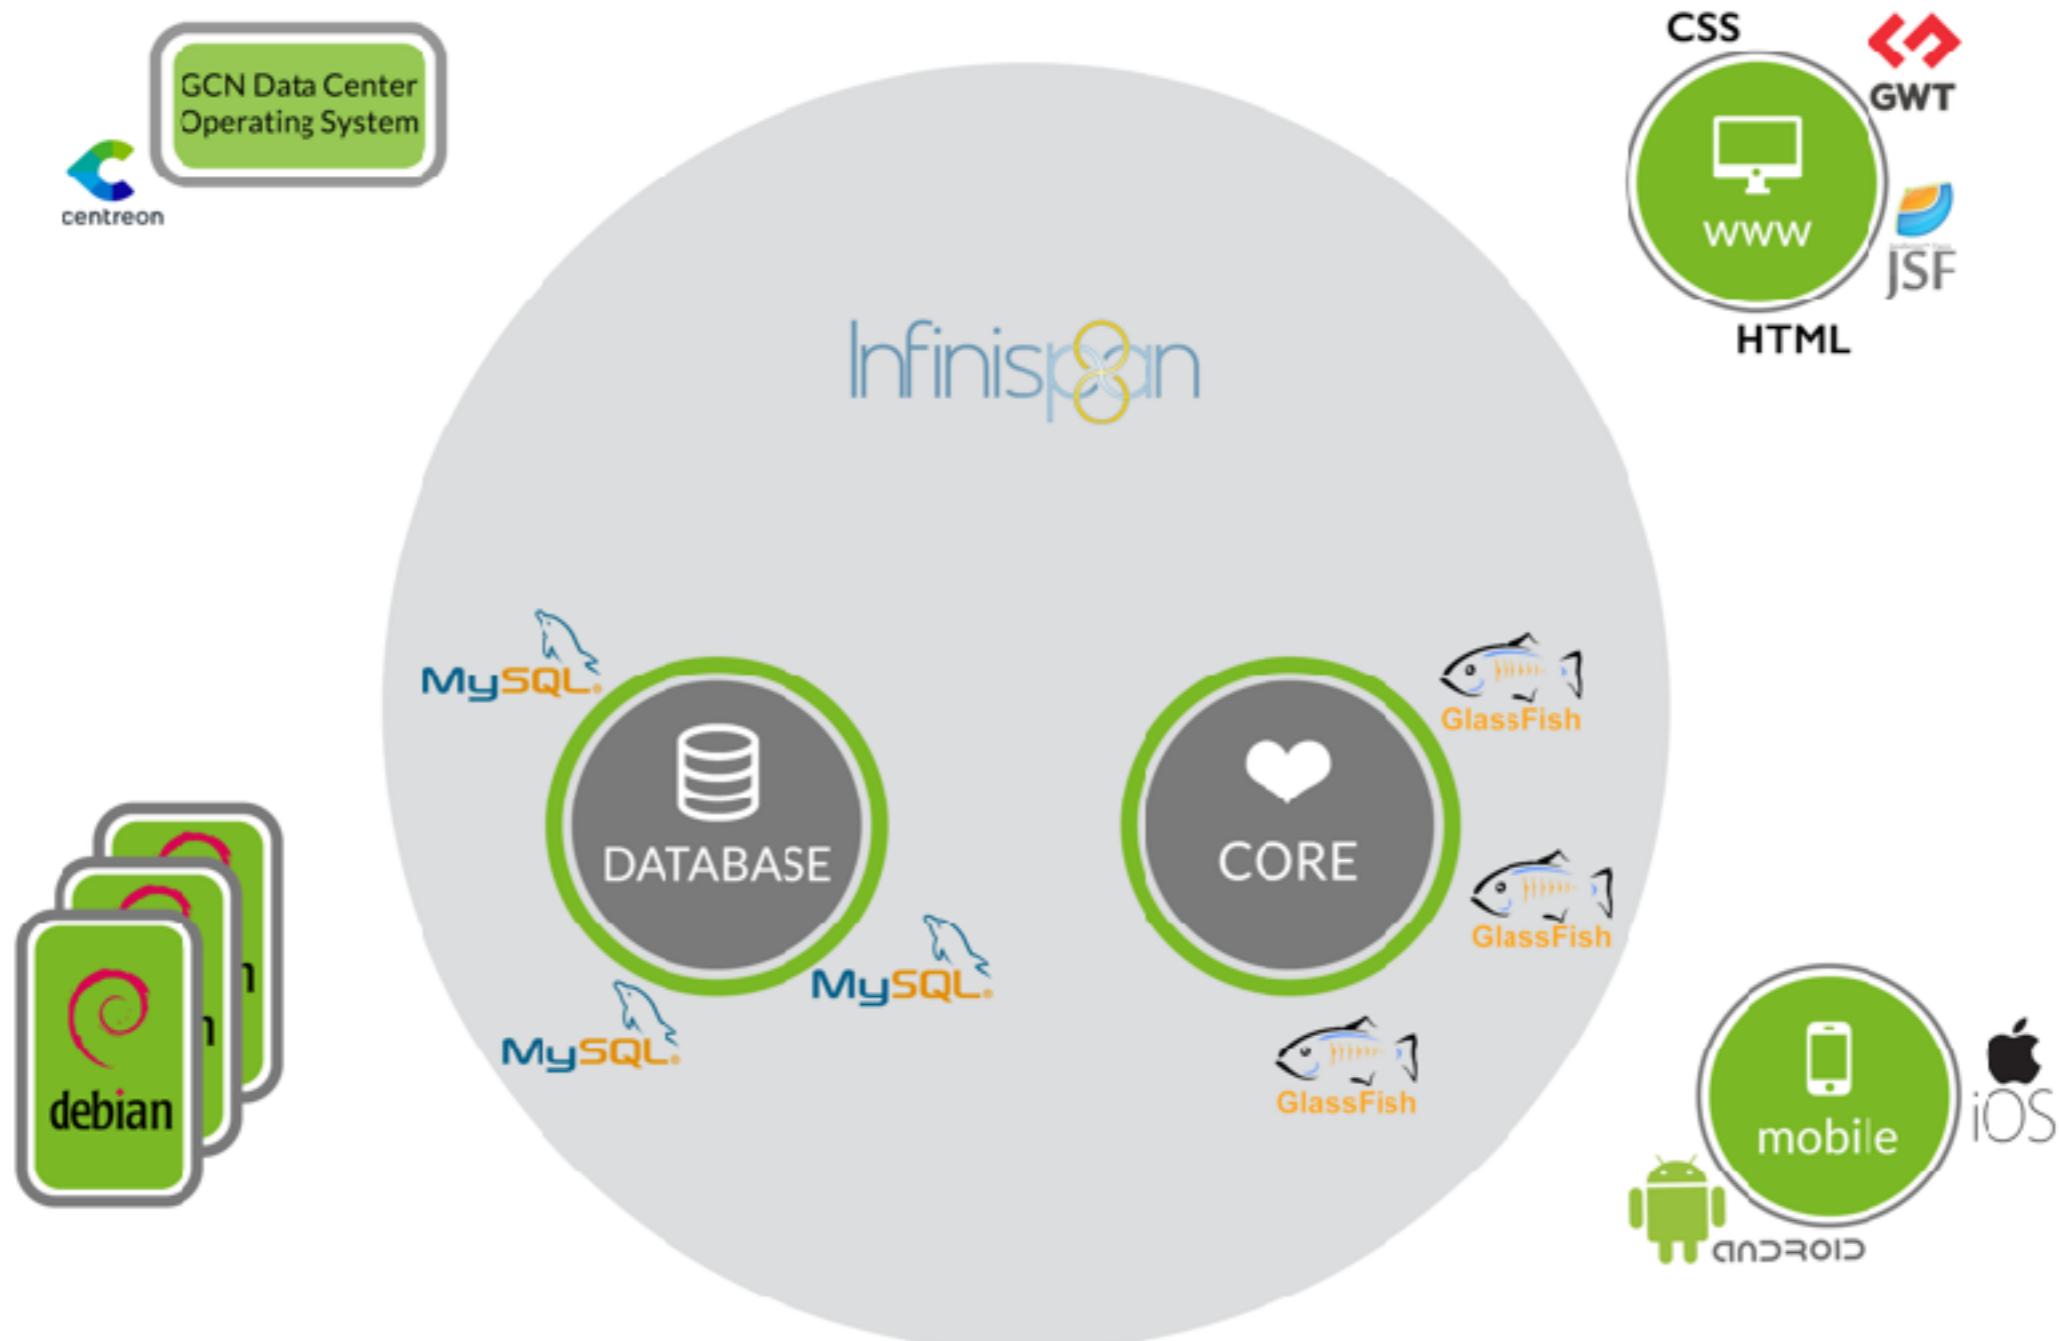
**GREENcom**  
NETWORKS



# EIBP Platform V1 ... after a while



# EIBP Platform V1 ... after a while



When you start spending more time optimising your internals than creating value for your customers, it's time to revise the architecture...

Let's  
~~revise~~  
~~rework~~  
reboot

# What do we need ?



- Time Series Storage / Management
- Metadata Storage / Management
- Scalability
- Replication / Data center awareness (data safety)
- Performance
- Security

- Speed of development
- Agility

## Time Series Storage:

- It will be NoSQL !!!

## Metadata Storage:

- SQL or NoSQL ? → NoSQL

## Which NoSQL database ?

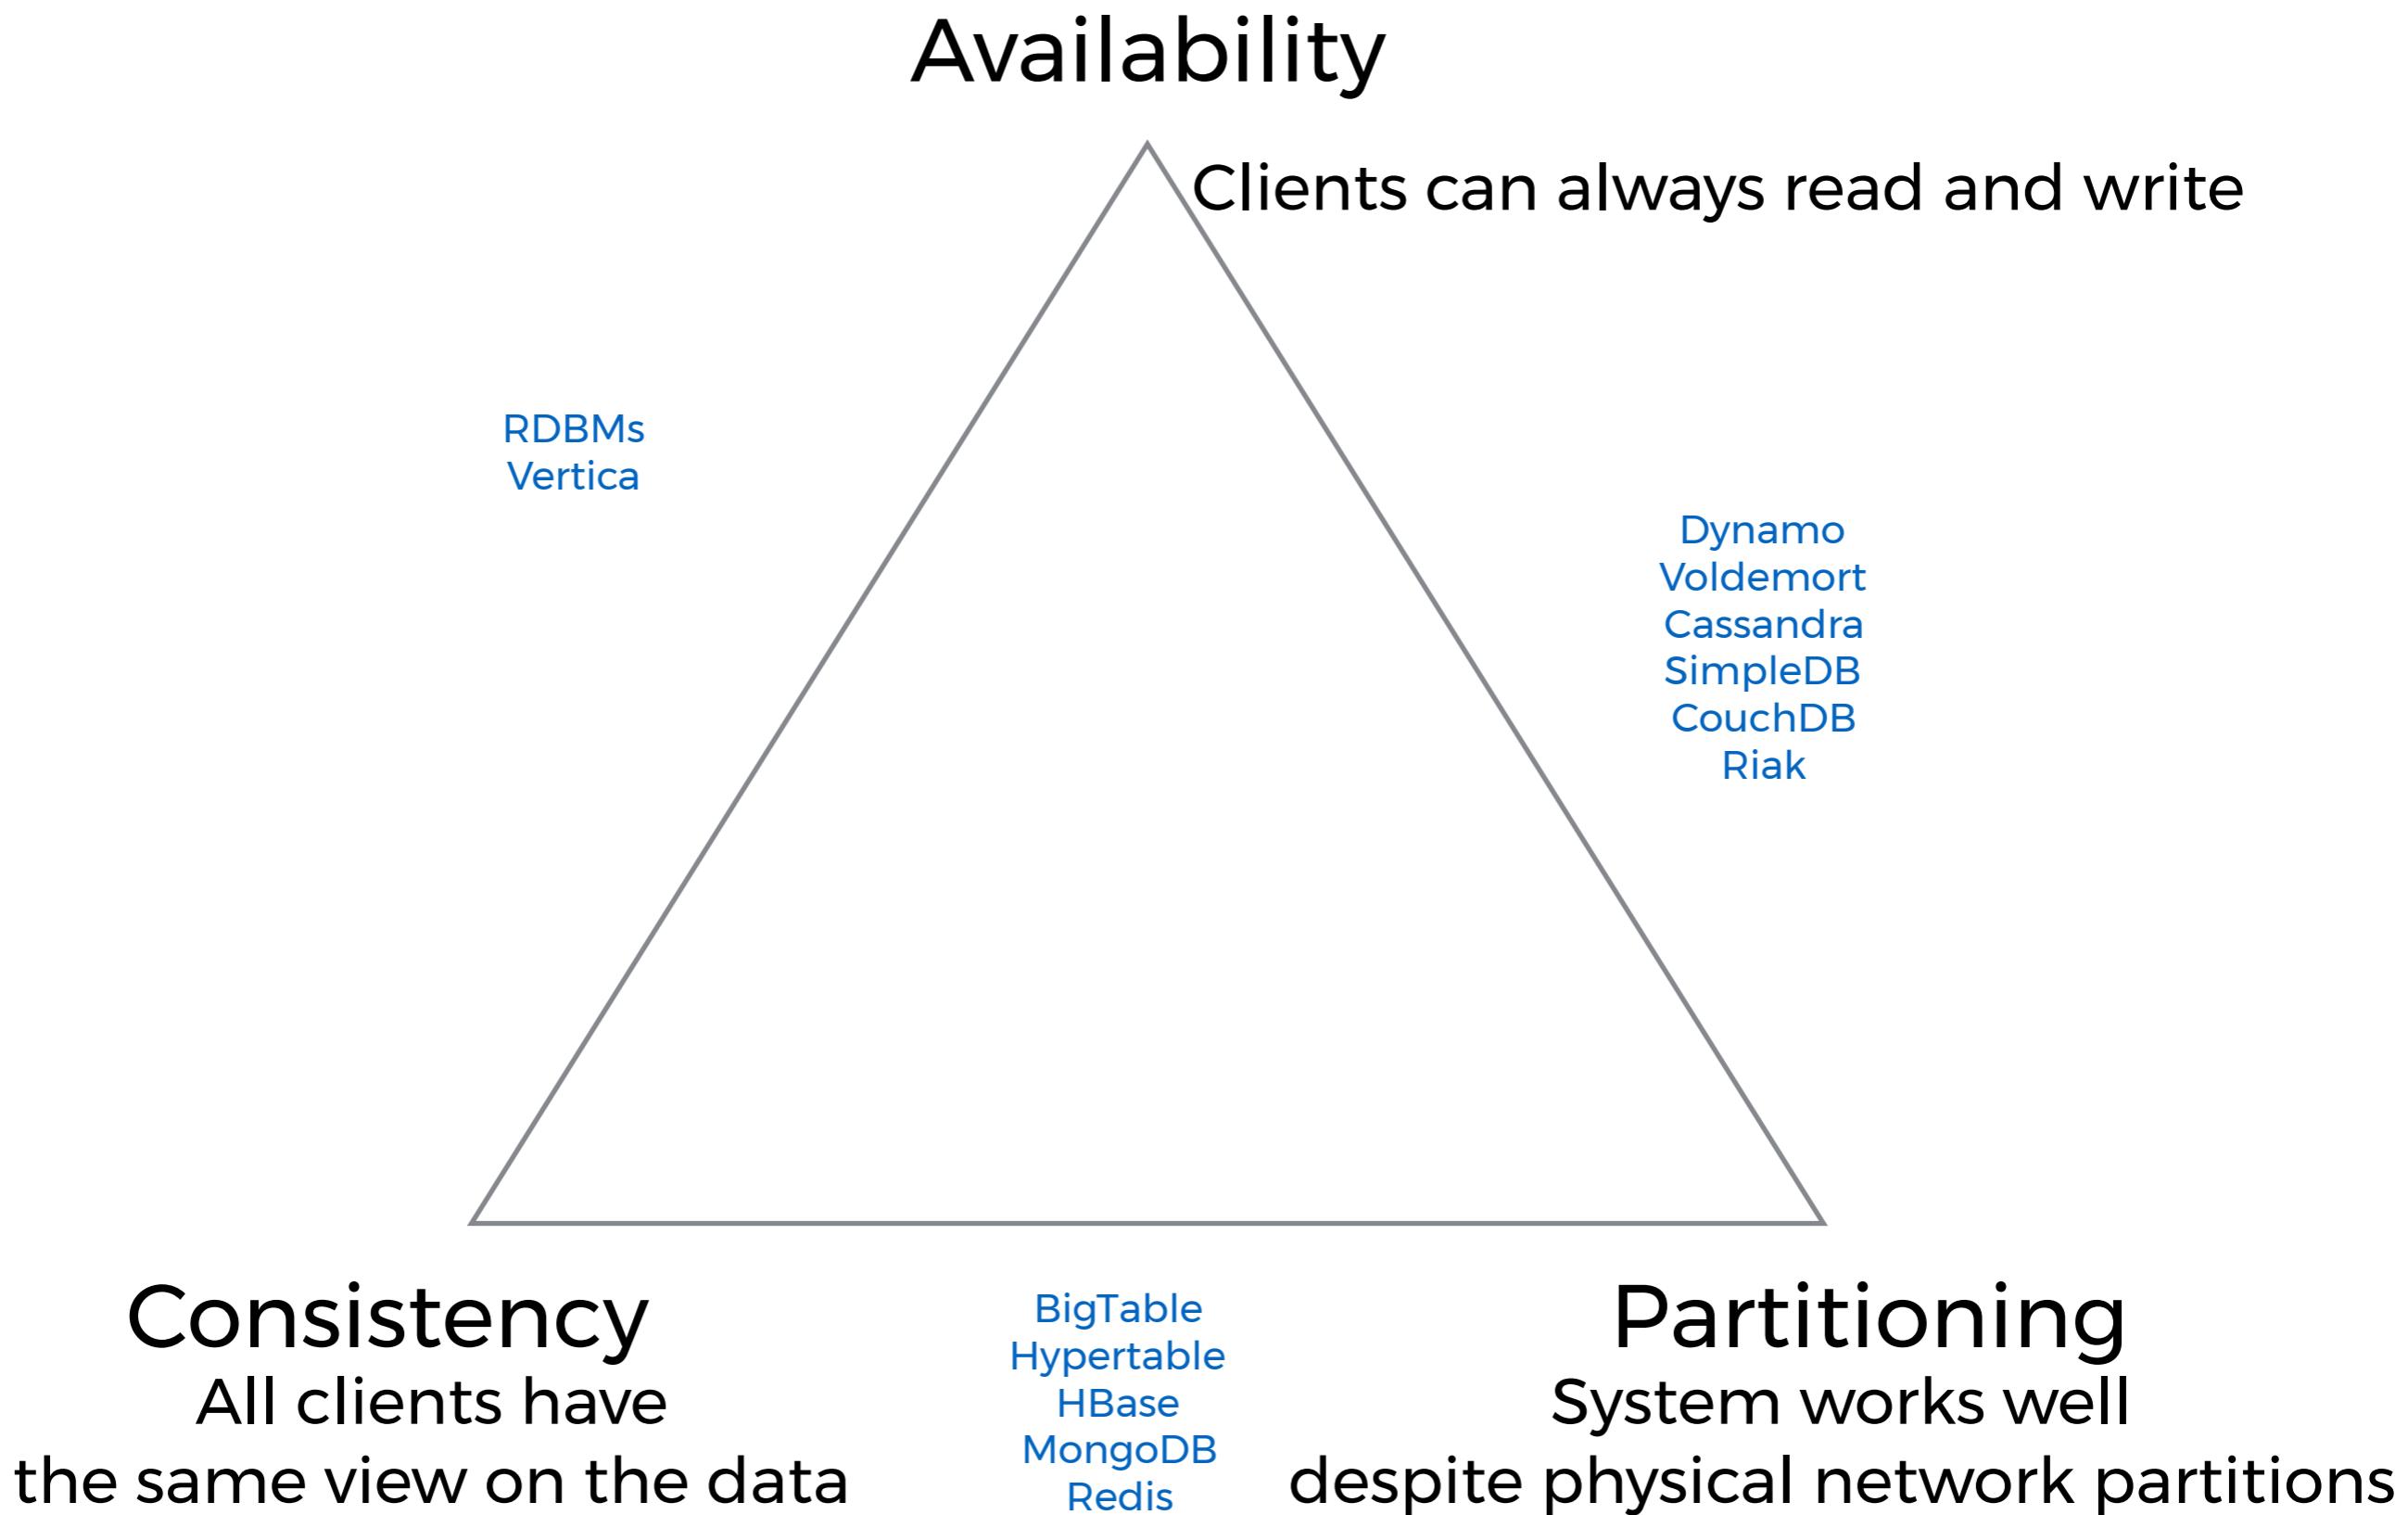
- Lack of maturity of TSDB at that time
- Need to store not only time series but also metadata
- CAP Theorem

# CAP theorem

In theoretical computer science, the CAP theorem, also known as Brewer's theorem, states that it is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- Consistency (all nodes see the same data at the same time)
- Availability (a guarantee that every request receives a response about whether it succeeded or failed)
- Partition tolerance (the system continues to operate despite arbitrary partitioning due to network failures)

# CAP theorem

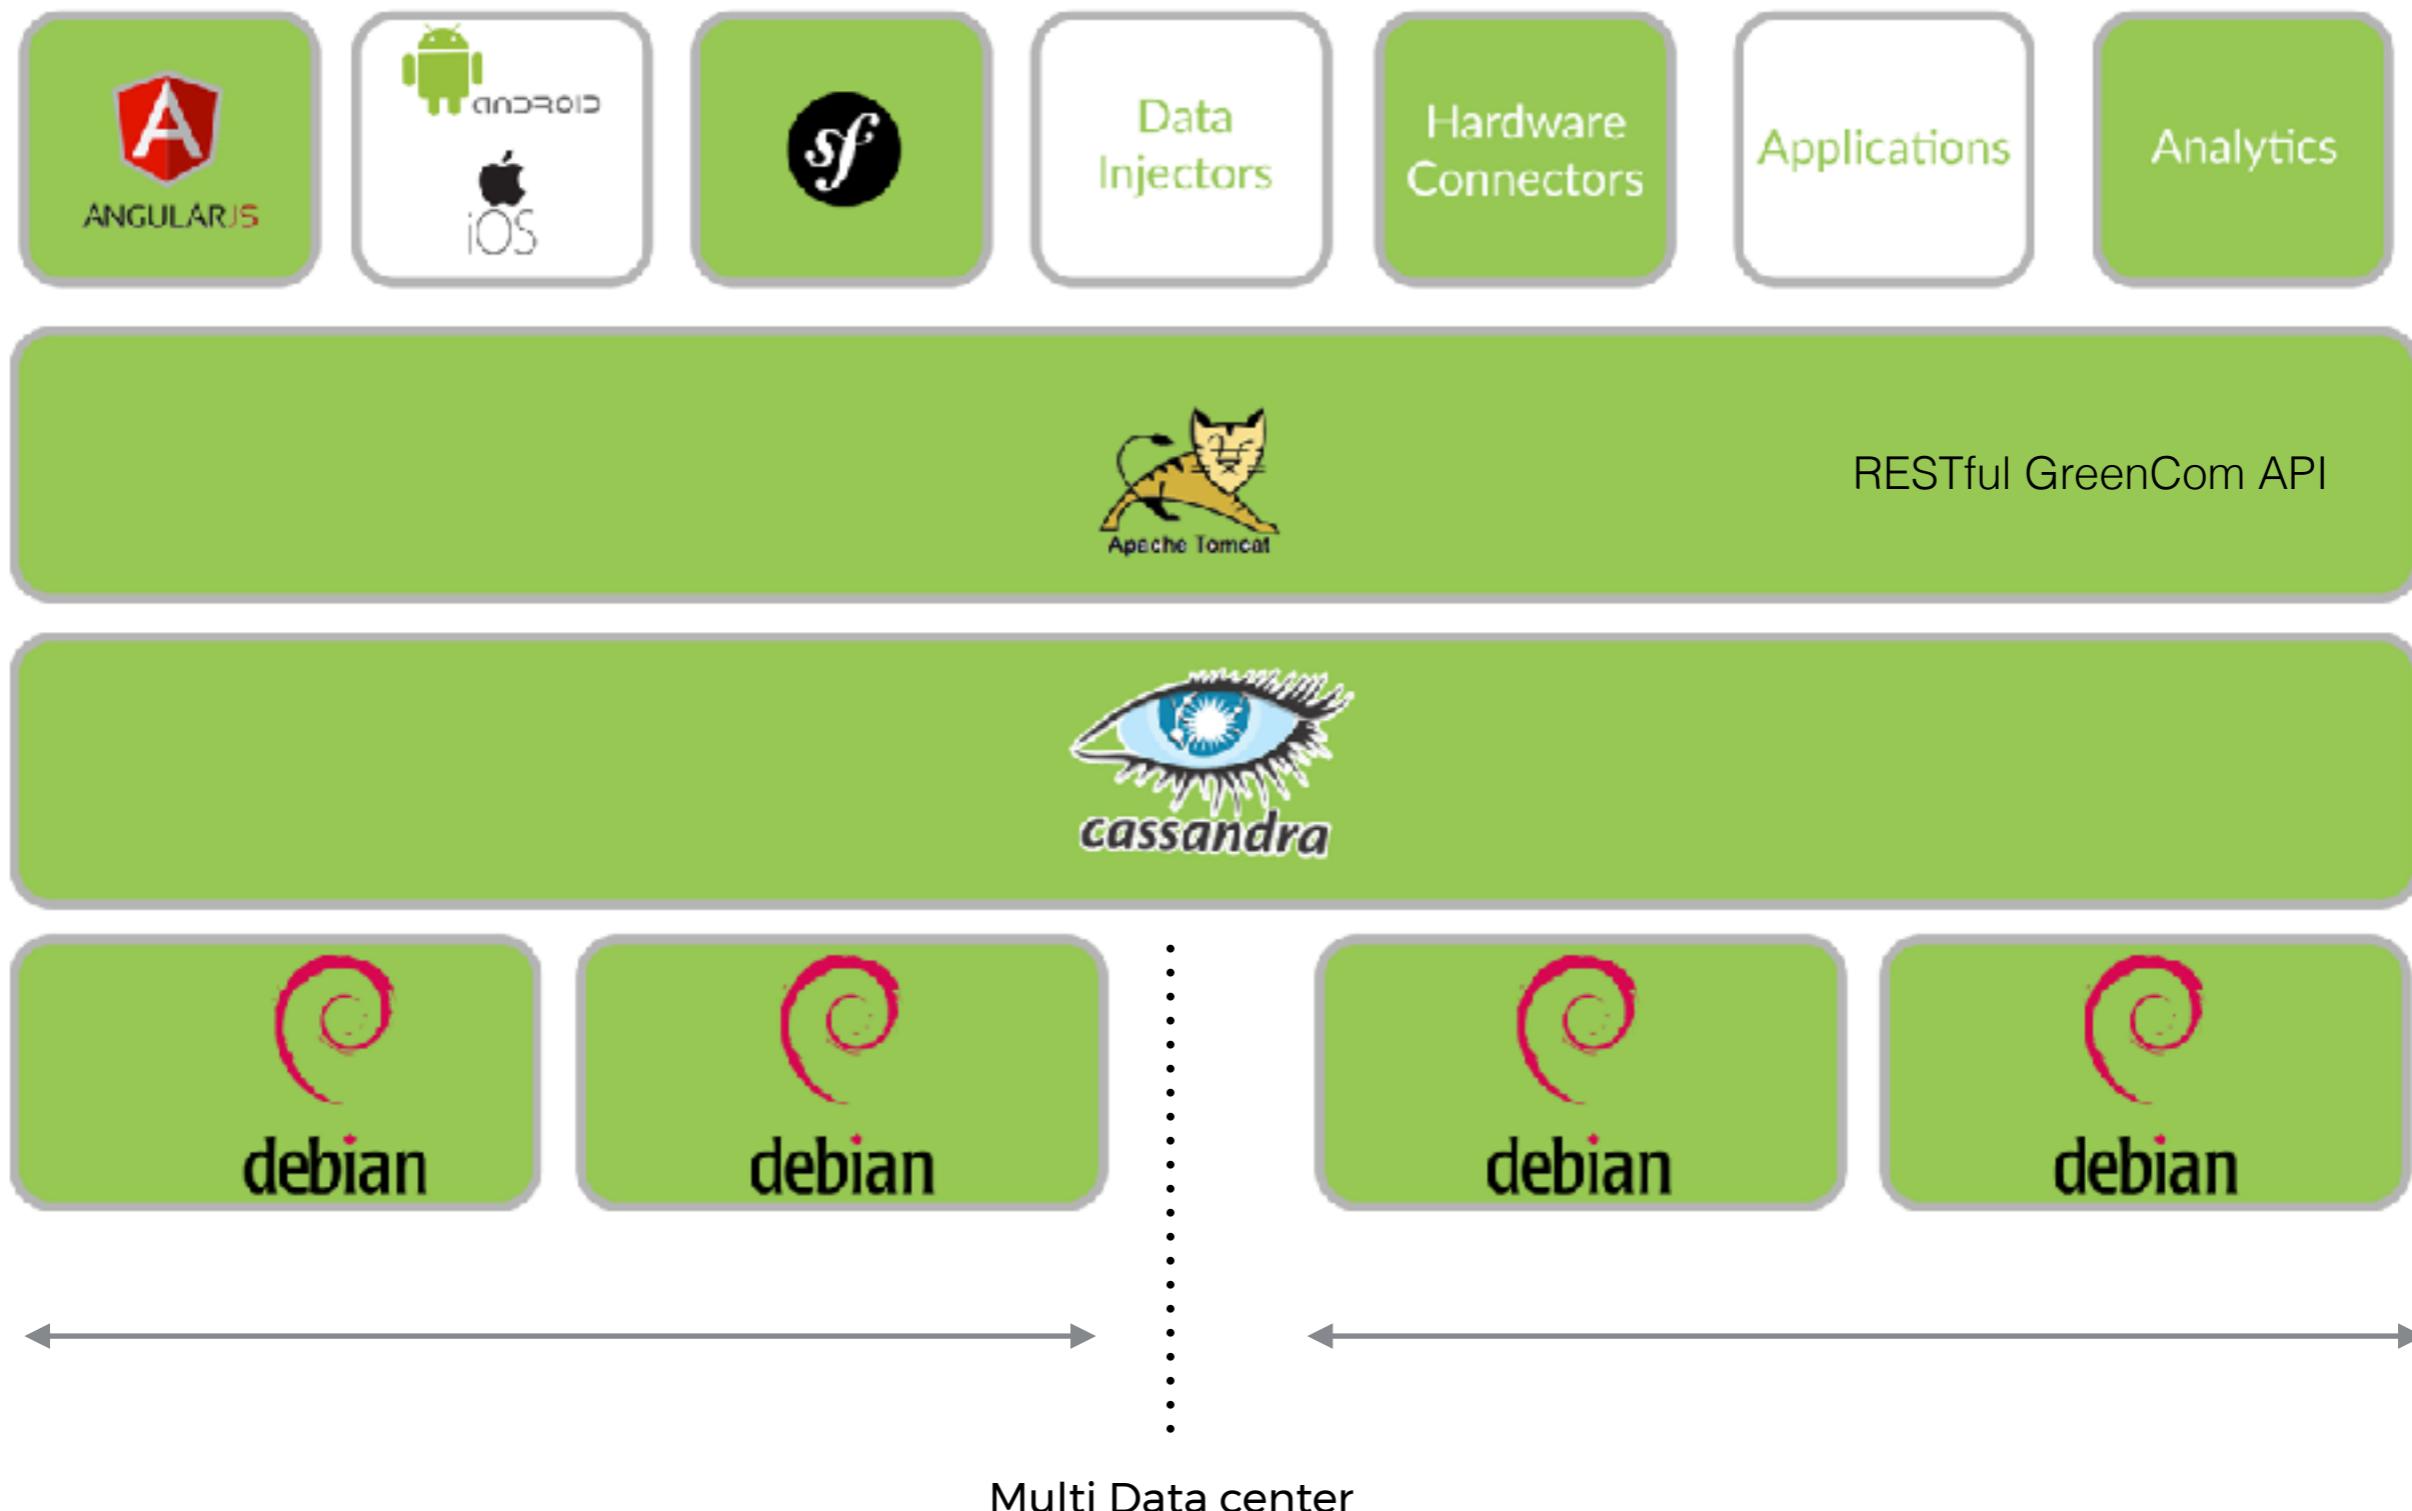


# Cassandra

Apache Cassandra is an **open source, distributed** database management system designed to handle **large amounts of data** across many **commodity servers**, providing **high availability** with **no single point of failure**. Cassandra offers robust support for clusters spanning **multiple datacenters**, with asynchronous masterless replication allowing low latency operations for all clients.

Cassandra also places a high value on **performance**.

*Apache Cassandra was initially developed at Facebook to power their Inbox Search feature.  
It was released as an open source project in 2008.*



- Simple architecture
- Agile
- API Based
- Resilient
- Built to scale
- Easy to maintain
- Easy to deploy

# What do we need ?



Time Series Storage / Management

Metadata Storage / Management

Scalability

Replication / Data center awareness (data safety)

Performance

Security

=> Cassandra

Speed of development

Agility

=> Micro Service Architecture

"Make each program do one thing well. To do a new job, build afresh rather than complicate old programs by adding new features."

"Expect the output of every program to become the input to another, as yet unknown, program."

*Doug McIlroy,  
inventor of Unix pipes and one of the founders of the Unix tradition*

## Micro-services architecture:

- Simple: small code base, easy to create, understand and maintain.
- Fits in the head of one developer
- The services are small - fine-grained as a singular business purpose similar to the Unix philosophy of "Do one thing and do it well"
- The services are elastic, resilient, composable, minimal, and complete.
- Easy to test
- Easy to deploy (each component can be deployed/upgraded separately; can leverage containers)
- Easy to scale : only the required component can be scaled ( != monolithic architecture )
- Use the right language/stack for the right task
- trivial lifecycle management : in a fast moving environment, micro services can be updated/modified/thrown away easily.
- Resilience: if a service fails, impact on the platform is limited

It's not only about:

- splitting
- replacing function call with network calls

But also about:

- Loose coupling (changing a service must not impact the others)
- High cohesion (related behaviour sits together)
- Technology agnostic API
- KISS principle (Keep It Simple Stupid) to ease consumption
- Keep scaling in mind (external persistence)

350+ services

1000+ instances

20 customers

- making heavy use of RESTful interfaces
- carrying JSON representation
- over HTTP / HTTPS

in a wide variety of languages and frameworks

# Heterogeneity (language / frameworks)



# Heterogeneity (language / frameworks)



A true language



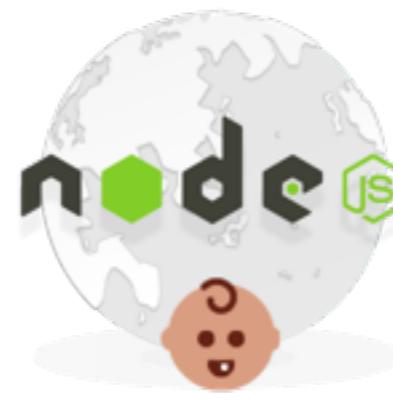
Fading



For old guys



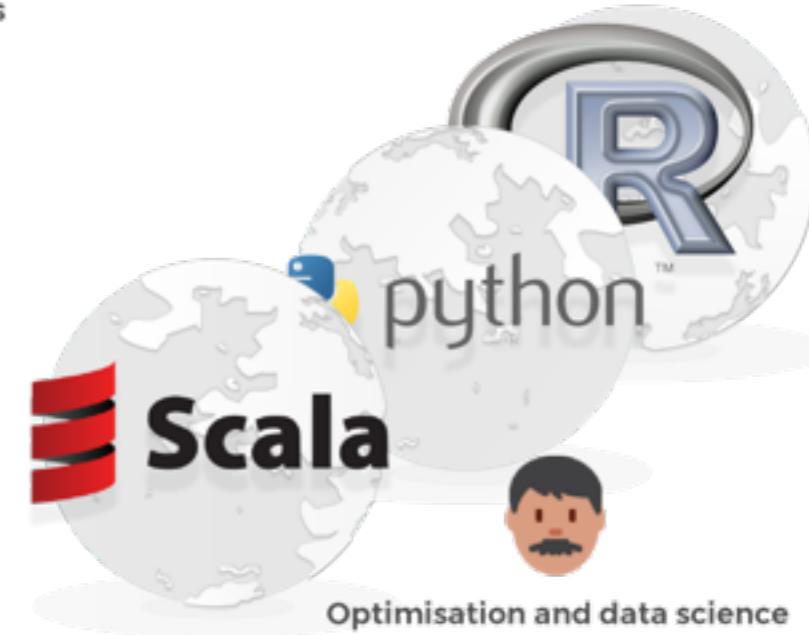
UI



For young guys



Nice logo



Optimisation and data science

# But... Not a silver bullet

- The organisation culture should embrace automation of deployment and testing.
- Adds complexity on the monitoring side
- Fault diagnostic can be complex
- The culture and design principles should embrace failure and faults
- Too early decomposition might jeopardise the entire system

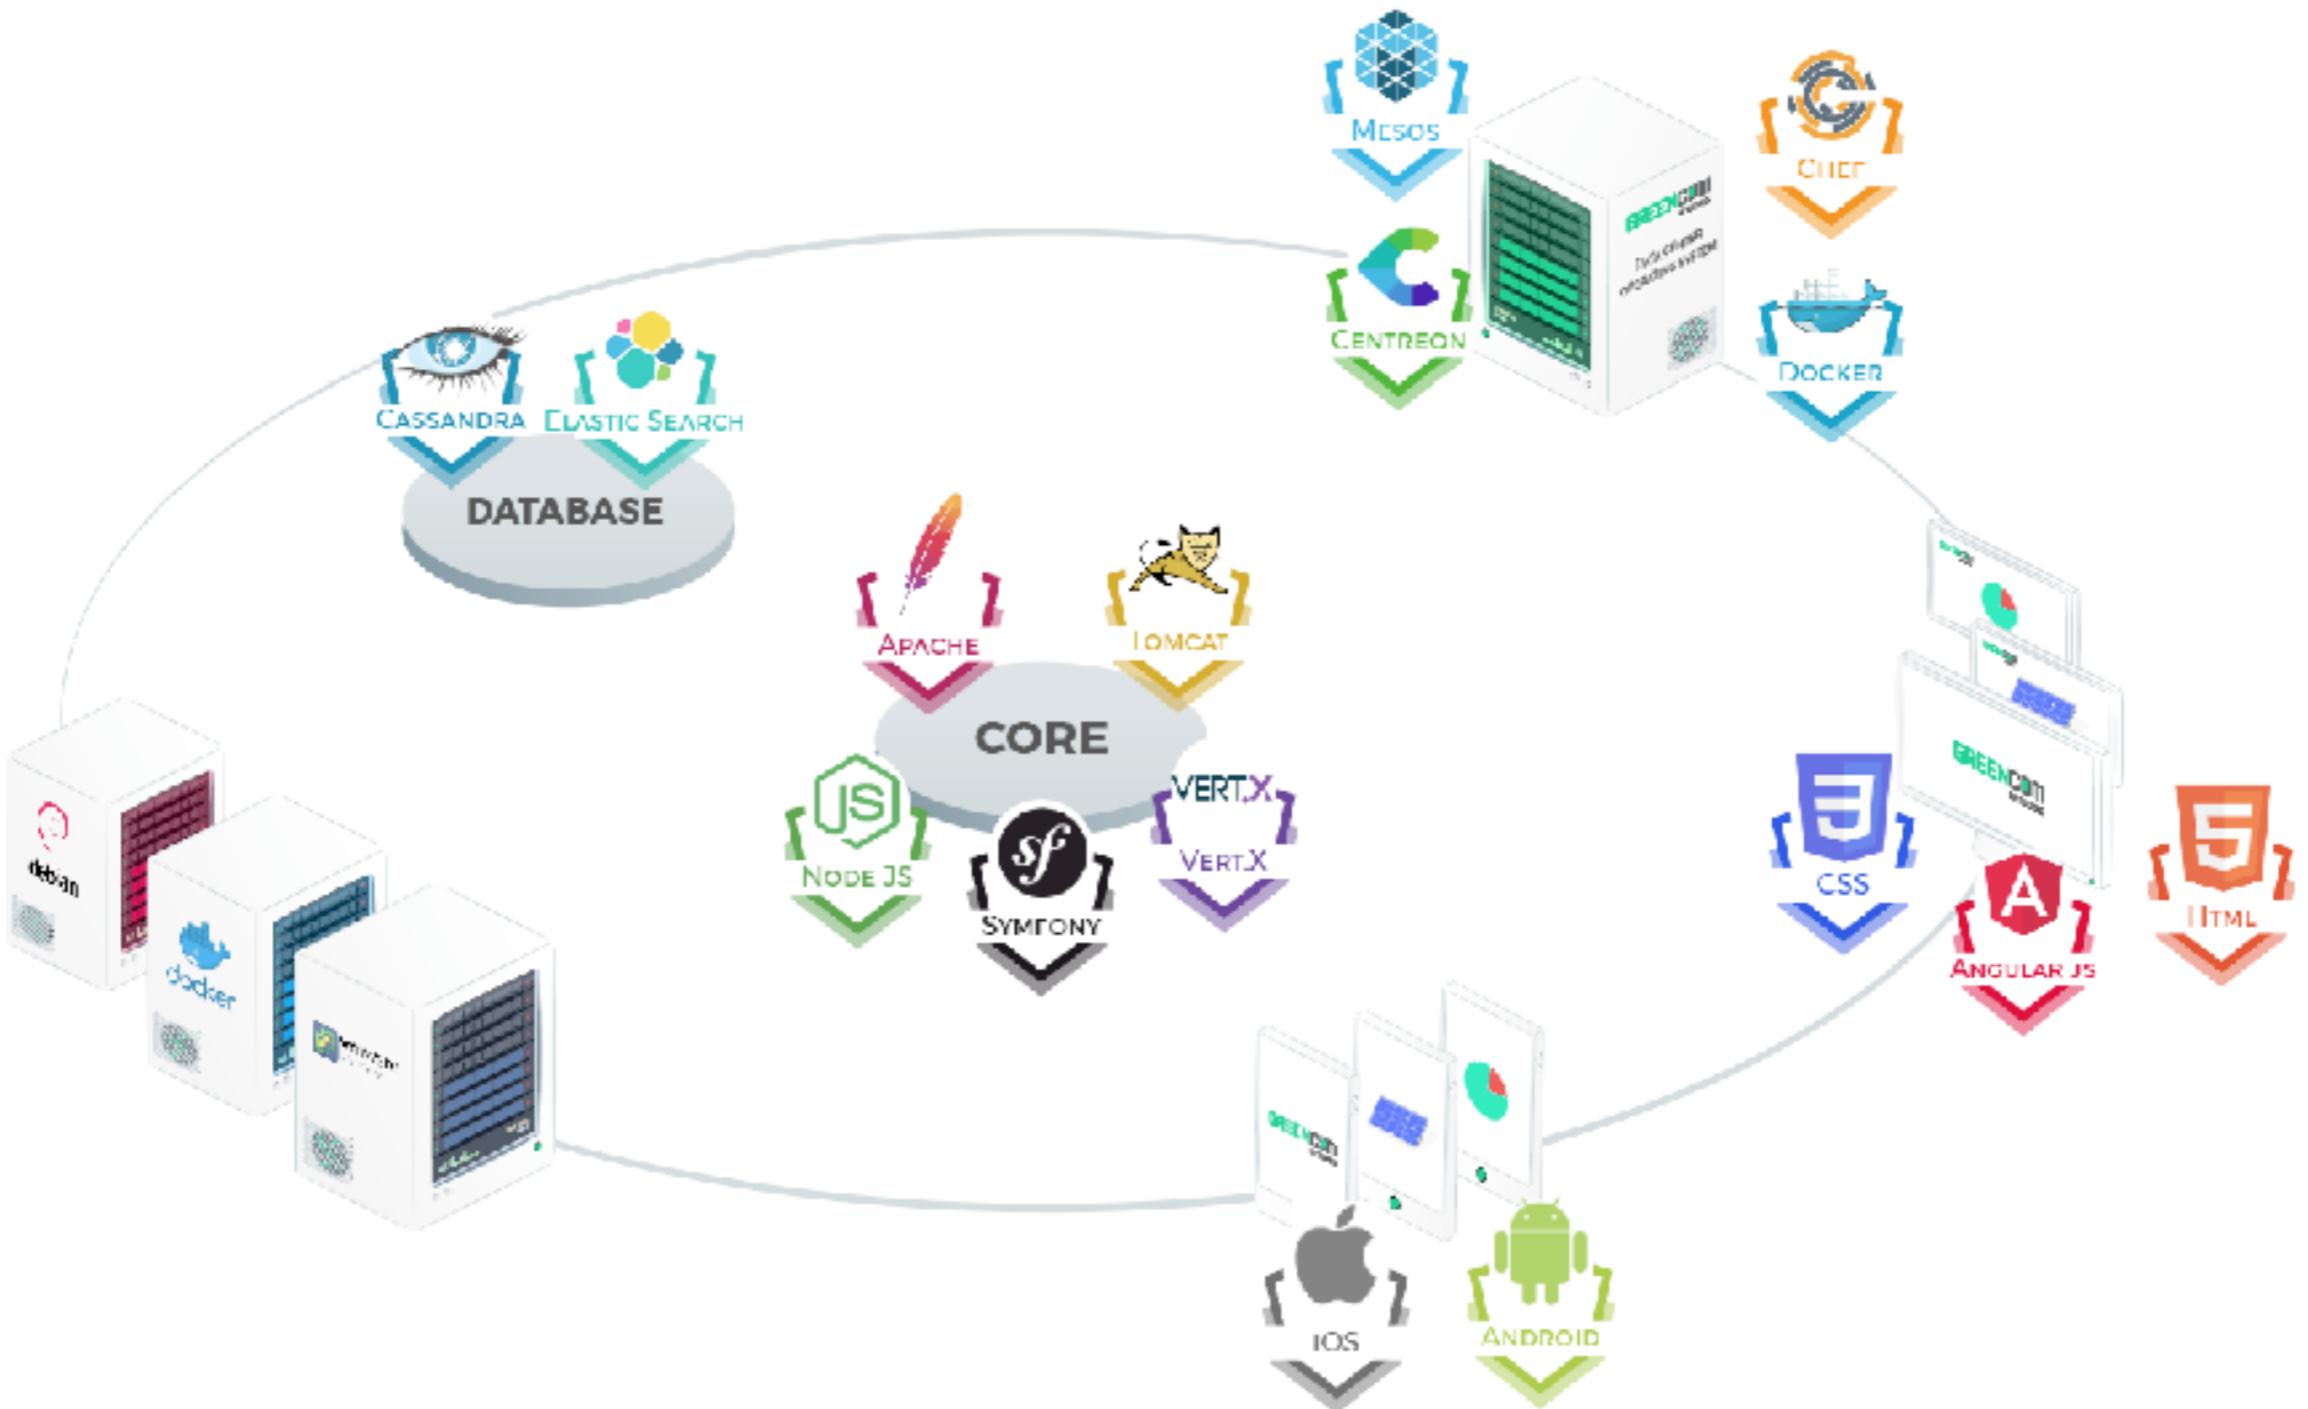


**Jenkins**



**JUnit**

# EIBP Platform V2+



**GREENCOM**  
NETWORKS



**That's all folks !**





Ehh,  
2 questions Doc



What if I want  
to  
react in real time ?

**Everything is event**

- Meter reading
- EV plugged
- Account Created

**Storing it creates data**

**But what if we want to react in real time ?**

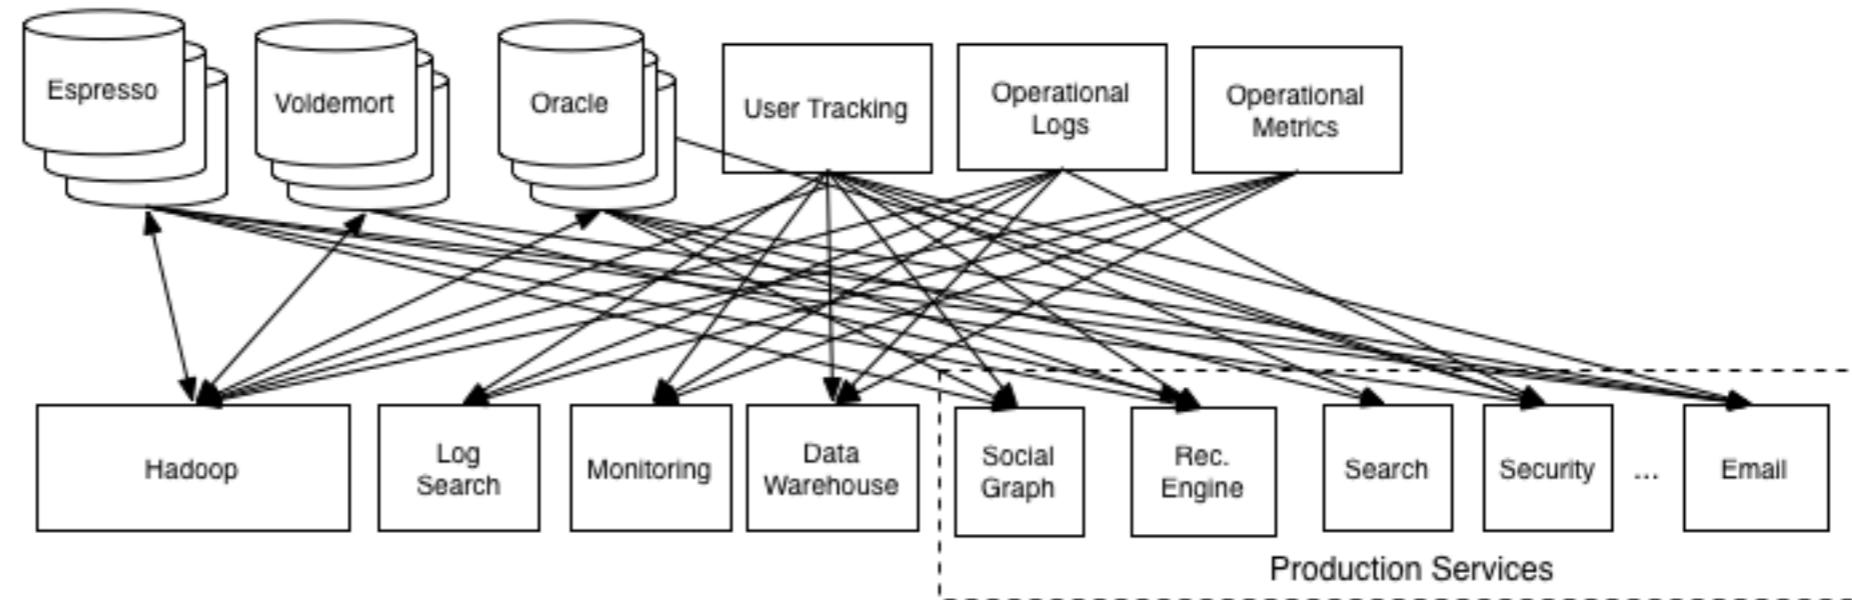
**Event is lost...**

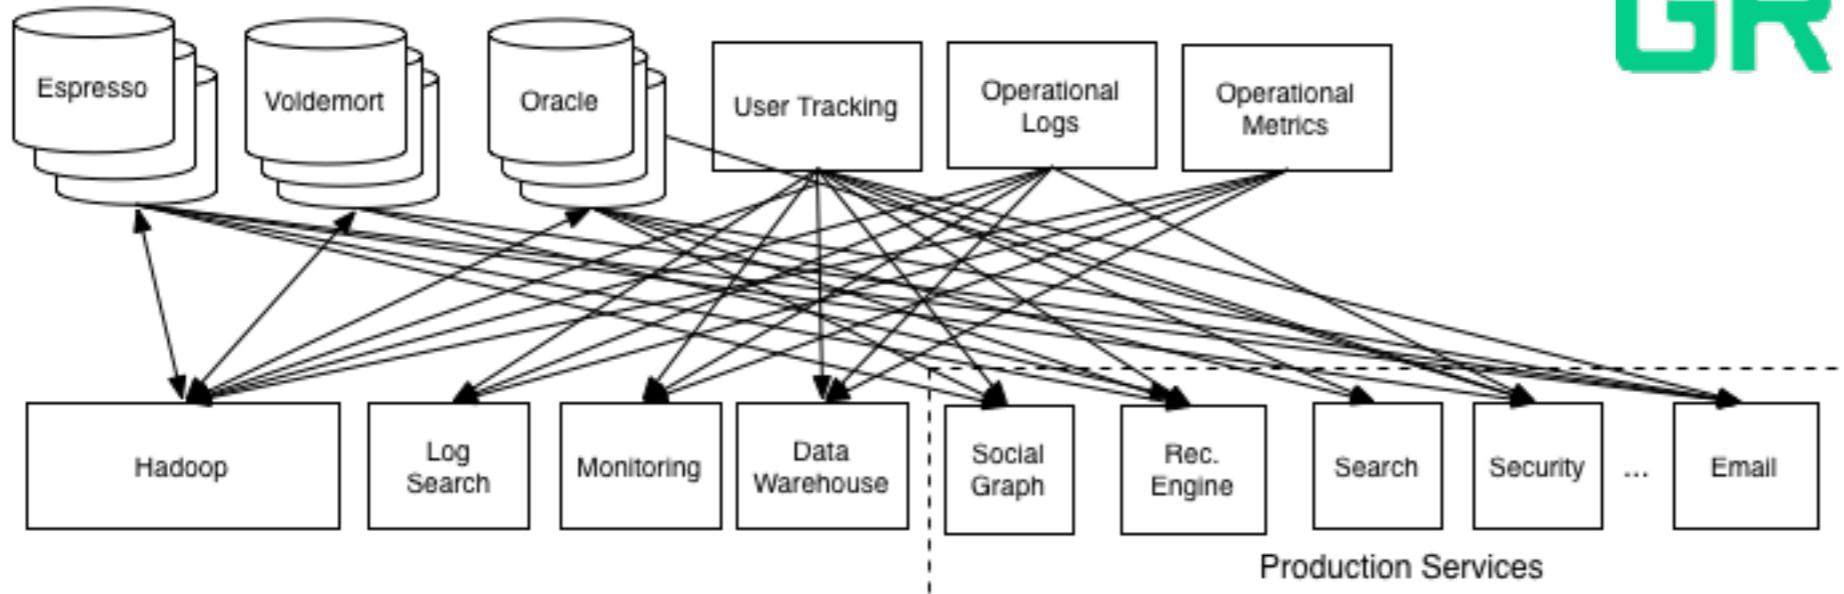
Option 1 : process the event while receiving it,  
hardwiring use cases (not flexible)

Option 2 : recreate the event later (beurk)

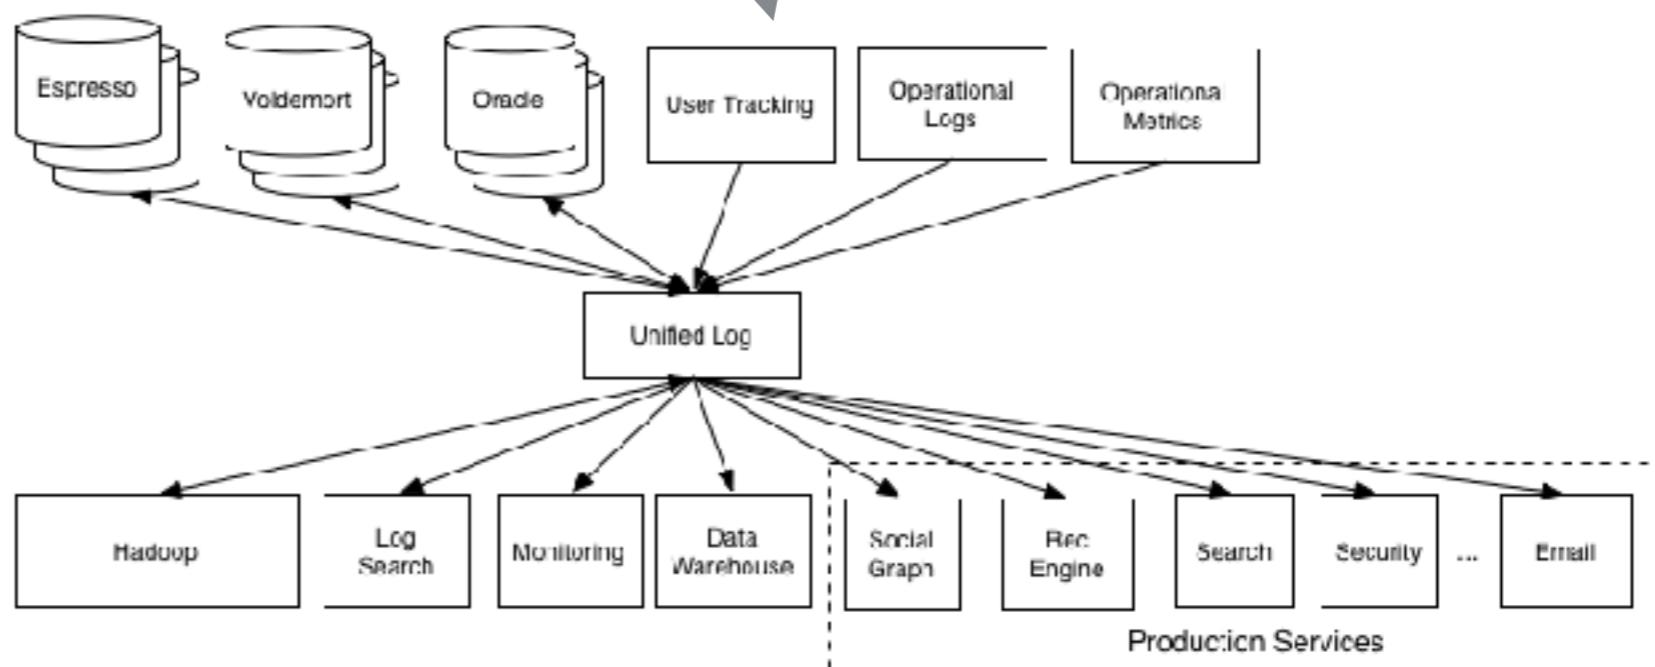
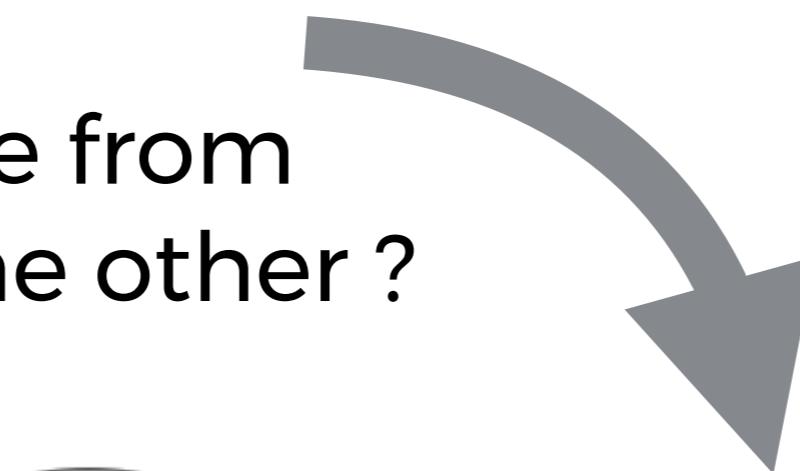


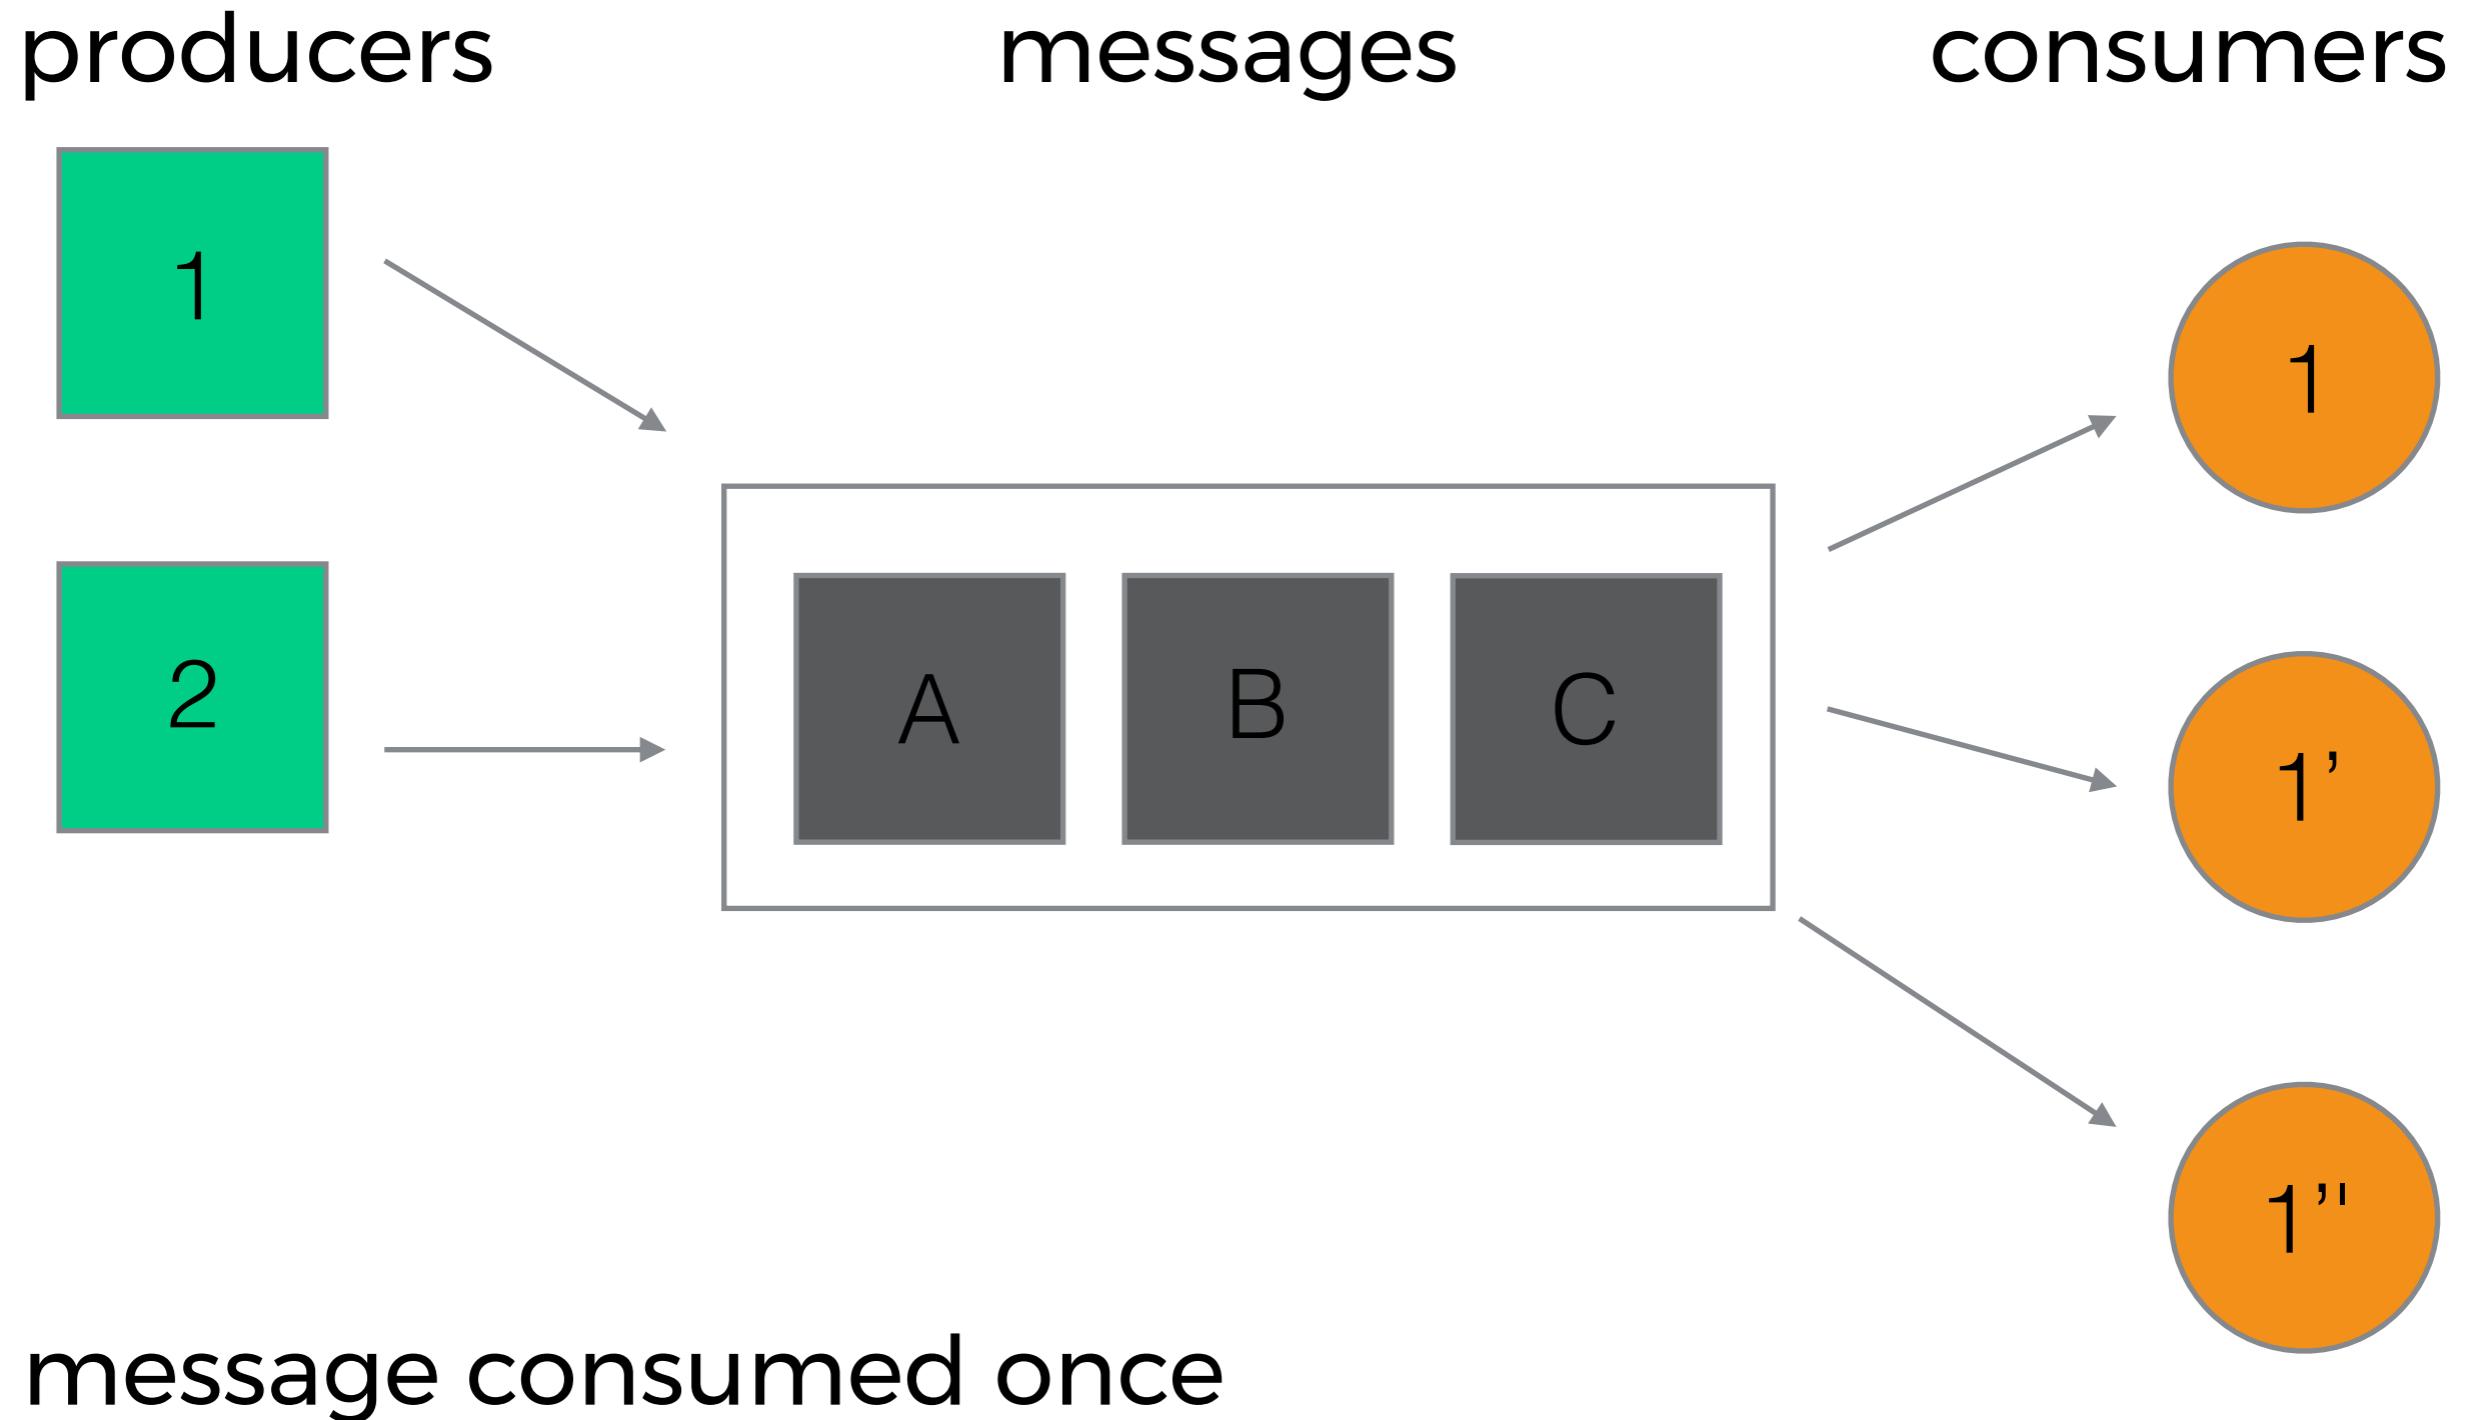
Ok.  
And what will your  
micro service architecture  
look like in few years.  
?





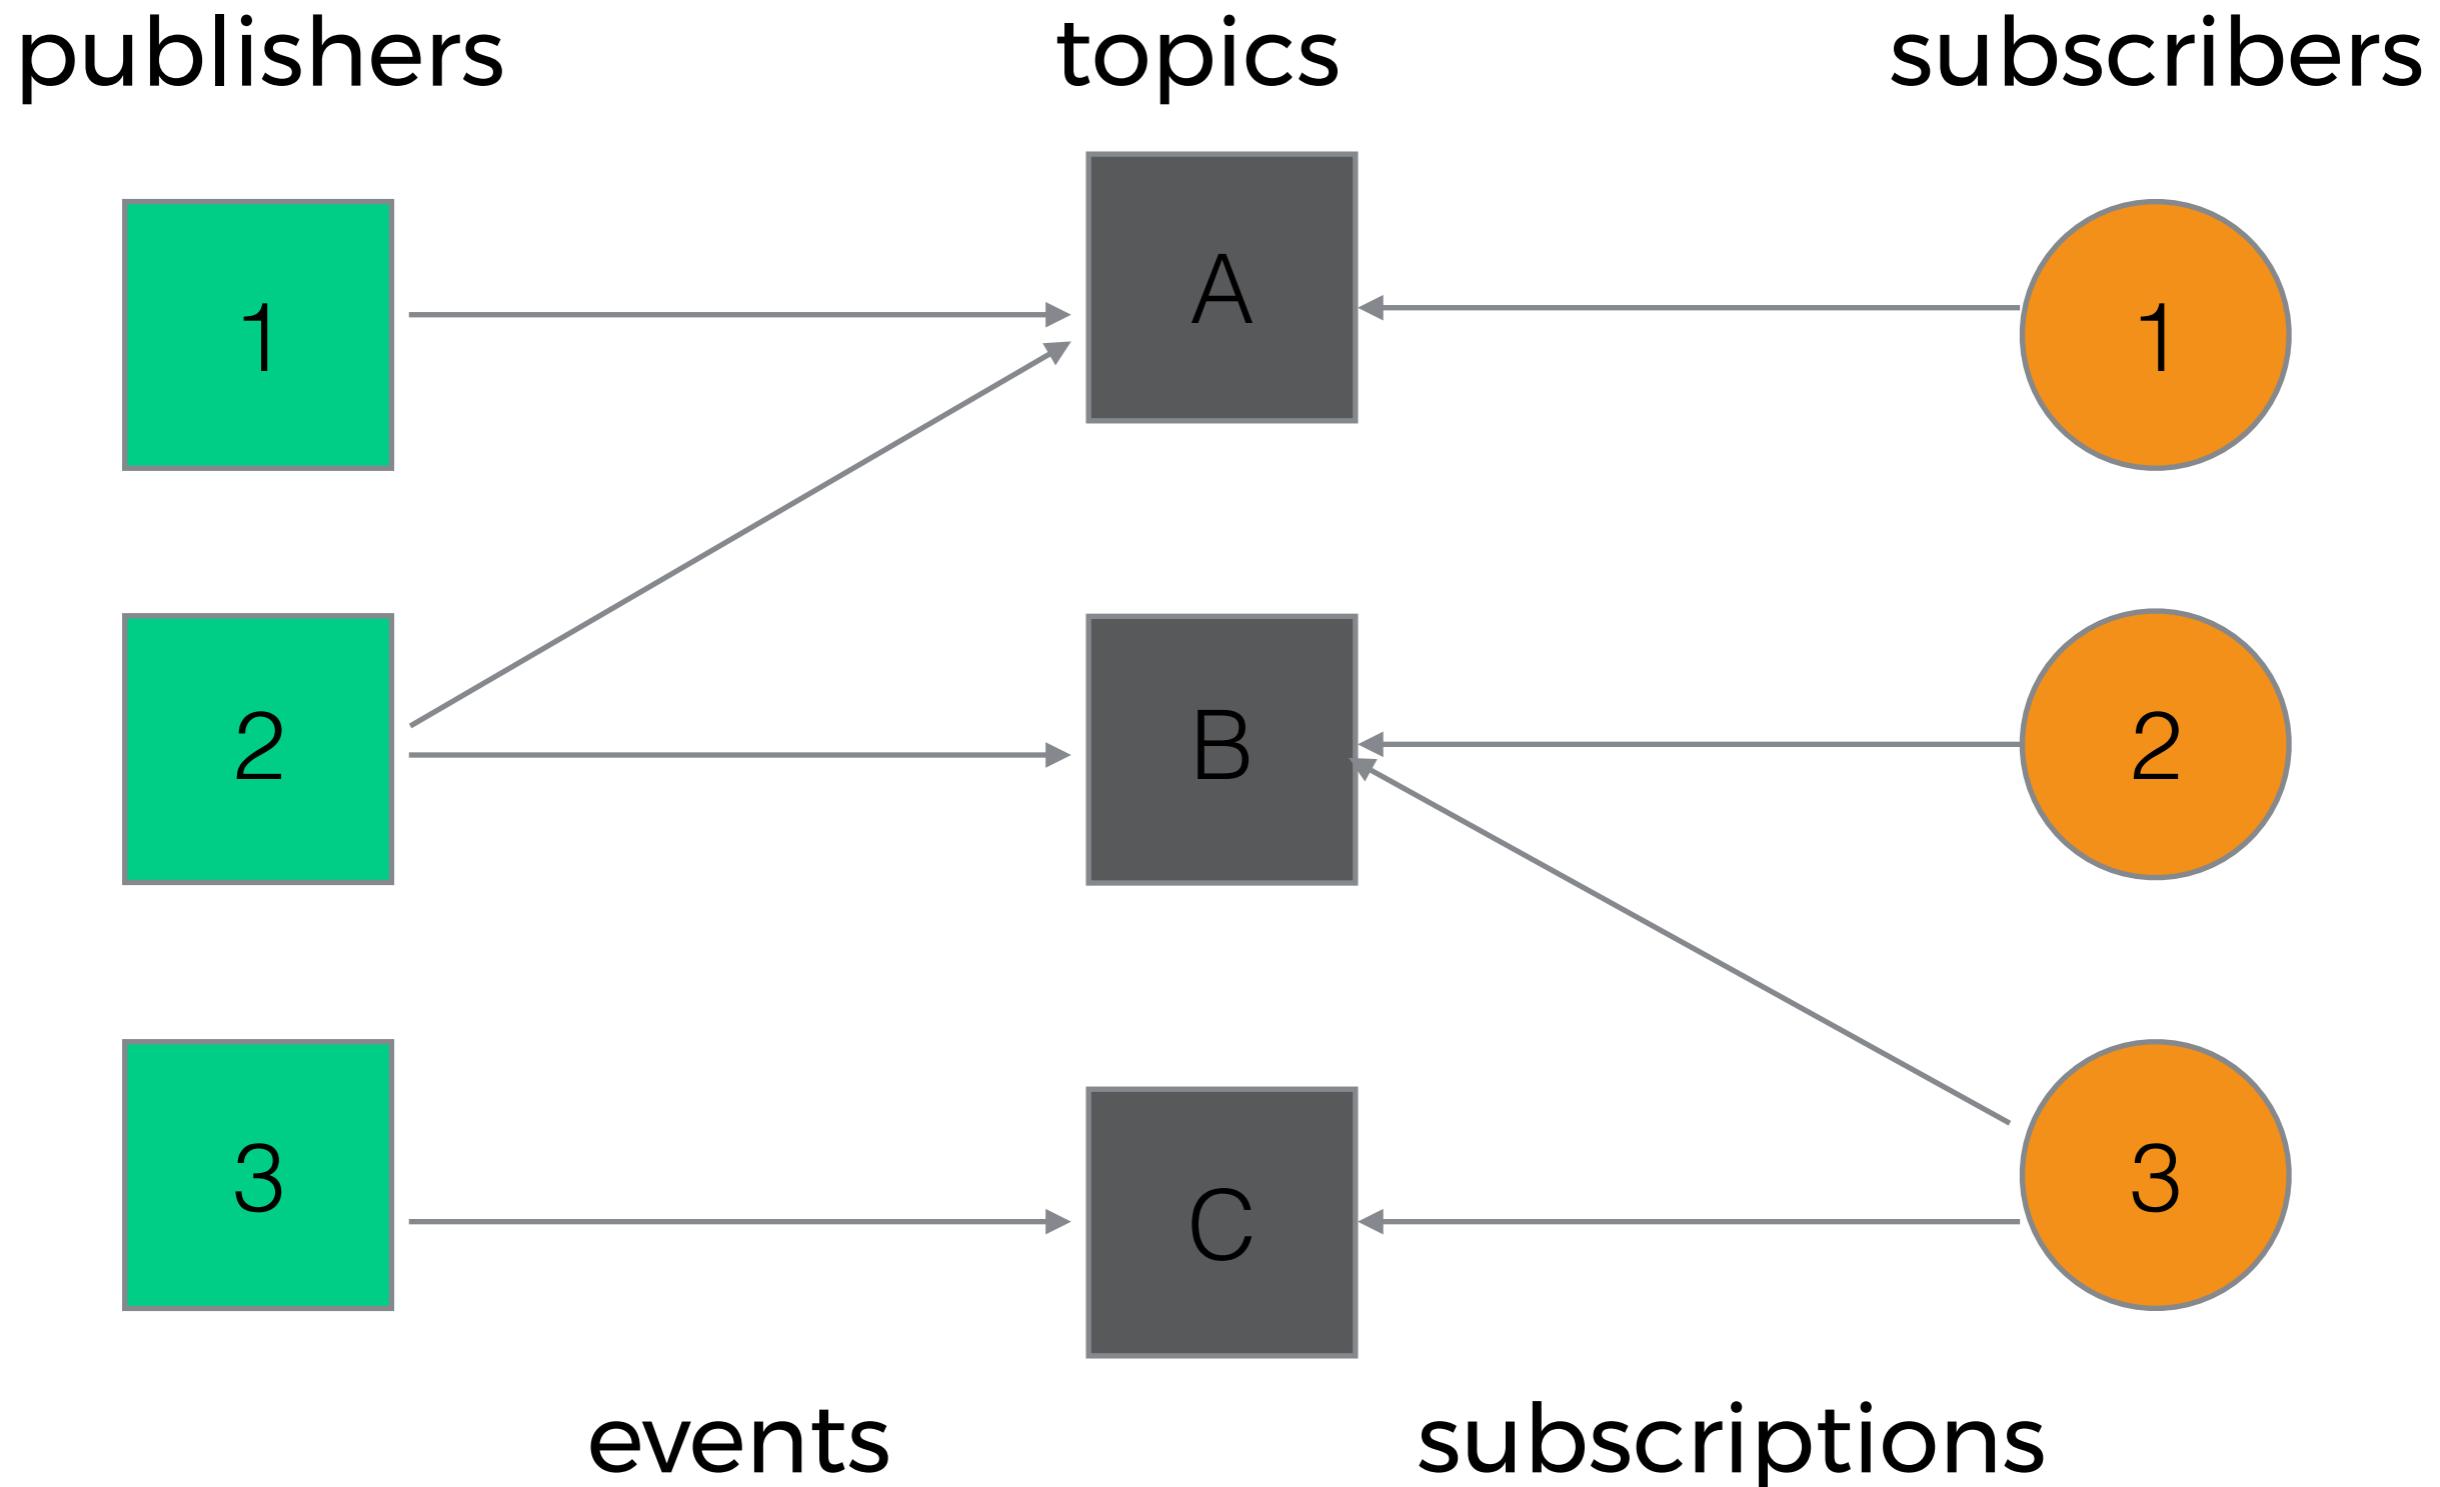
How do we move from  
1 architecture to the other ?





# Publish Subscribe

**GREENCOM**  
NETWORKS



Ok, we need a messaging system then...

## Name dropping

ActiveMQ

RabbitMQ

ZeroMQ\*

HornetQ

...

Kafka

Kerstrel (RIP)

DistributedLog

...

...

SNS

SQS

MessageHub

EventHub

IoTHub

The target solution will have to be:

- Fast (performant)
- Reliable (no SPOF)
- Agile (easy to deal with)
- Scalable (horizontal scalability)
- Cheap (operationally)
- Secure



# Apache Kafka

## A high-throughput distributed messaging system.

Apache Kafka is publish-subscribe messaging rethought as a distributed commit log.

### **Fast**

A single Kafka broker can handle hundreds of megabytes of reads and writes per second from thousands of clients.

### **Scalable**

Kafka is designed to allow a single cluster to serve as the central data backbone for a large organization. It can be elastically and transparently expanded without downtime. Data streams are partitioned and spread over a cluster of machines to allow data streams larger than the capability of any single machine and to allow clusters of co-ordinated consumers

### **Durable**

Messages are persisted on disk and replicated within the cluster to prevent data loss. Each broker can handle terabytes of messages without performance impact.

### **Distributed by Design**

Kafka has a modern cluster-centric design that offers strong durability and fault-tolerance guarantees.

## Scale: Kafka at Microsoft (Ads, Bing, Office)

Kafka Brokers	1000+ across 5 Datacenters
Operating System	Windows Server 2012 R2
Hardware Spec	12 Cores, 32 GB RAM, 4x2 TB HDD (JBOD), 10 GB Network
Incoming Events	1 million per sec, (90 Billion per day, 100 TB per day)
Outgoing Events	5 million per sec, (1 Trillion per day, 500 TB per day)
Kafka Topics/Partitions	50+/5000+
Kafka version	0.8.1.1 (3 way replication)

Scaling November 20, 2015



Neha Narkhede @nehanarkhede · 20 Nov 2015

. @apachekafka powers @Microsoft Bing, Ads and Office. 1 trillion messages/day, 1000+ brokers, 5 million/sec peak



269

213

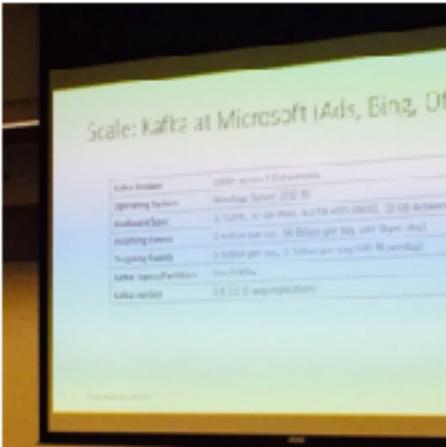
...

# Kafka



**Octave Klabo / Oles** @olesovhcom · 27 Nov 2015

> 5 millions per sec here .. with lot of BigHG 24x600Go SAS each .. :)



**Neha Narkhede** @nehanarkhede

. @apachekafka powers @Microsoft Bing, Ads and Office.  
1 trillion messages/day, 1000+ brokers, 5 million/sec peak



4



7

...

@linkedin:  
400 nodes,  
18k topics,  
220 billions msg/day (pic 3.2 millions msg/s)

May 2014

# Kafka deployments



400 nodes,  
18k topics,  
220 billions msg/day  
peak 3.2 millions msg/s



50 clusters  
4,000 nodes  
3 AWS regions  
1 trillion msg/day



1000 nodes  
5 datacenters  
1 trillion msg/day  
peak 5 millions msg/s

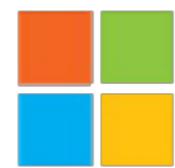
# Kafka deployments



400 nodes,  
18k topics,  
220 billions msg/day  
peak 3.2 millions msg/s



50 clusters  
4,000 nodes  
3 AWS regions  
1 trillion msg/day



1000 nodes  
5 datacenters  
1 trillion msg/day  
peak 5 millions msg/s



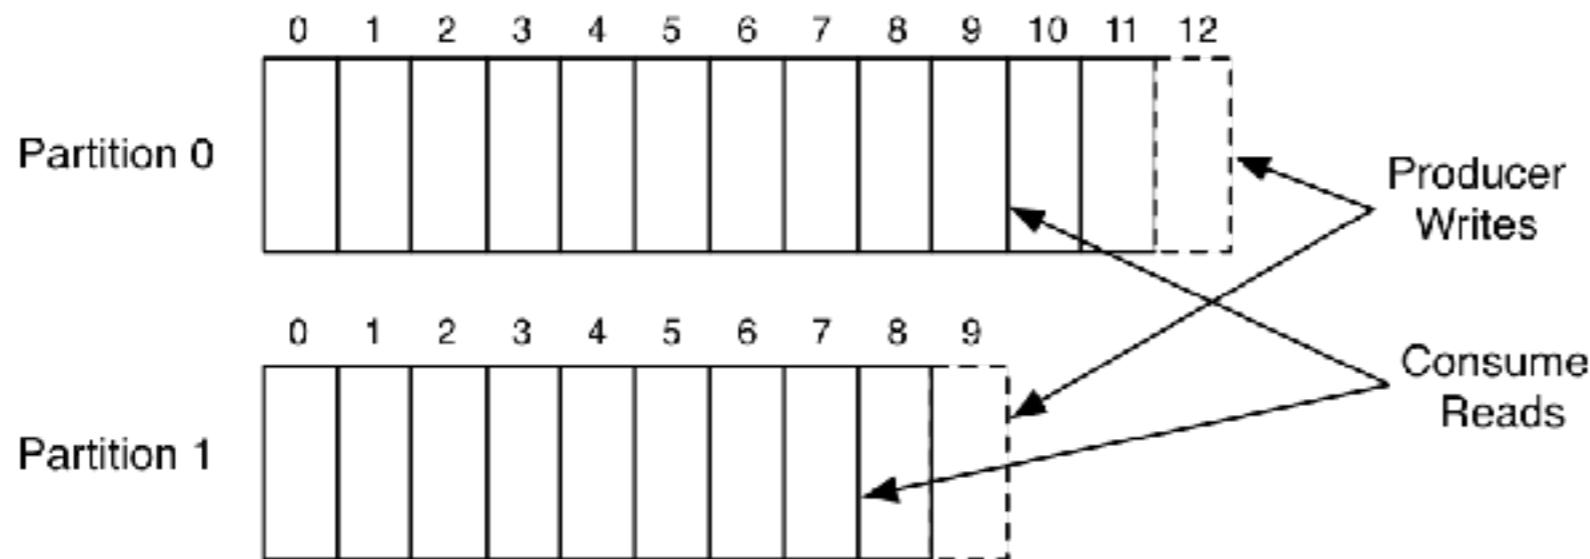
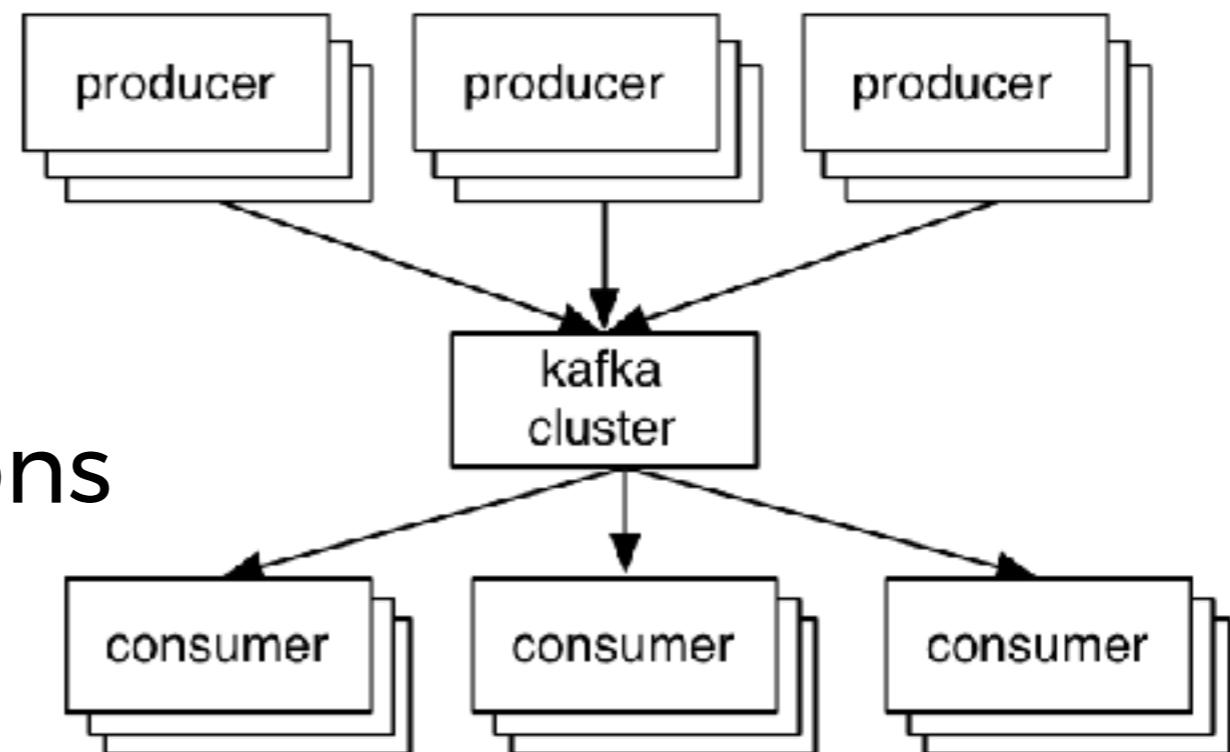
4 clusters  
8 nodes in production,  
1.8k topics,  
500 millions msg/day  
peak 5000 msg/s



# Kafka concepts

producers produce  
consumers consume :-o

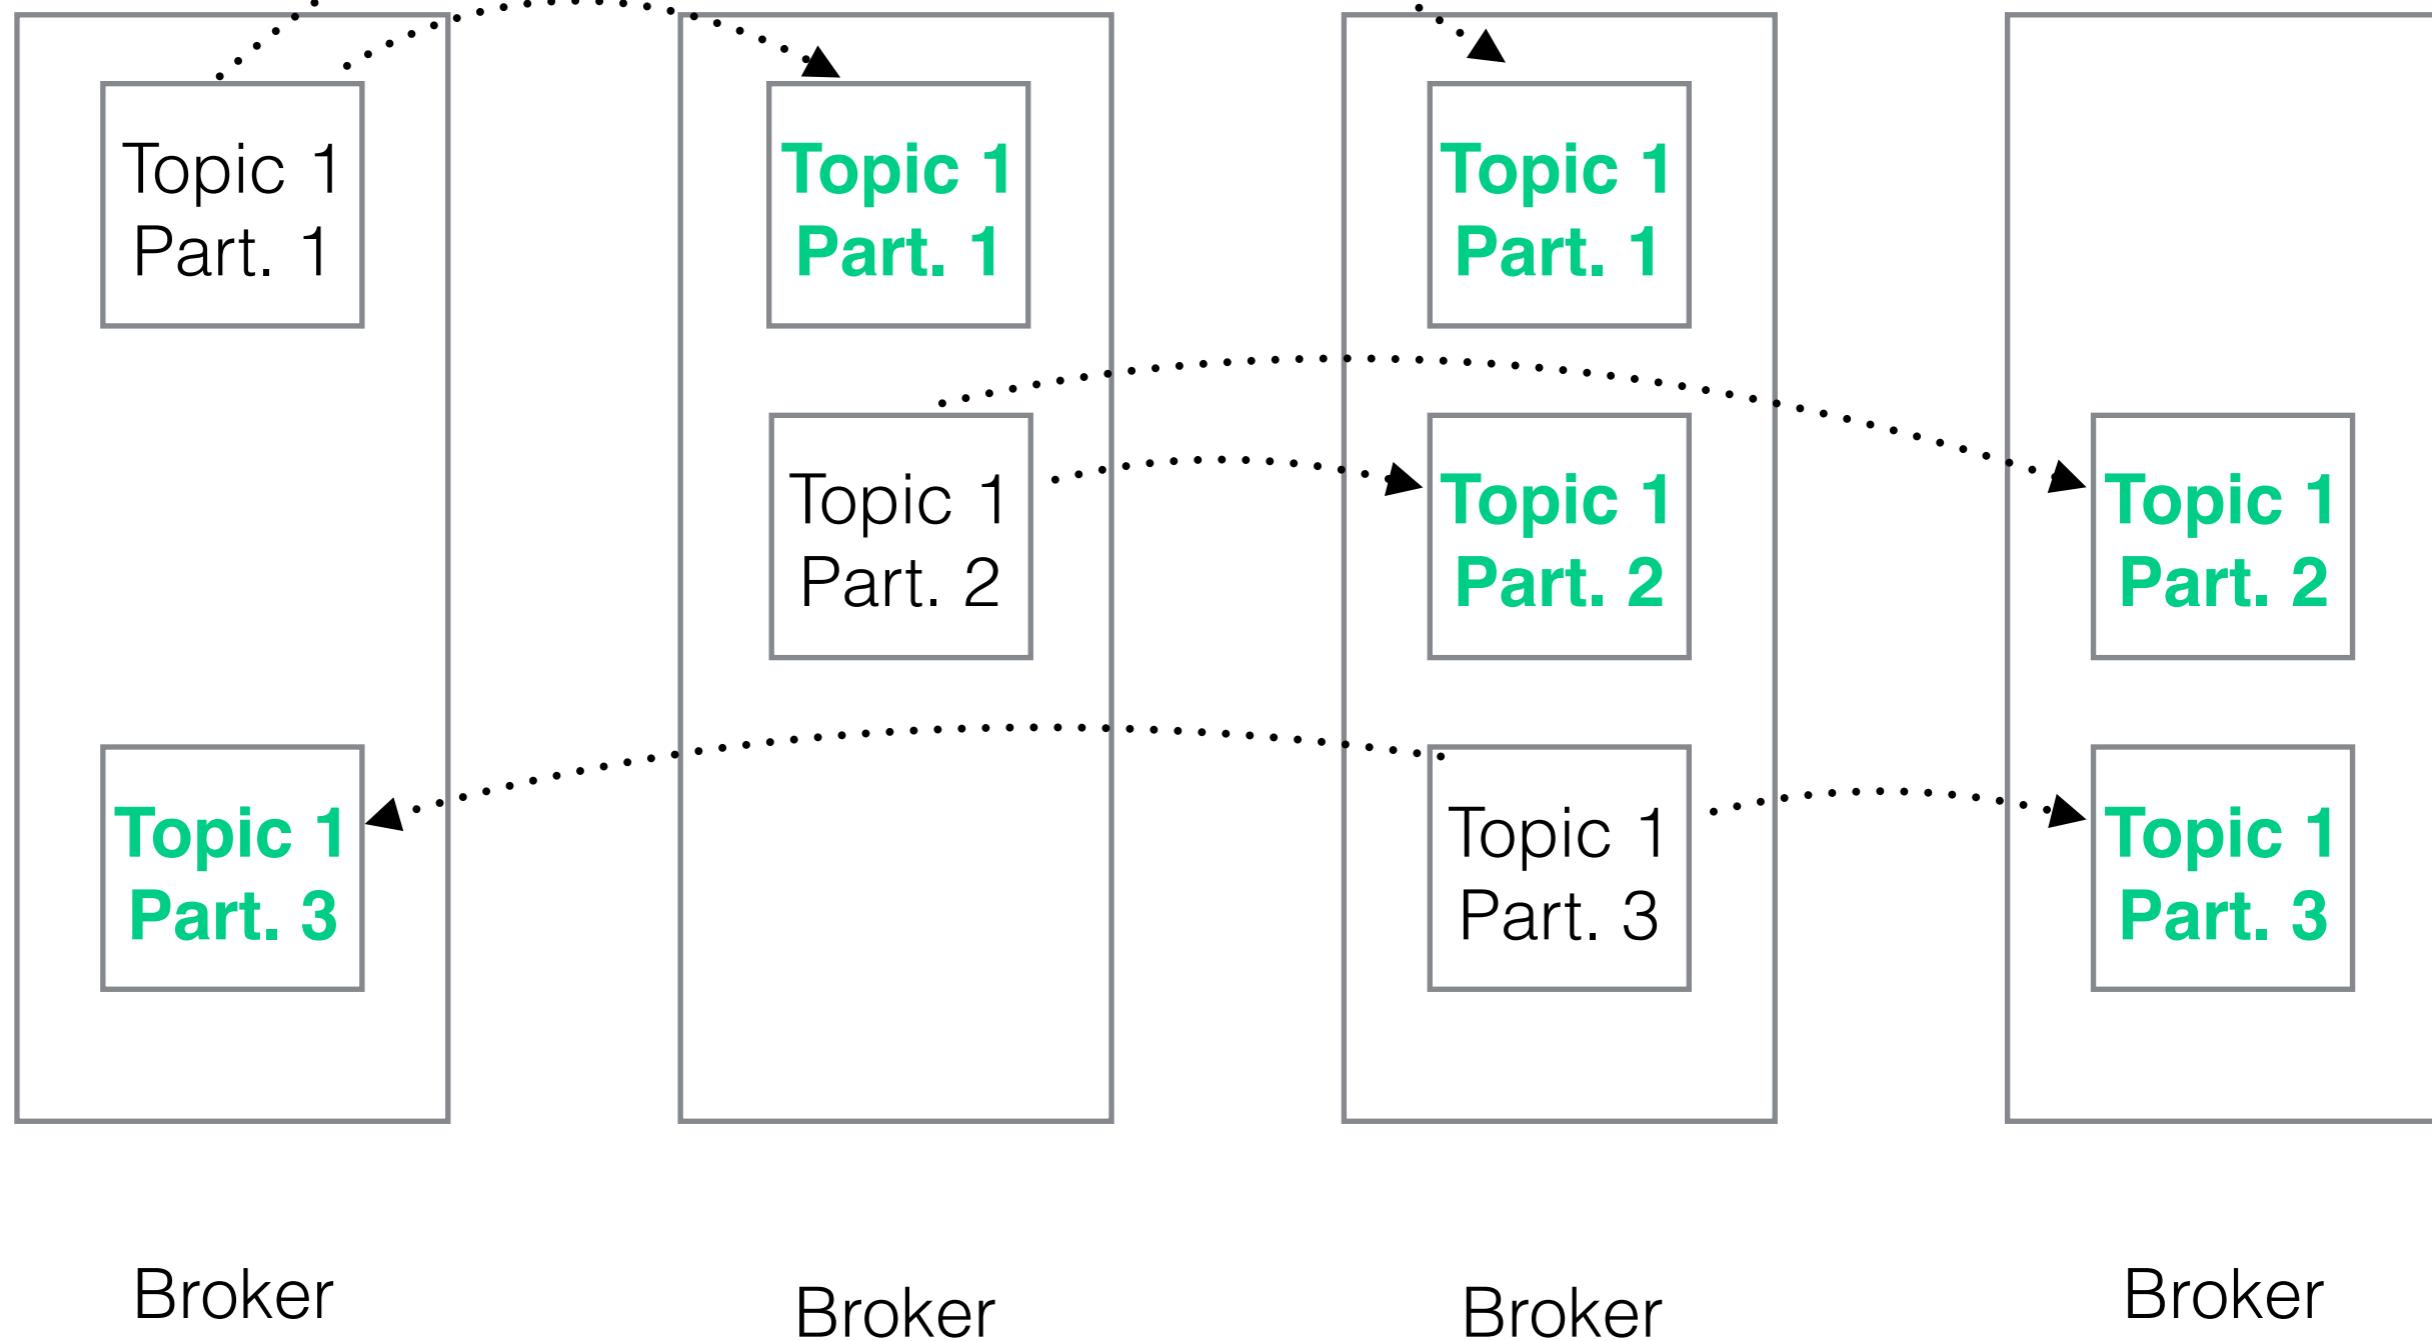
- data are stored in topics
- topics are split in partitions
- partitions are replicated



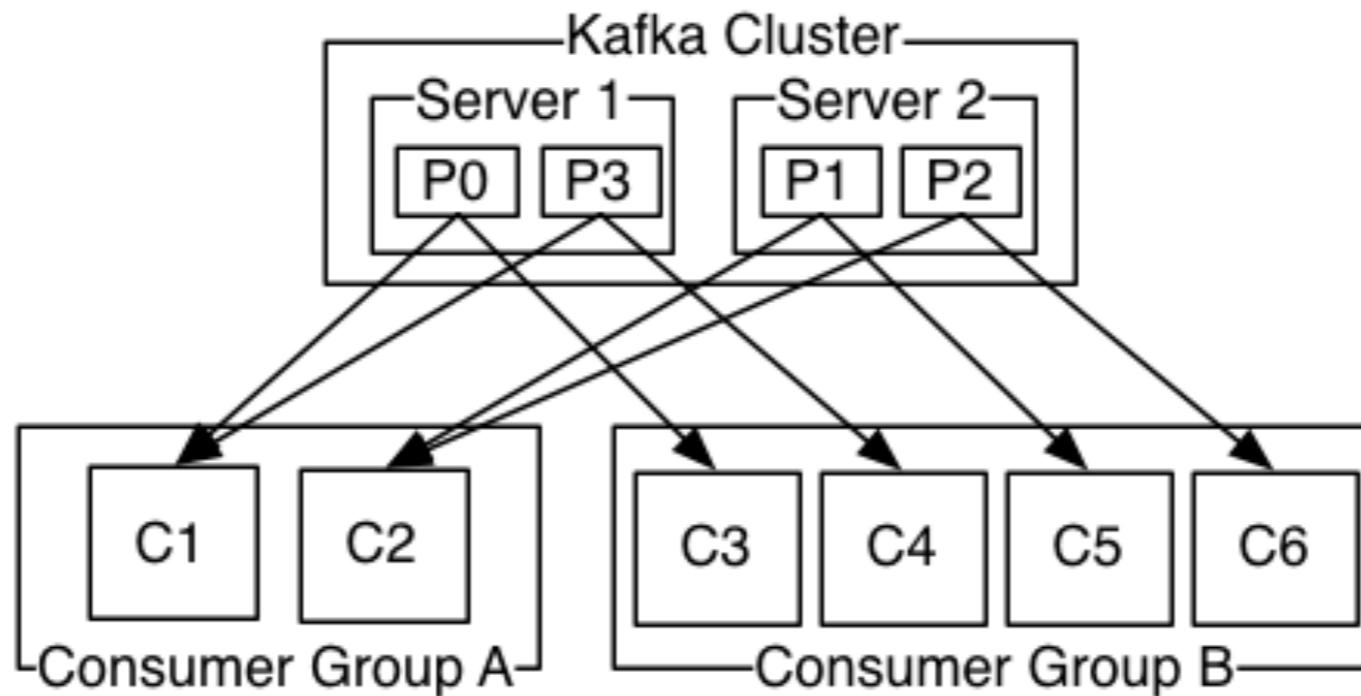
# Kafka concepts

Leader

**Follower**



# Kafka concepts



Number of partitions for a topic fixes the maximum number of consumers for a consumer group.

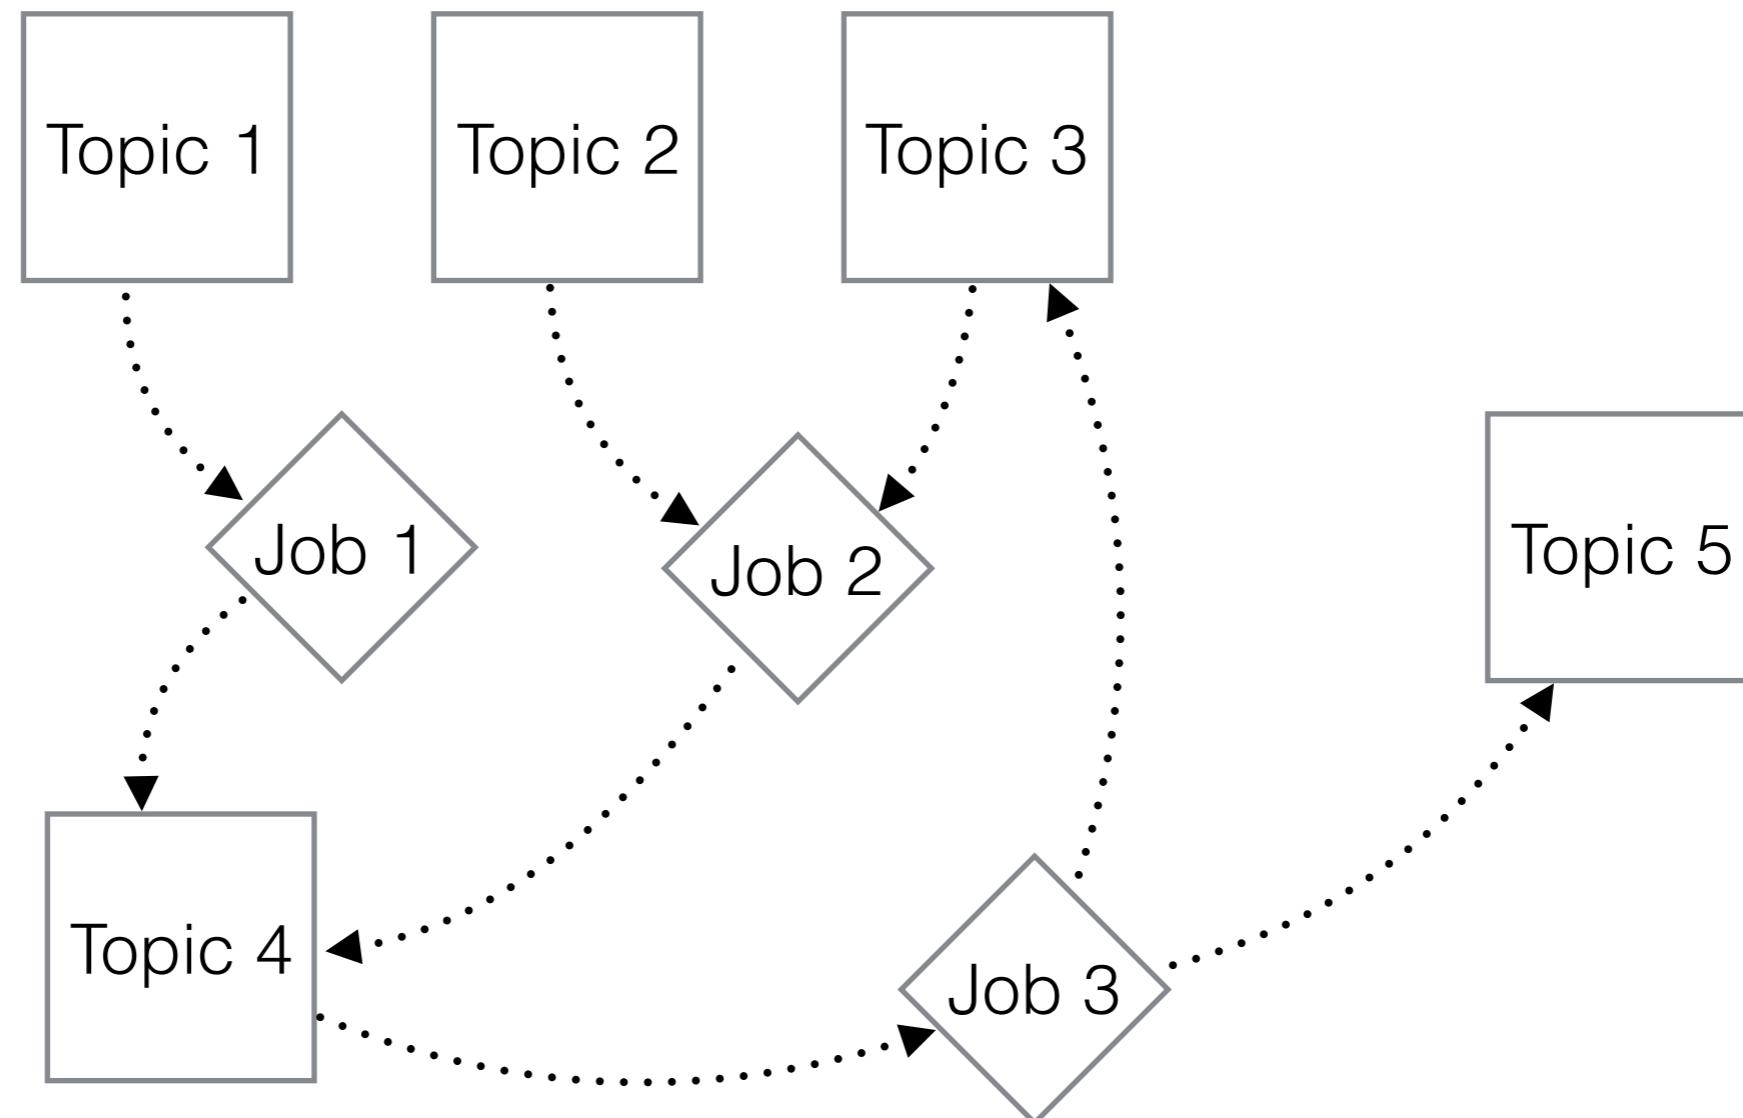


Seb, don't forget to say few words regarding:

- how to scale (cluster & app)
- ordering
- keying
- retention

And now...

Now that we have a way of producing/consuming messages we can create complex event processing pipelines



Producing to Kafka or consuming from Kafka is easy.  
Client library exist for most of the languages

```
> bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test
This is a message
This is another message
```

```
> bin/kafka-console-consumer.sh --zookeeper localhost:2181 --topic test --from-beginning
This is a message
This is another message
```

But for some high throughput complex task. we might want to rely on existing framework / applications

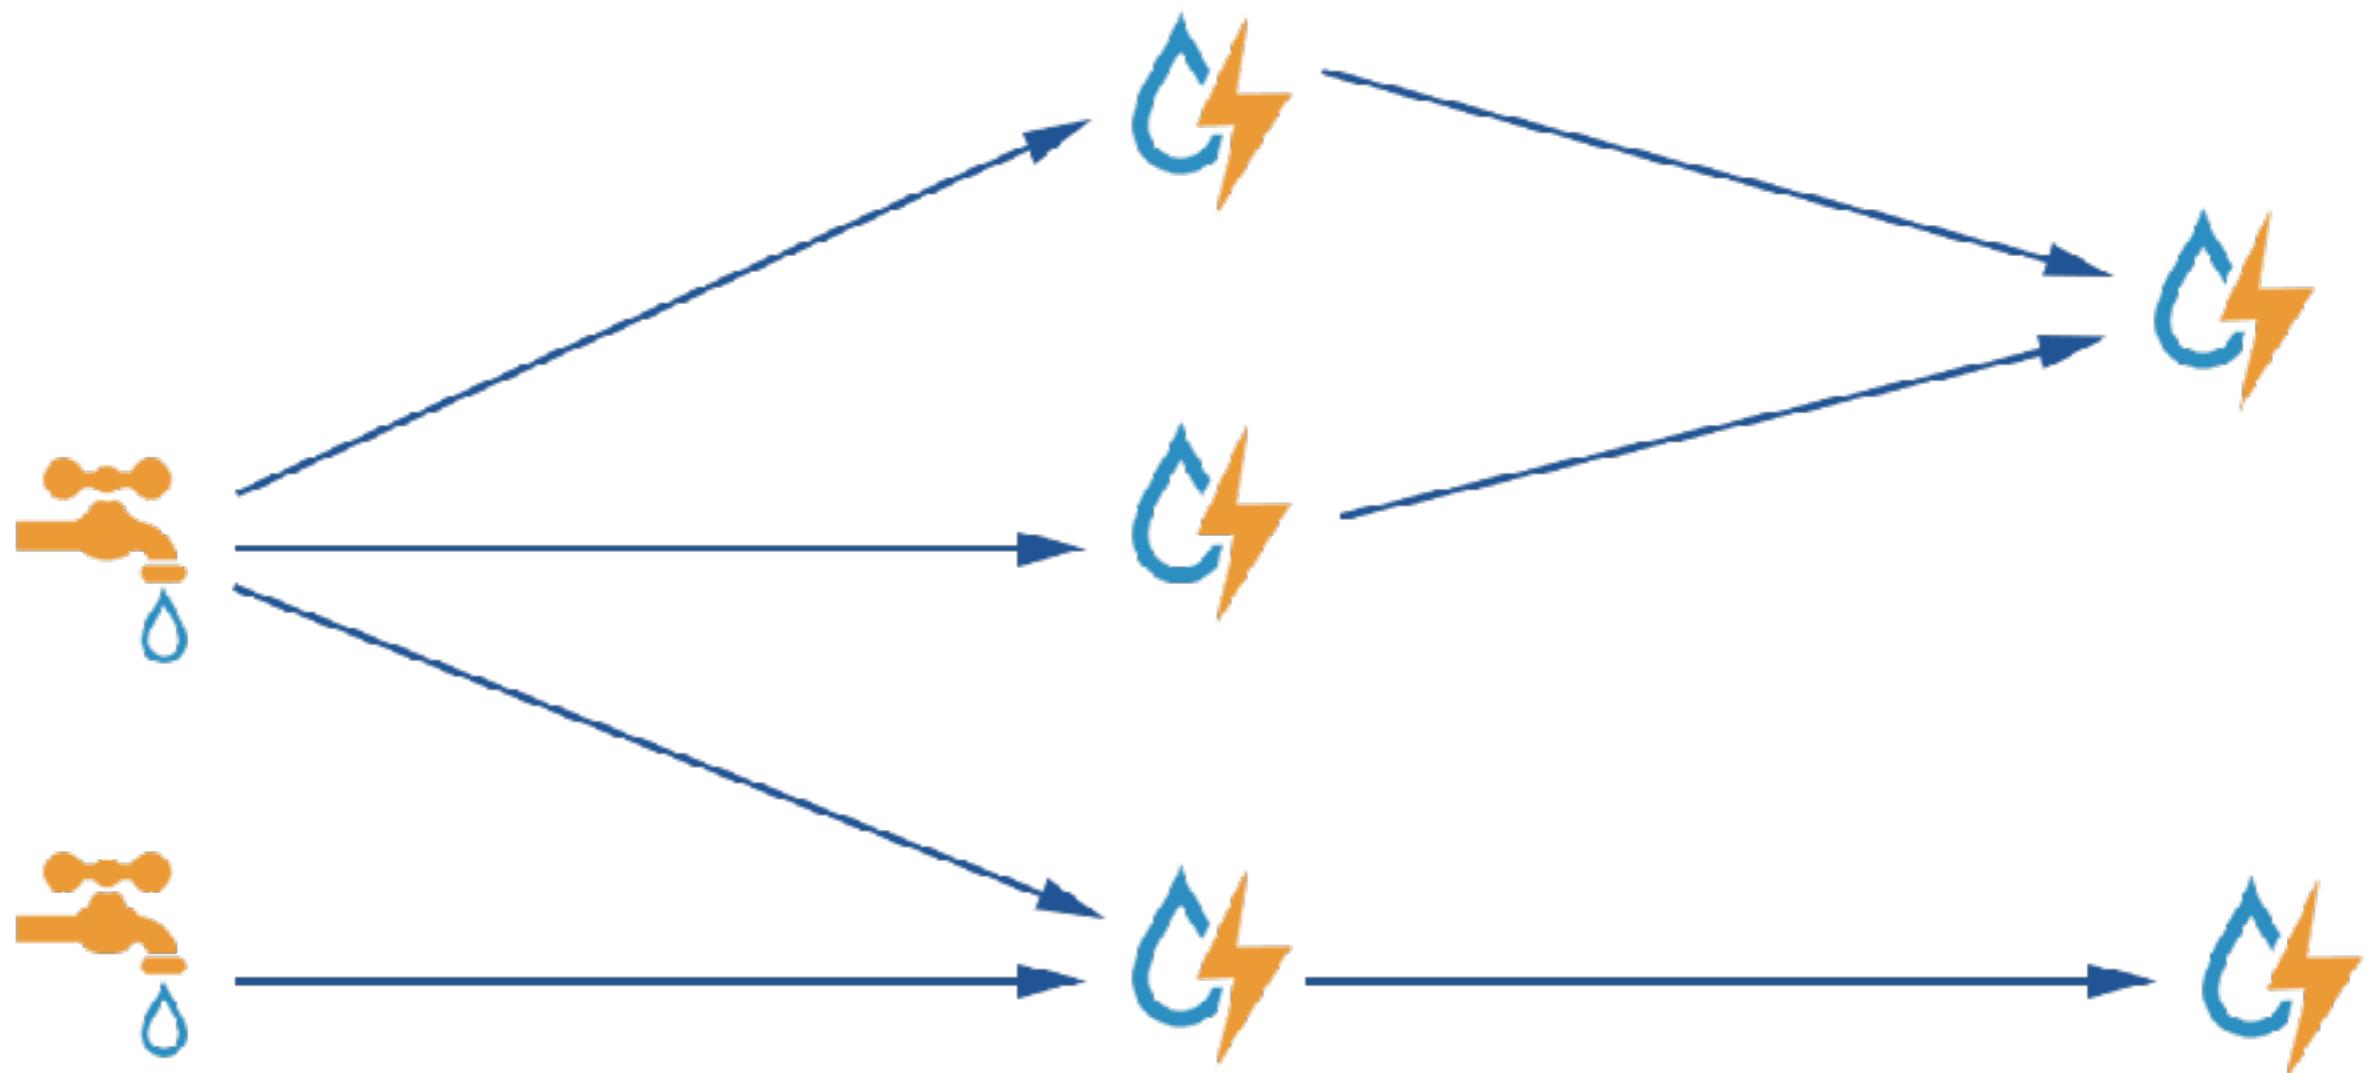


Apache Storm is a free and **open source distributed realtime computation system**. Storm makes it easy to reliably process unbounded streams of data, doing for realtime processing what Hadoop did for batch processing. Storm is simple, can be used with any programming language, **and is a lot of fun to use!**

Storm has many use cases: realtime analytics, online machine learning, continuous computation, distributed RPC, ETL, and more. Storm is **fast**: a benchmark clocked it at **over a million tuples processed per second per node**. It is **scalable, fault-tolerant**, guarantees your data will be processed, and is easy to set up and operate.

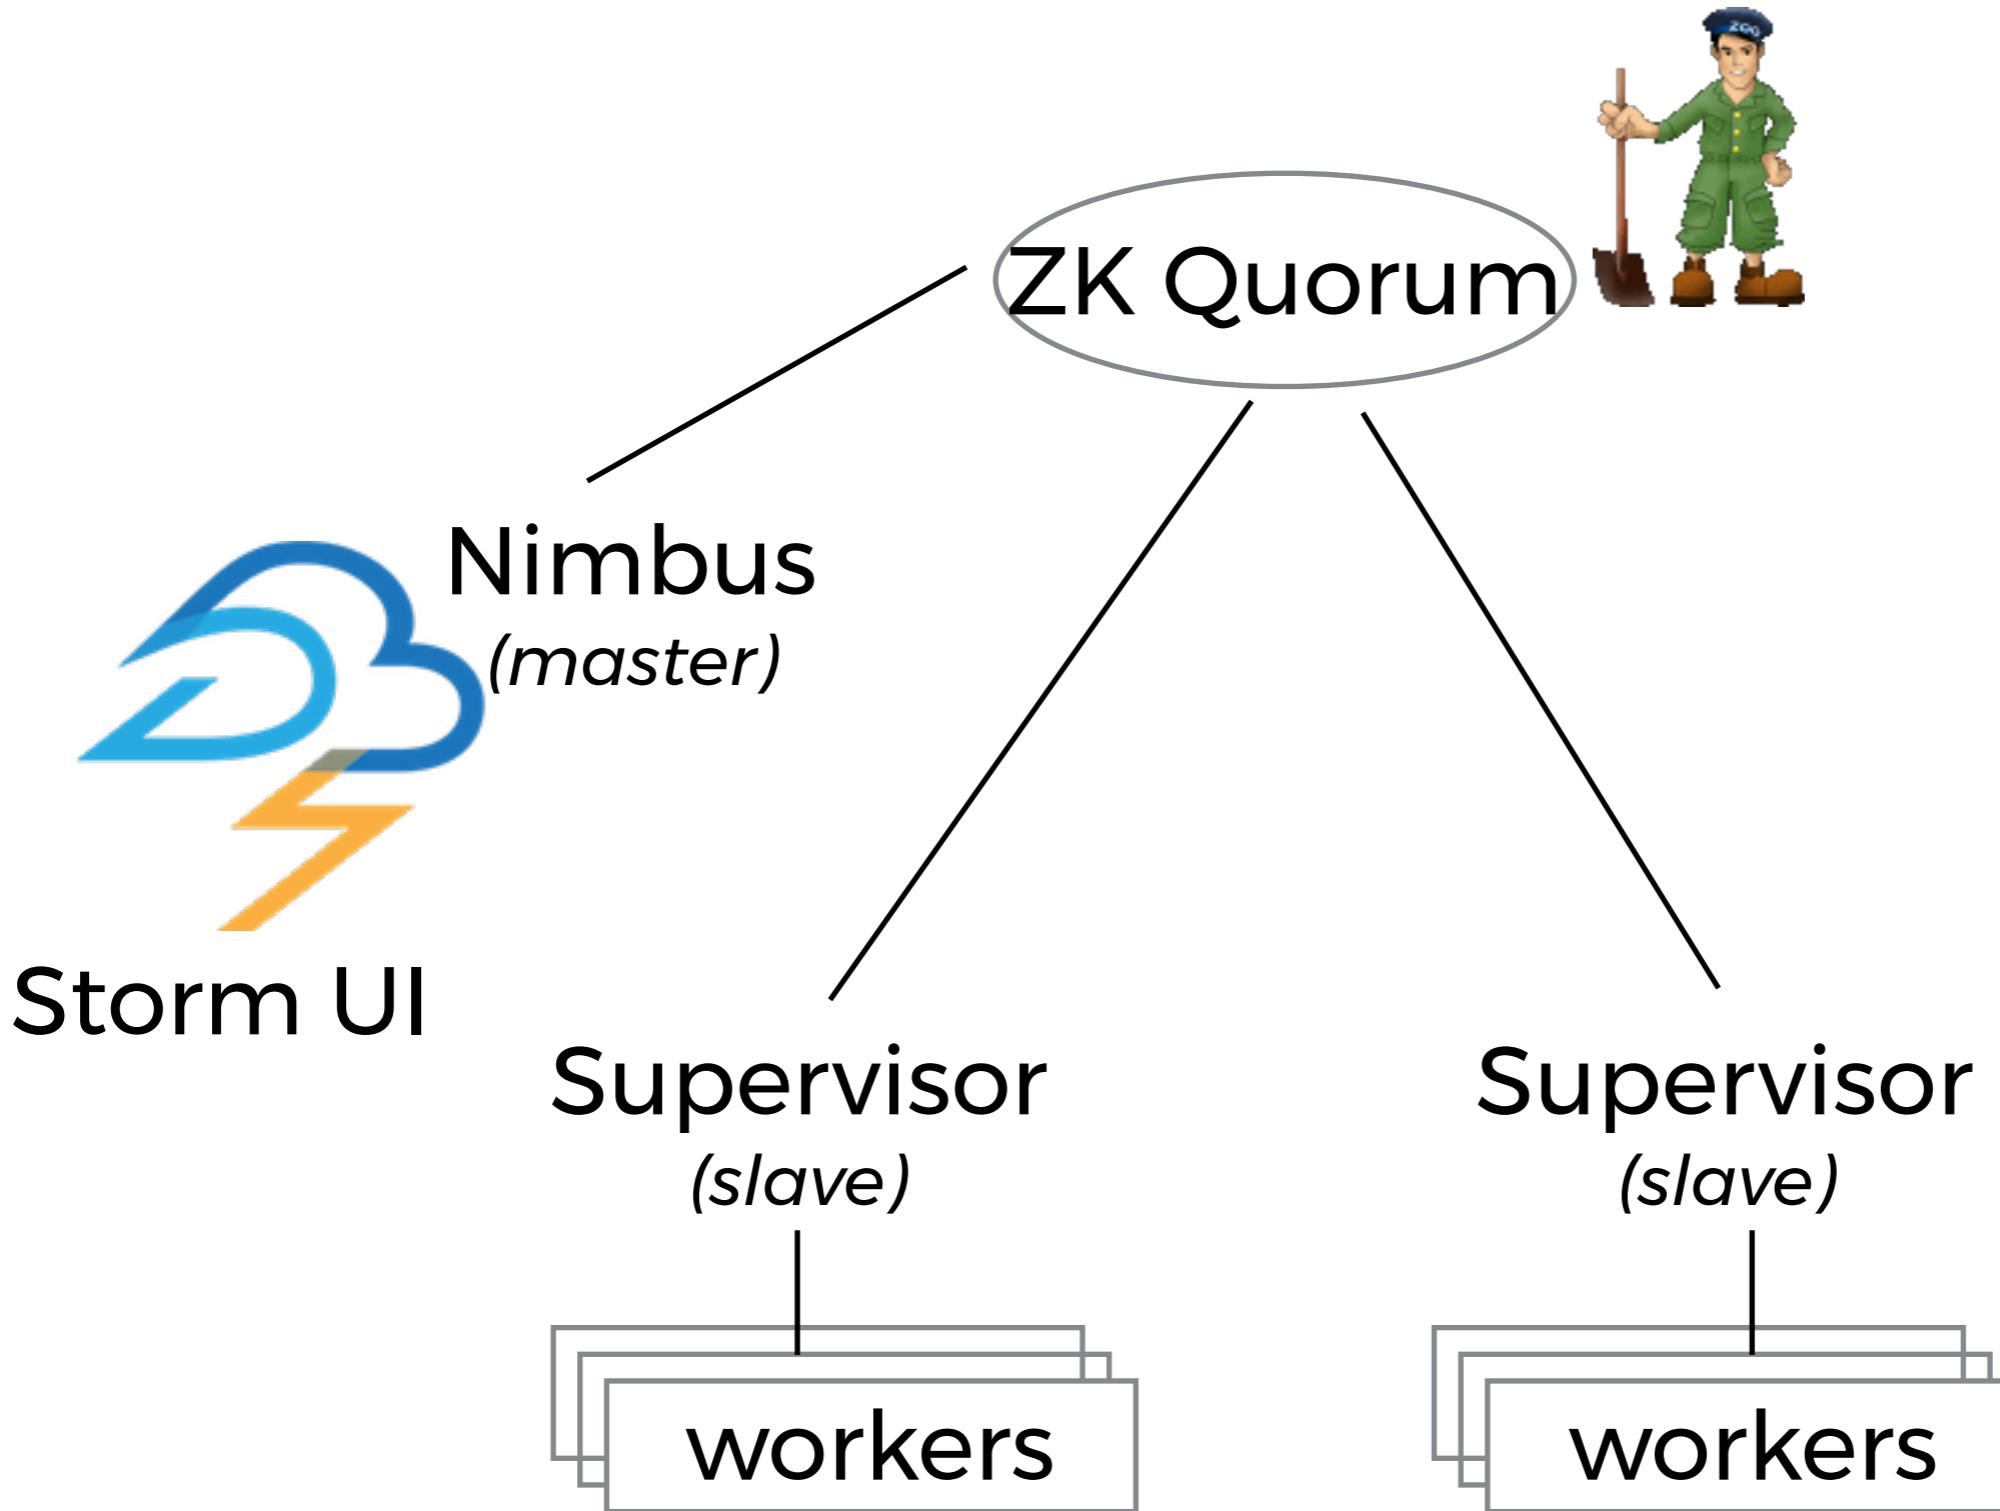
Storm integrates with the queueing and database technologies you already use. A Storm topology consumes streams of data and processes those streams in arbitrarily complex ways, repartitioning the streams between each stage of the computation however needed. Read more in the tutorial.

# Storm concepts



<http://storm.apache.org/>

# Storm architecture





## **Apache Storm**

1M+ messages/second on a 10 nodes cluster

mid-2017 @gcn:

*3 clusters*

*170 topologies*

# Do we really need that ?

100 000 households

100 endpoints/ site

15 minutes interval  
datapoints

**960 000 000** direct raw events /  
day

>10 000 events / second

Only talking about DIRECT  
events.

final figures will be more around  
50 000 events / second

but our reality is closer to:



100 000 households  
100 endpoints/ site  
30 seconds interval  
datapoints

**28 800 000 000** direct raw events /  
day  
> 350K events / second

Still wanna use MS Access ? :-)

some other frameworks exist / can be used



Apache Spark is a fast and general engine for large-scale data processing.



Akka is an open-source toolkit and runtime simplifying the construction of concurrent and distributed applications on the JVM.



Heron is a distributed streaming processing engine developed at Twitter. According to the creators at Twitter, the scale and diversity of Twitter data has increased, and Heron is a real-time analytics platform to process streaming.



Apache Samza is a distributed stream processing framework. It uses Apache Kafka for messaging, and Apache Hadoop YARN to provide fault tolerance, processor isolation, security, and resource management.



**STORM**

**Apache Storm**

1M+ messages/second on a 10 nodes cluster

mid-2017 @gcn:  
*3 clusters*  
*170 topologies*

did you forget  
to tell us something?





**STORM**

**Apache Storm**

1M+ messages/second on a 10 nodes cluster

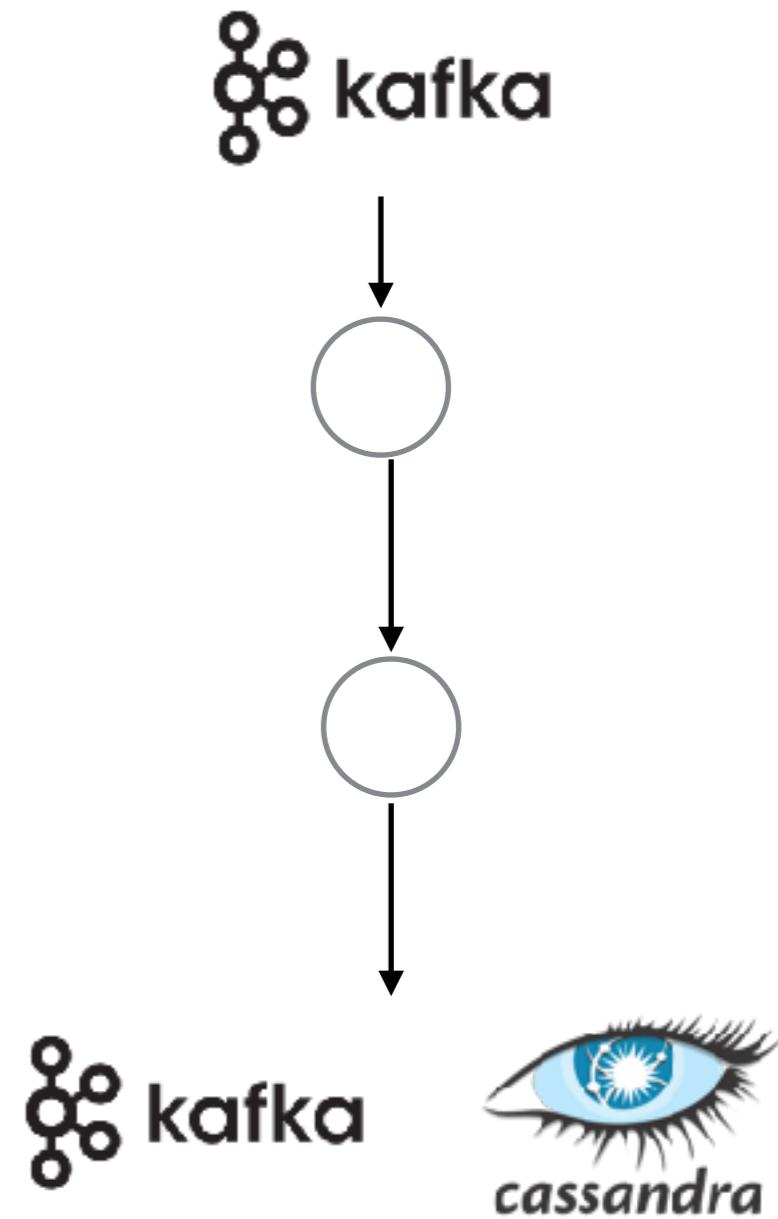
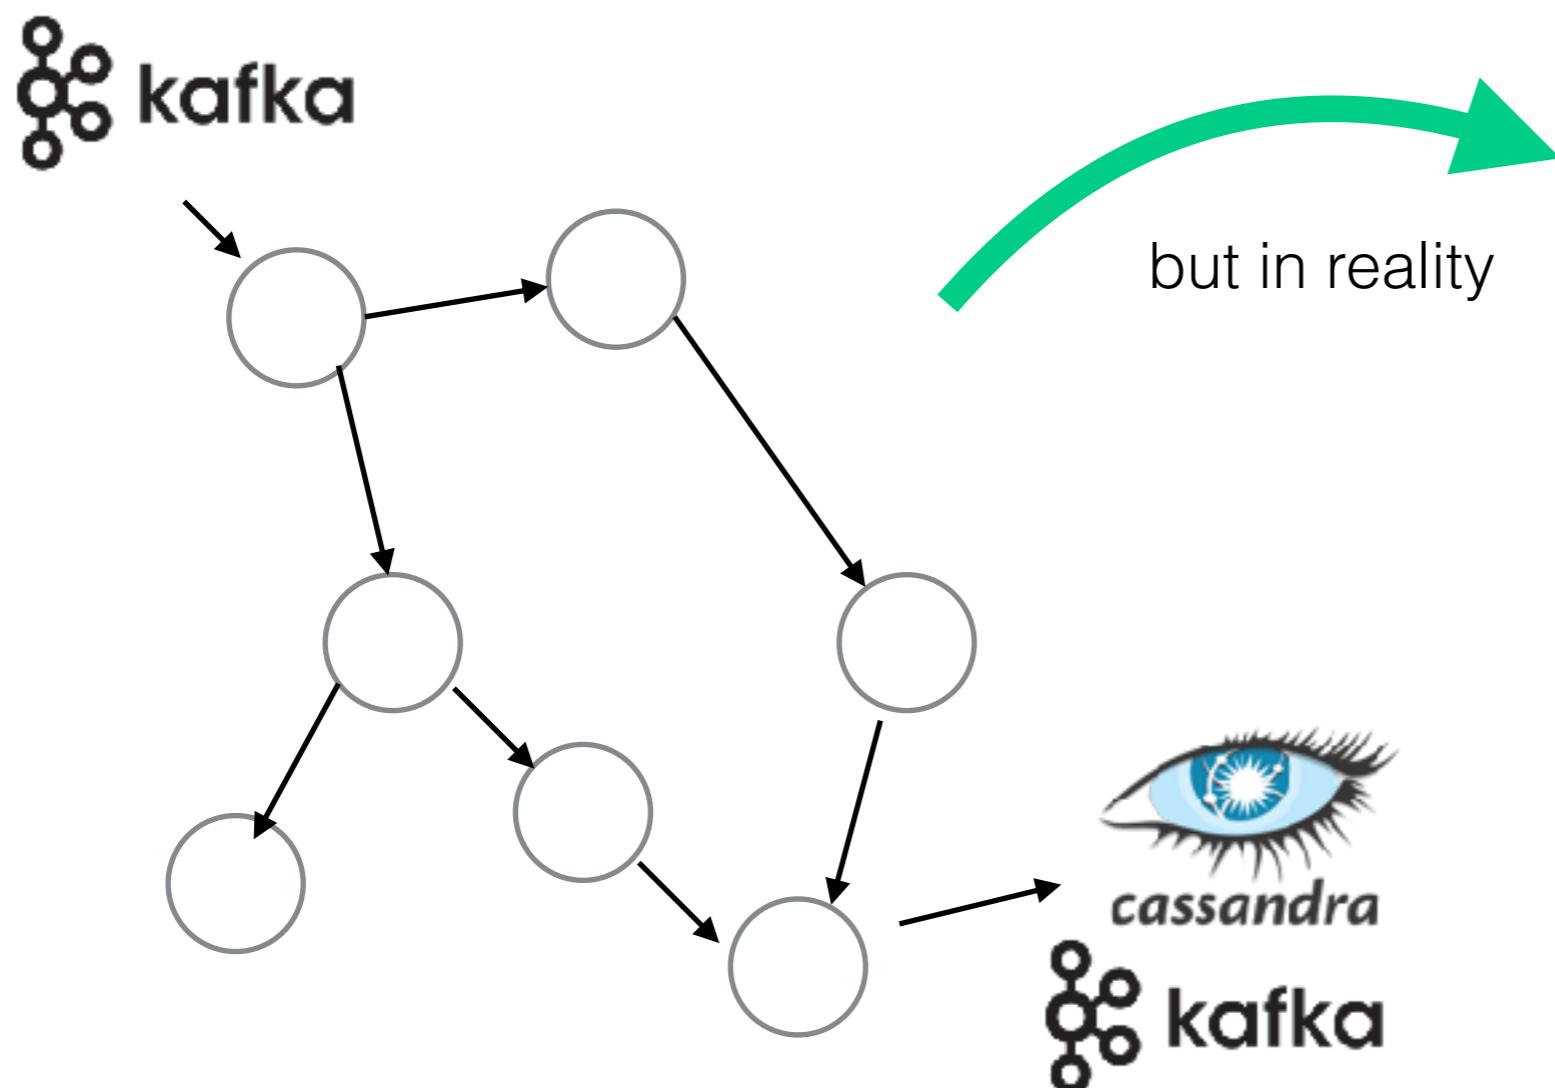
end 2017 @gcn:  
*0 clusters*  
*0 topologies*

oups...  
but why ?



# Why ?

## Un-needed complexity due to poor DAG



why ?



## Logging

Where are my logs ?

Not so dynamic scalability (in our case)

topology.yaml

# Why ?



## Yet an additional cluster



So now?

So what are you using now ?



Elastic, highly scalable, fault-tolerant  
Deploy to containers, VMs, bare metal, cloud  
Equally viable for small, medium, & large use cases  
Write standard Java applications  
No separate processing cluster required

# Code Sample



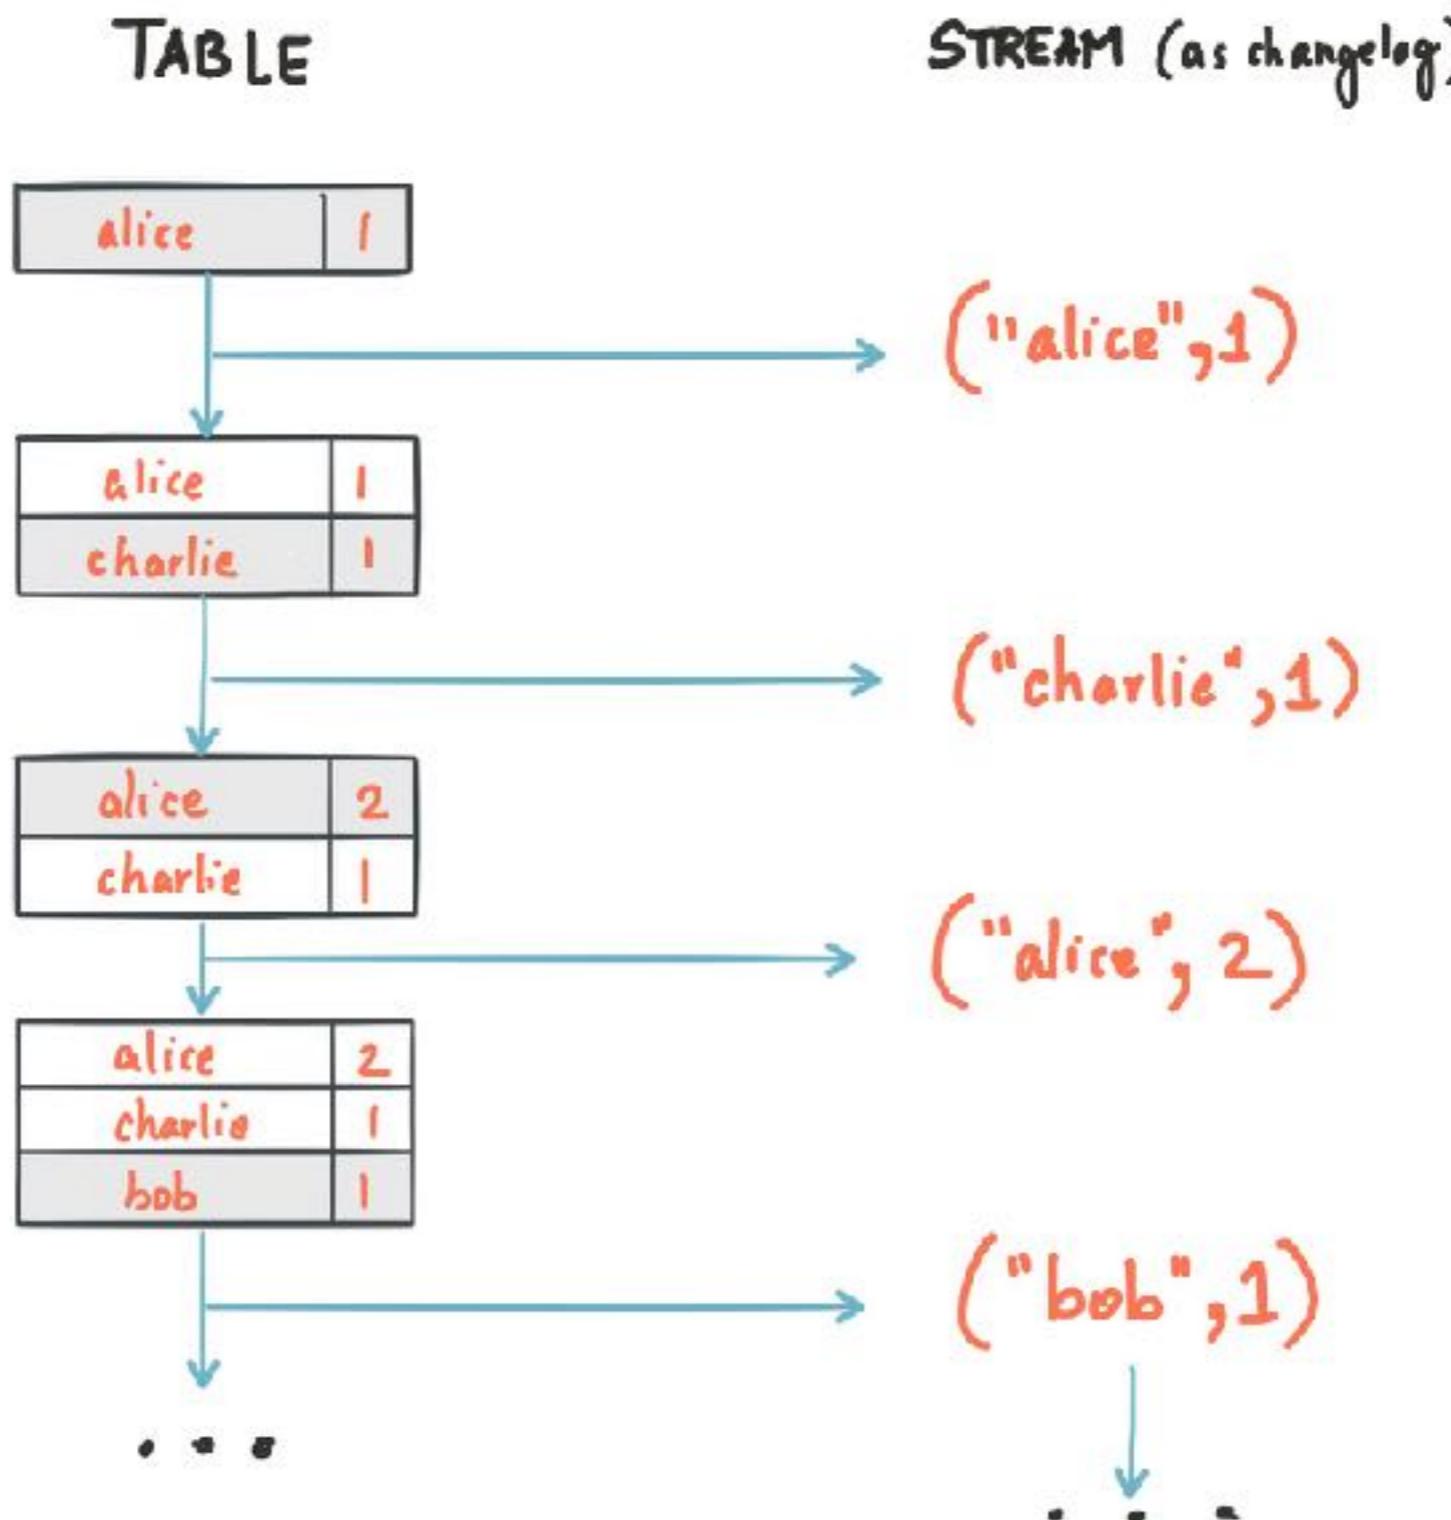
```
1.     public static void main(final String[] args) throws Exception {
2.         Properties config = new Properties();
3.         config.put(StreamsConfig.APPLICATION_ID_CONFIG, "simple-application");
4.         config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "10.46.2.91:9092,10.46.2.92:9092,10.46.2.93:9092");
5.         config.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
6.         config.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass());
7.         config.put(StreamsConfig.DEFAULT_TIMESTAMP_EXTRACTOR_CLASS_CONFIG, WallclockTimestampExtractor.class);
8.
9.
10.        KStreamBuilder builder = new KStreamBuilder();
11.        KStream<String, String> messageStream = builder.stream("testcorp.datapoints");
12.        messageStream.filter(new Predicate<String, String>() {
13.
14.            @Override
15.            public boolean test(String key, String value) {
16.                return (key != null && key.contains("Power"));
17.            }
18.        }).foreach(new ForeachAction<String, String>() {
19.
20.            @Override
21.            public void apply(String key, String value) {
22.                System.out.println("Key: " + key);
23.                System.out.println("Value: " + value);
24.            }
25.
26.        });
27.
28.        KafkaStreams streams = new KafkaStreams(builder, config);
29.        streams.start();
30.
31.        Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
32.    }
```

# Code Sample



```
1.     public static void main(final String[] args) throws Exception {
2.         Properties config = new Properties();
3.         config.put(StreamsConfig.APPLICATION_ID_CONFIG, "simple-application");
4.         config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "10.46.2.91:9092,10.46.2.92:9092,10.46.2.93:9092");
5.         config.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
6.         config.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass());
7.         config.put(StreamsConfig.DEFAULT_TIMESTAMP_EXTRACTOR_CLASS_CONFIG, WallclockTimestampExtractor.class);
8.
9.         KStreamBuilder builder = new KStreamBuilder();
10.        KStream<String, String> messageStream = builder.stream("testcorp.datapoints");
11.        messageStream.filter(new Predicate<String, String>() {
12.
13.            @Override
14.            public boolean test(String key, String value) {
15.                return (key != null && key.contains("Power"));
16.            }
17.        ).to("aNewTopic");
18.
19.        KafkaStreams streams = new KafkaStreams(builder, config);
20.        streams.start();
21.
22.        Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
23.    }
```

# Stream & Tables



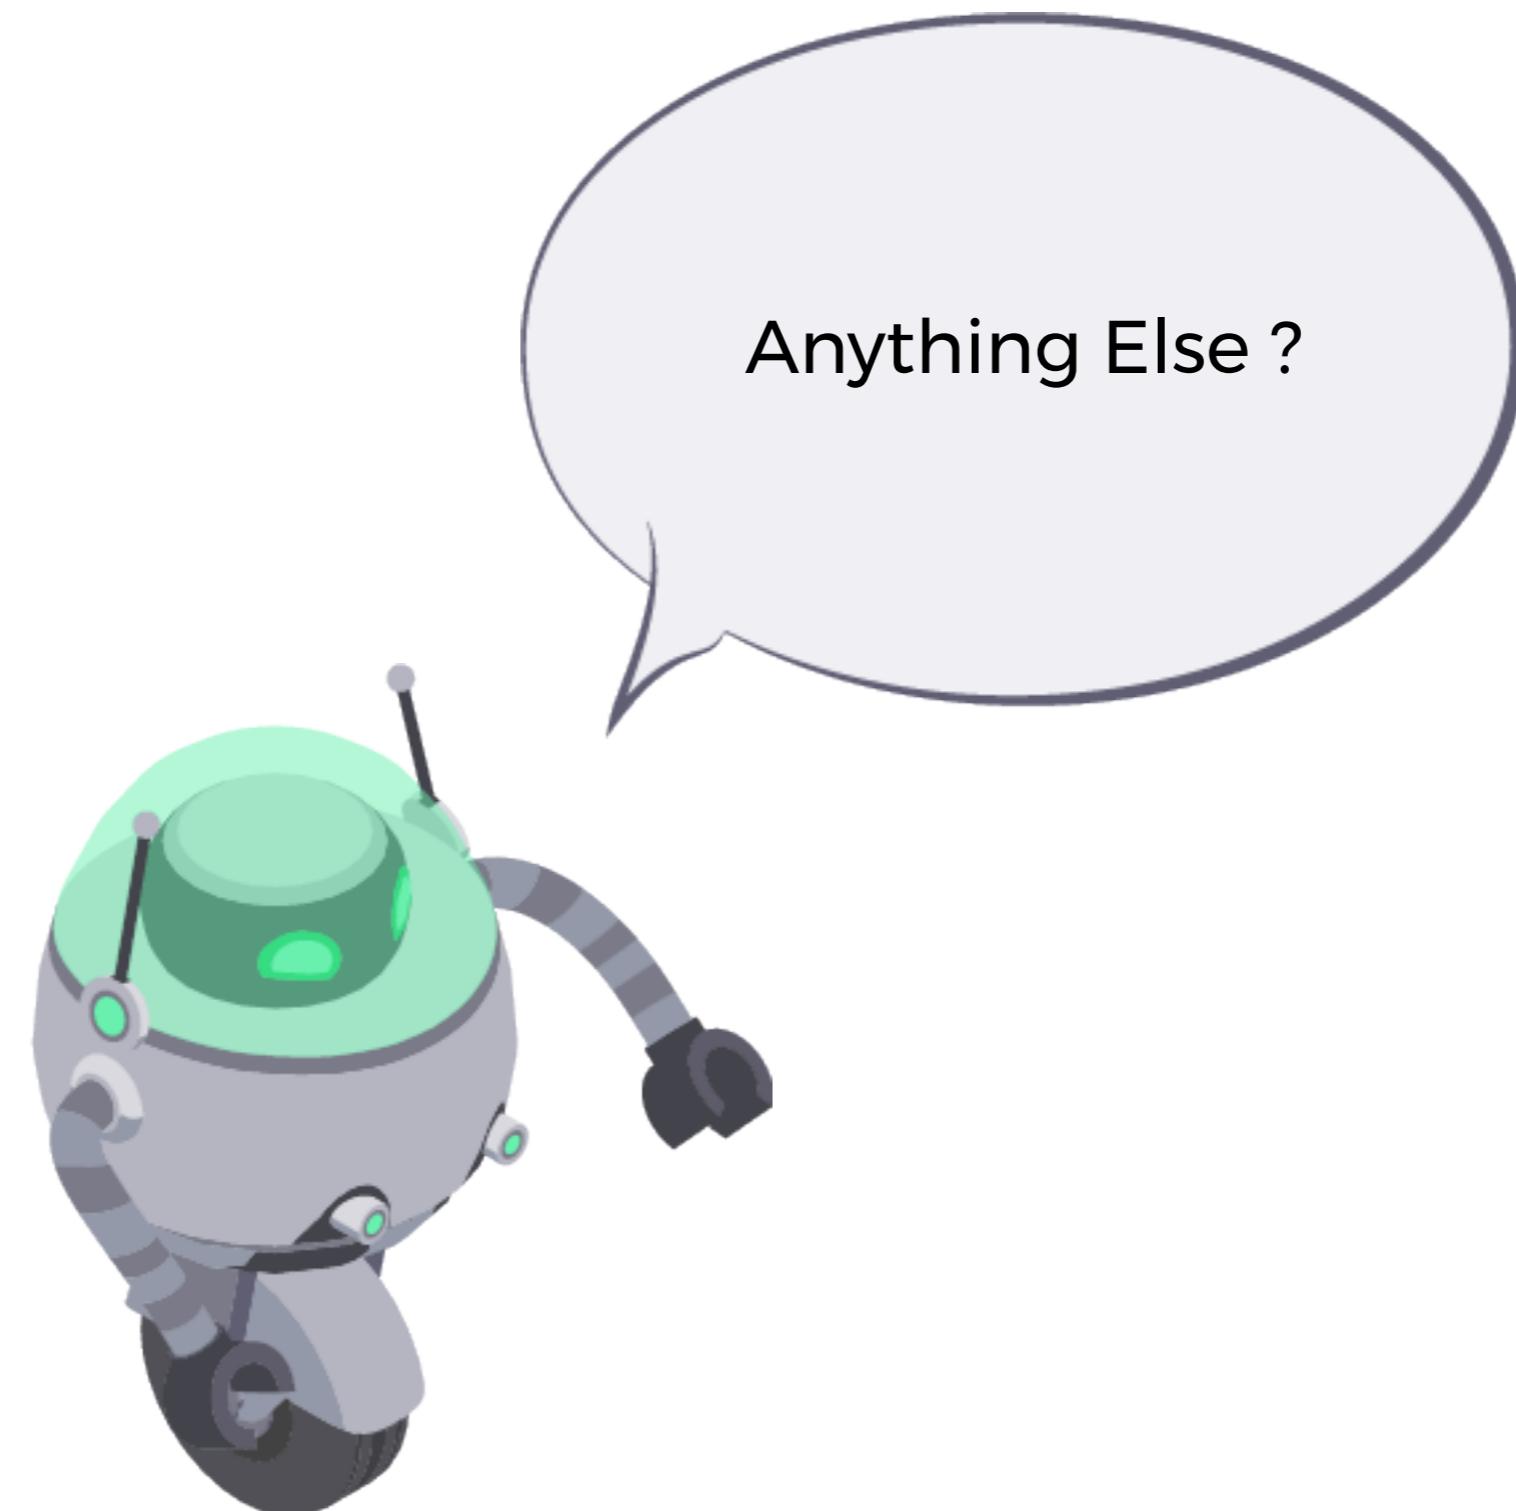
# Stream & Tables

```
1.     public static void main(final String[] args) throws Exception {
2.         Properties config = new Properties();
3.         config.put(StreamsConfig.APPLICATION_ID_CONFIG, "simple-application");
4.         config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "10.46.2.91:9092,10.46.2.92:9092,10.46.2.93:9092");
5.         config.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
6.         config.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass());
7.         config.put(StreamsConfig.DEFAULT_TIMESTAMP_EXTRACTOR_CLASS_CONFIG, WallclockTimestampExtractor.class);
8.
9.         KStreamBuilder builder = new KStreamBuilder();
10.        KStream<String, String> messageStream = builder.stream("testcorp.datapoints");
11.        KTable<String, String> ktable = builder.table("sensor-location");
12.
13.        messageStream.leftJoin(ktable, new ValueJoiner<String, String, String>() {
14.
15.            @Override
16.            public String apply(String value1, String value2) {
17.                return value1 + " measured at " + value2;
18.            }
19.        }).to("aNewTopic");
20.
21.        KafkaStreams streams = new KafkaStreams(builder, config);
22.        streams.start();
23.
24.        Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
25.    }
```

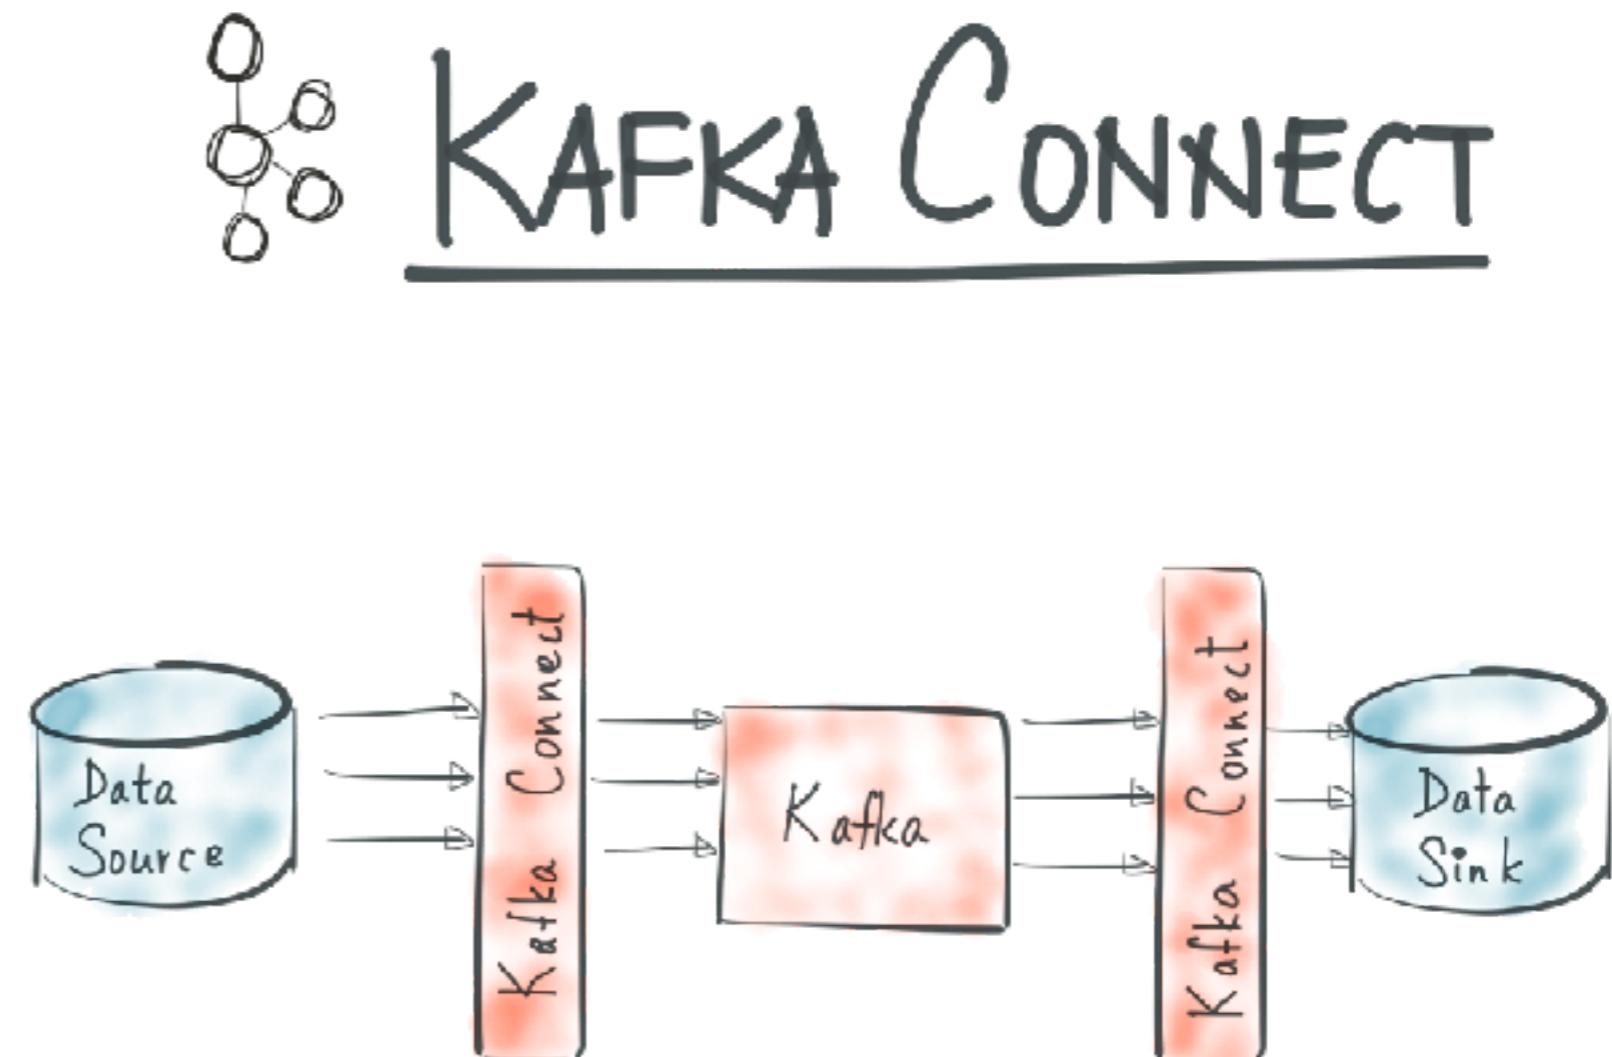
# Stream & Tables



```
1.     @Override
2.     protected KStreamBuilder kStreamBuilder( ) {
3.         final KStreamBuilder builder = new KStreamBuilder( );
4.         builder.globalTable( getCustomerTopic( "endpoints" ), STATE_STORE_NAME );
5.         return builder;
6.     }
7.
8.     private Endpoint retrieveEndpointFromStore( KafkaStreams stateStore, String uri ) {
9.         if( stateStore == null )
10.             return null;
11.         try {
12.             try {
13.                 Object jsonEdp = stateStore
14.                     .store( EntitiesStateStore.STATE_STORE_NAME, QueryableStoreTypes.keyValueStore( ) ).get( uri );
15.                 return gson.fromJson( jsonEdp.toString( ), Endpoint.class );
16.             } catch( NullPointerException npe ) {
17.                 populateGlobalTable( eibp, EndpointLink.fromURI( uri ) );
18.             }
19.         } catch( Exception ex ) {
20.             log.warn( ExceptionUtils.getStackTrace( ex ) );
21.         }
22.     }
23. }
```



- s3
- Elasticsearch
- HDFS
- JDBC
- SAP HANA
- Cassandra
- CoAP
- FTP
- Github
- InfluxDB
- MongoDB
- MQTT
- Twitter
- ...
- 



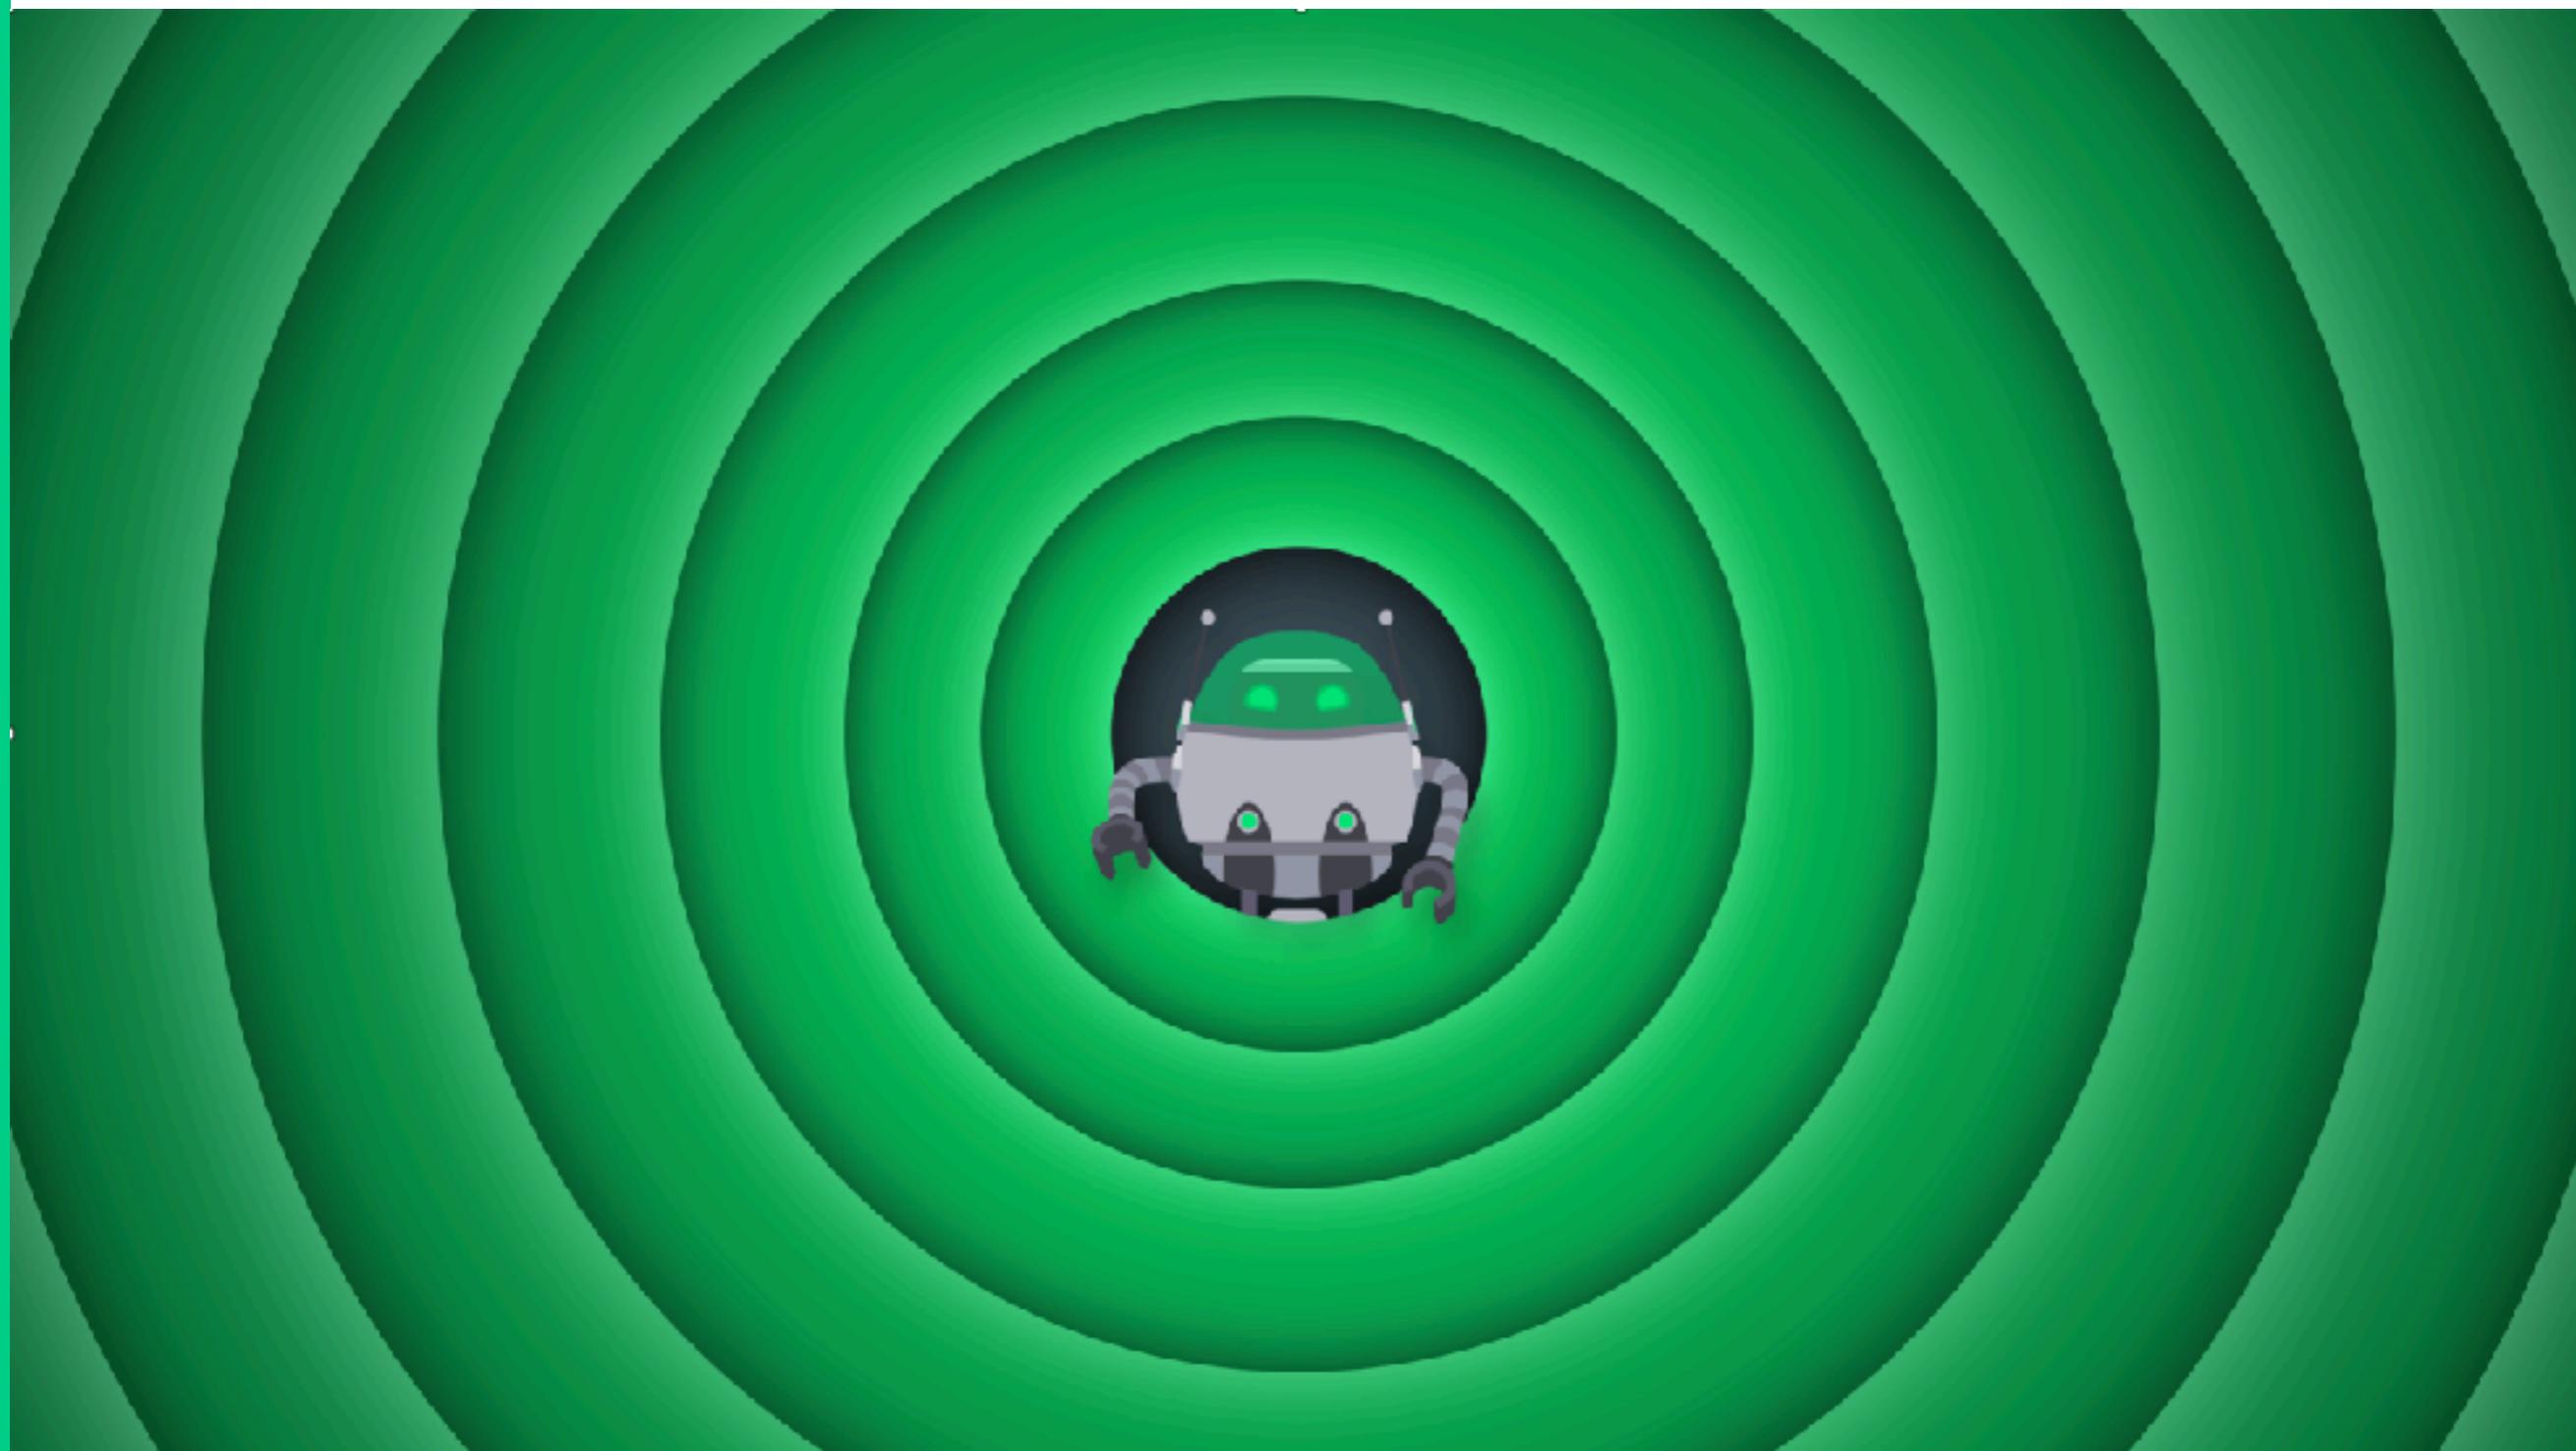
<https://www.confluent.io/product/connectors/>

# Exactly once semantic Message headers

KSQL is an open source streaming SQL engine for Apache Kafka.

```
CREATE TABLE possible_fraud AS  
SELECT card_number, count(*)  
FROM authorization_attempts  
WINDOW TUMBLING (SIZE 5 SECONDS)  
GROUP BY card_number  
HAVING count(*) > 3;
```

```
CREATE STREAM vip_actions AS  
SELECT userid, page, action  
FROM clickstream c  
LEFT JOIN users u ON c.userid = u.user_id  
WHERE u.level = 'Platinum';
```



**Oh... Wait...**

**How does all of this  
scale/survive ?**

**From pet to cattle**



✓ scales  
✓ survives



✗ scales  
✗ survives



✓ scales  
✓ survives



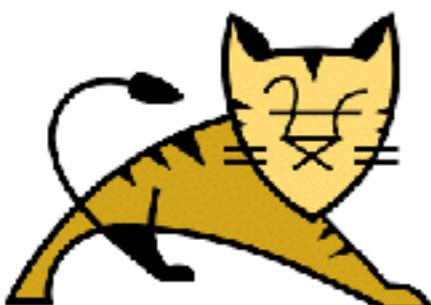
✓ scales  
✓ survives



✓ scales  
✓ survives

## Misc Stuff

✗ scales  
✗ survives



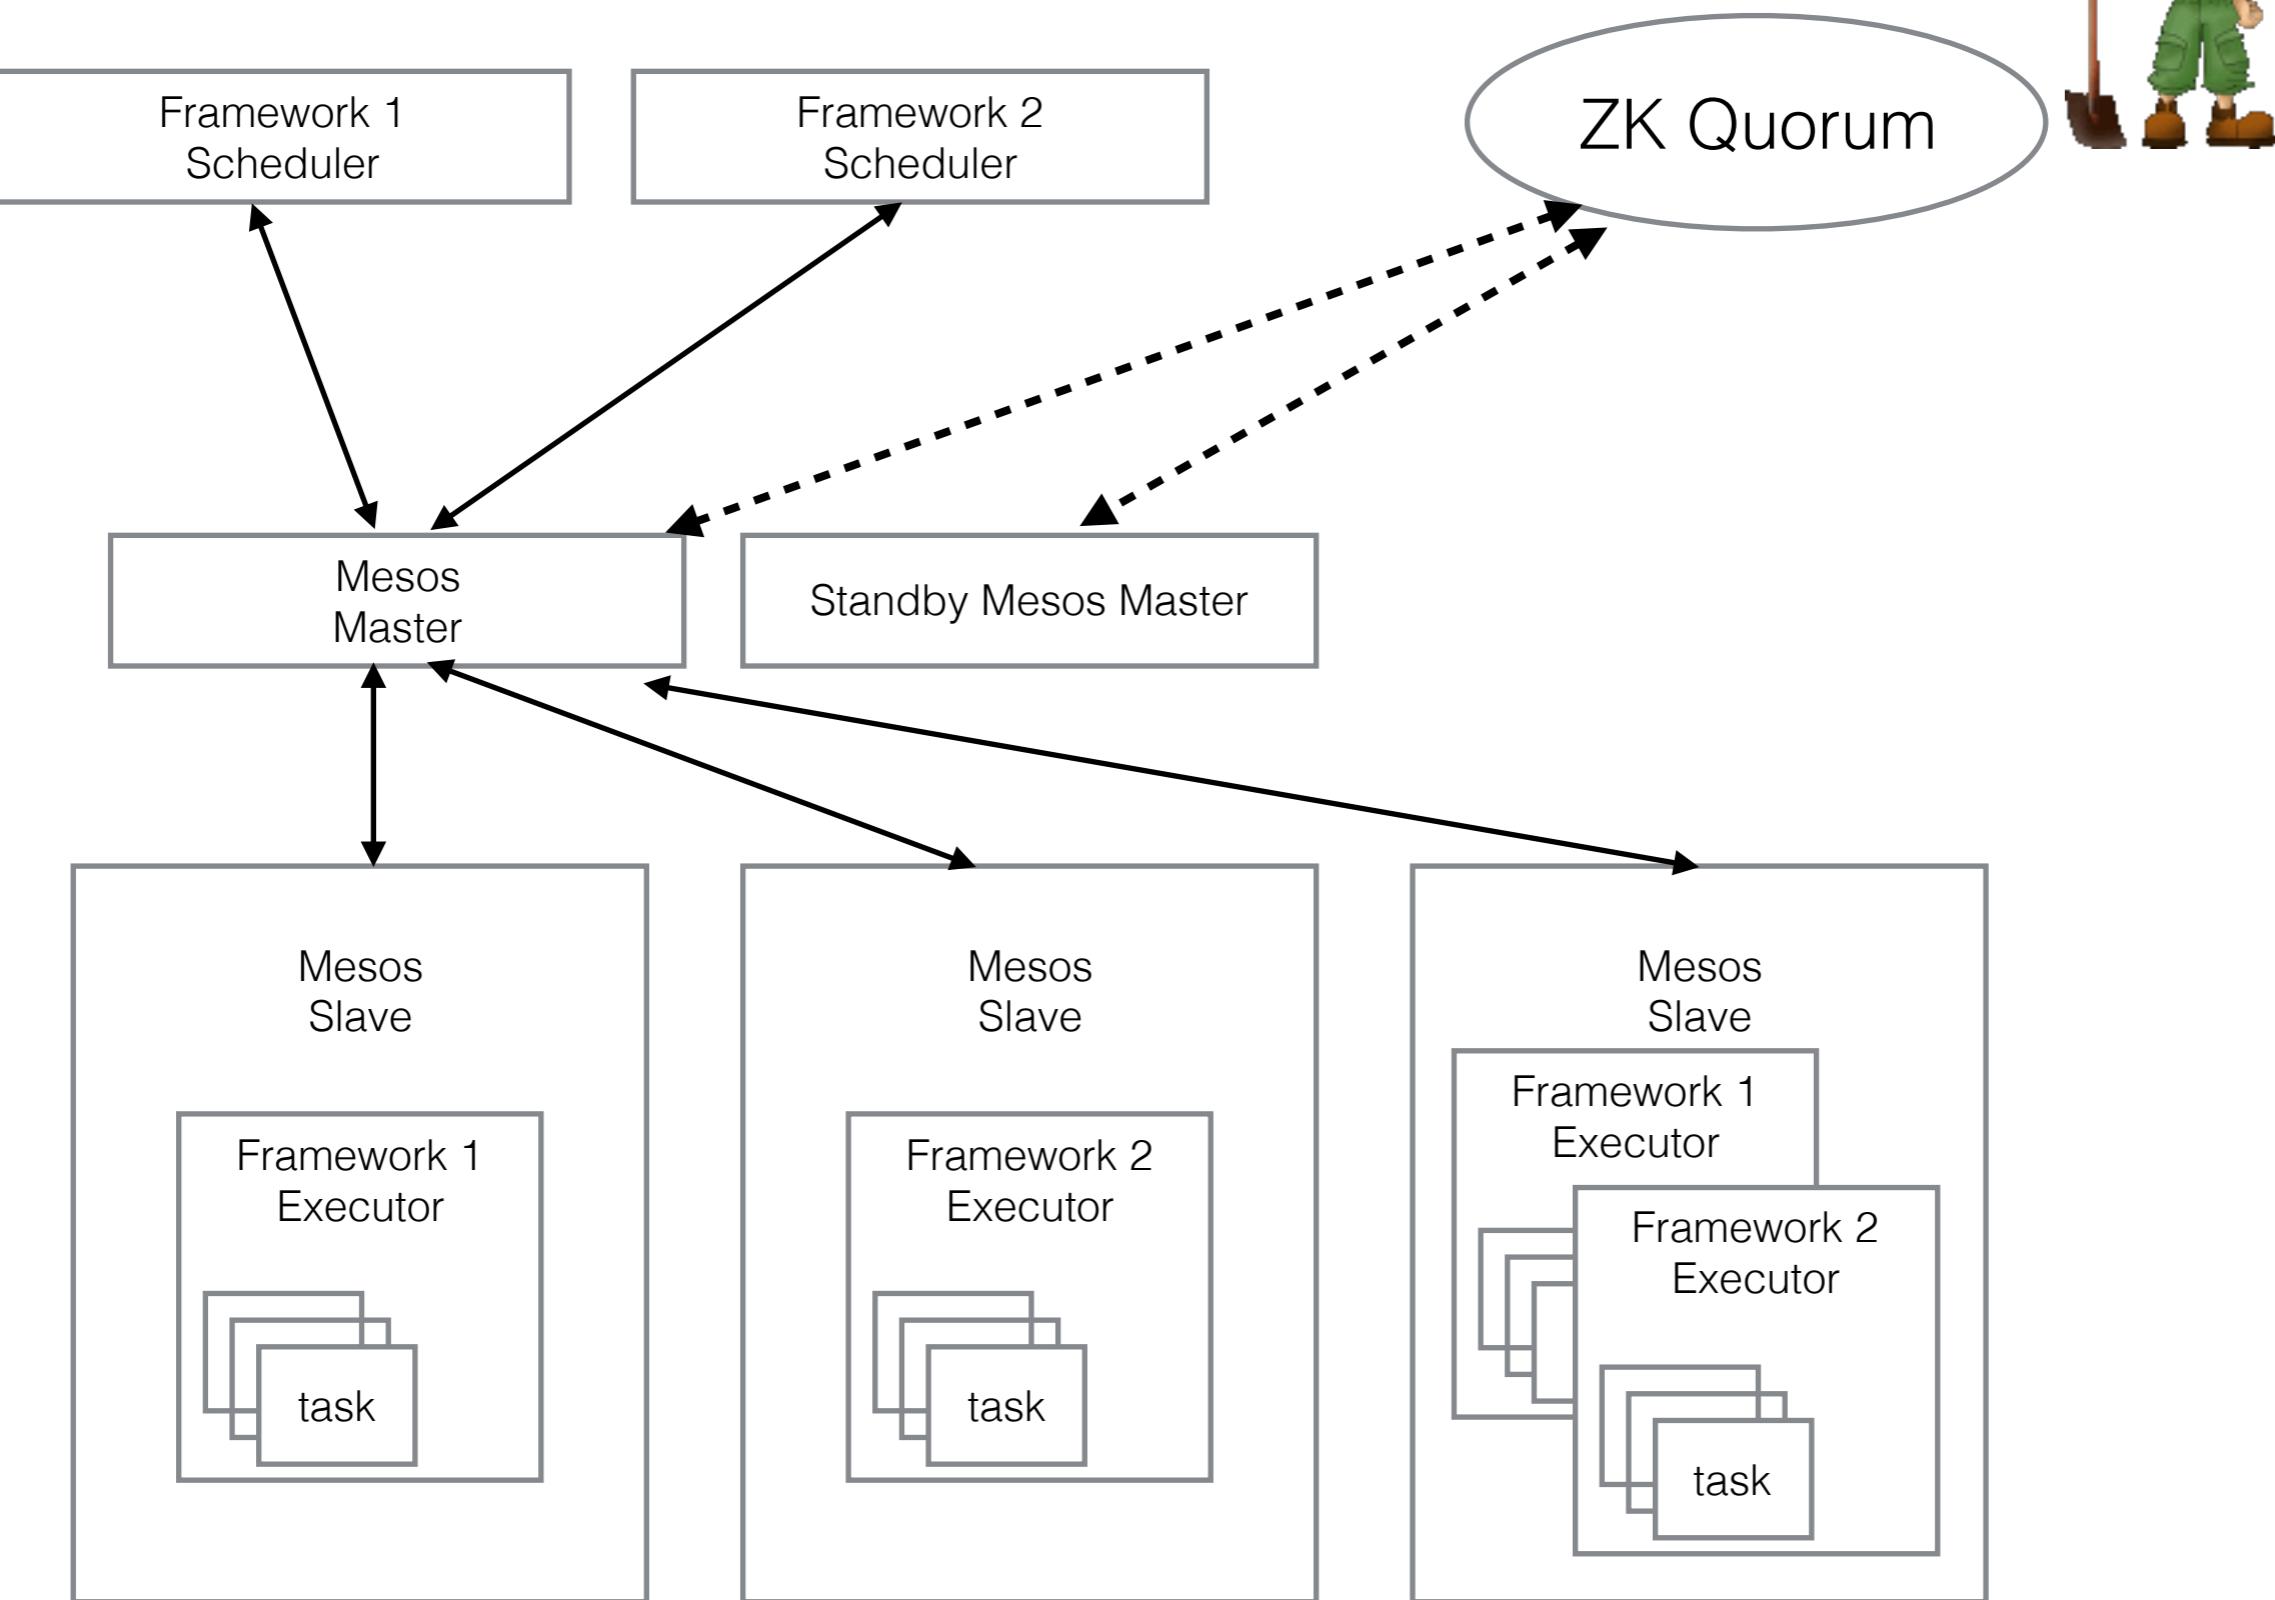
✓ scales  
✓ survives



# MESOS

Apache Mesos is an open-source cluster manager. It "provides efficient resource isolation and sharing across distributed applications, or frameworks". The software enables resource sharing in a fine-grained manner, improving cluster utilization.

# Mesos architecture



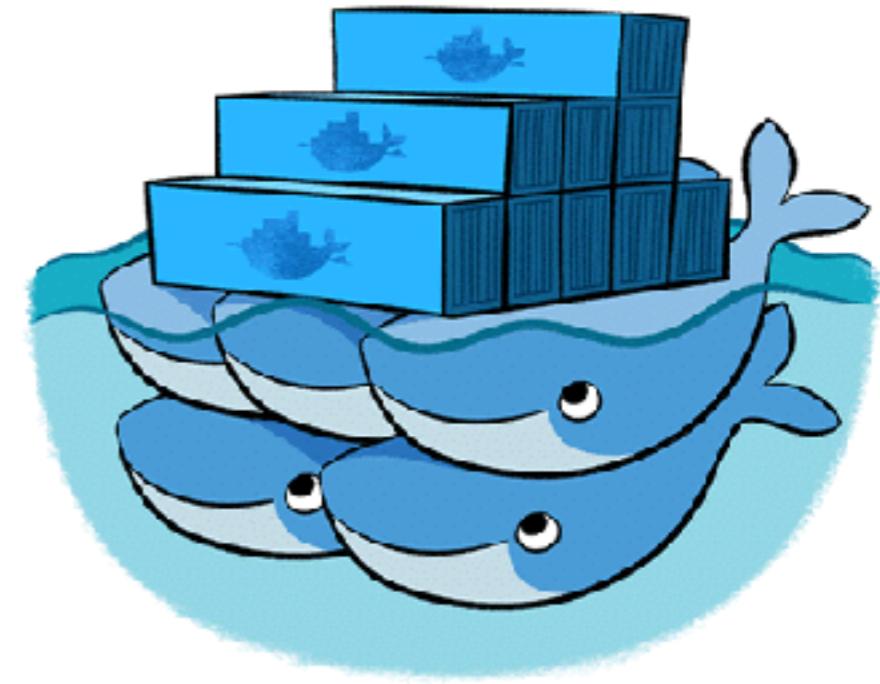
Full abstraction of physical hardware or virtual machine.

Data center is view as a pool of resources

Resources	CPUs	GPUs	Mem	Disk
Total	120	0	552.3 GB	32.8 TB
Allocated	94.7	0	413.9 GB	323.4 GB
Offered	0	0	0 B	0 B
Idle	25.3	0	138.4 GB	32.5 TB

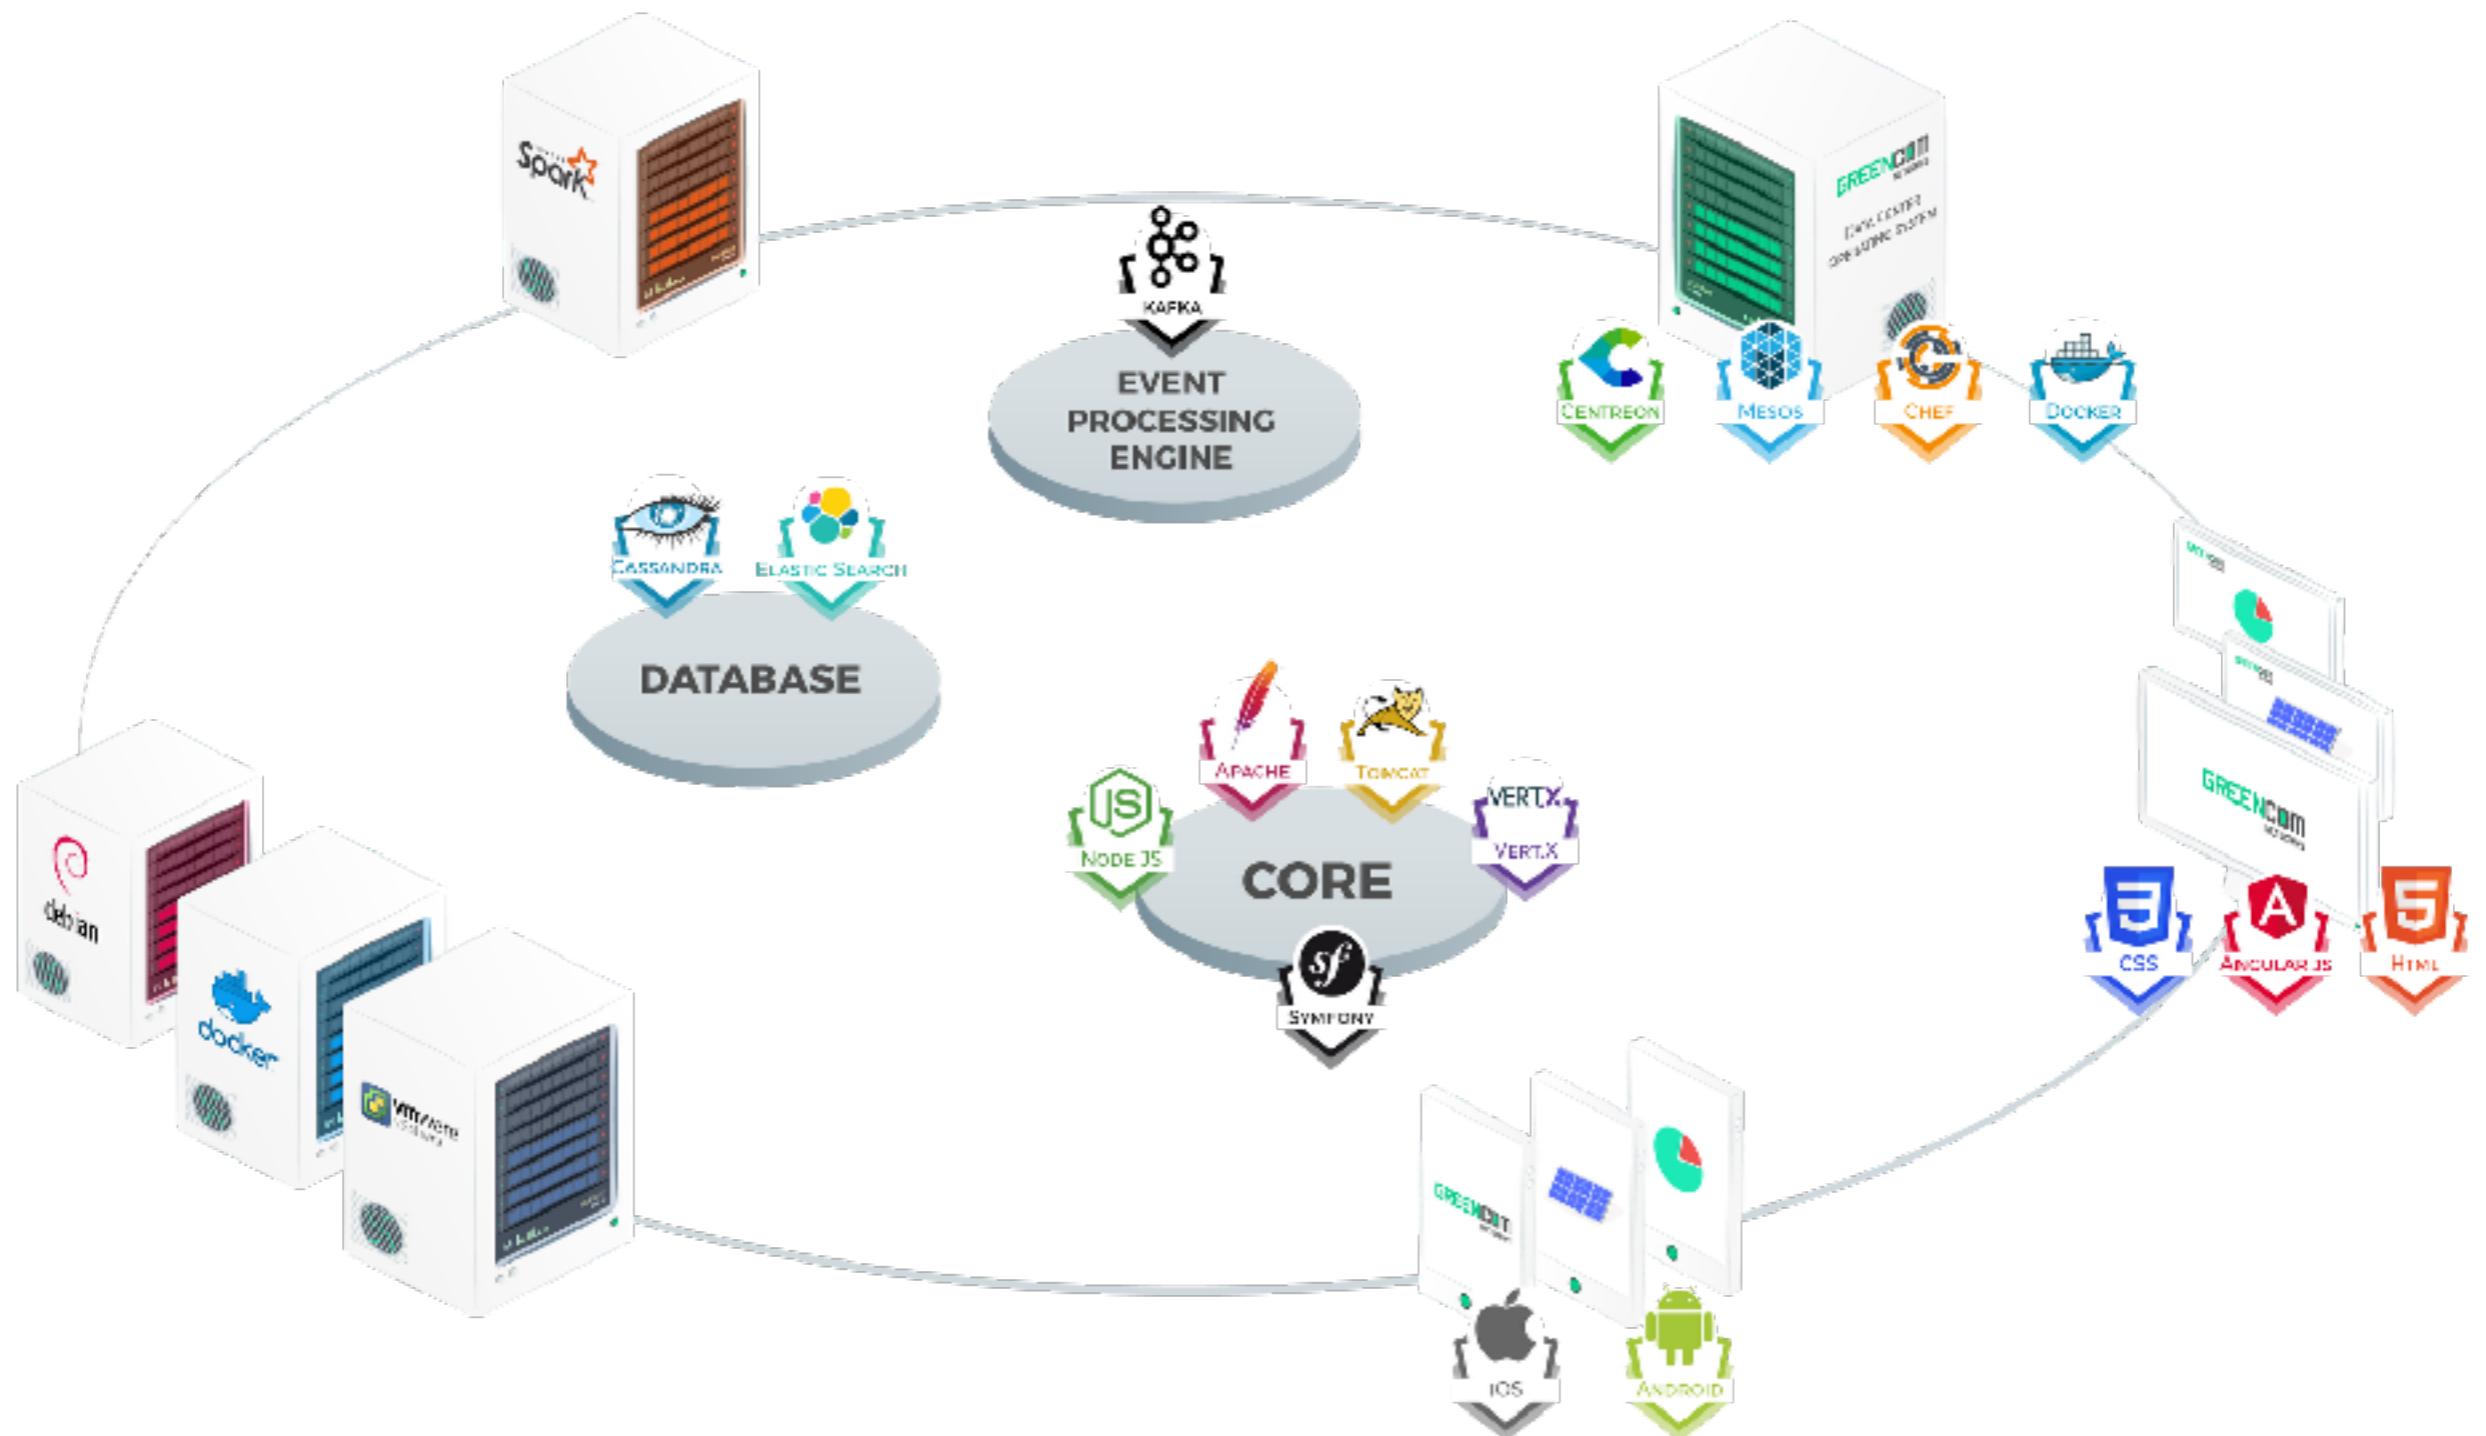
Frameworks negotiate with mesos masters to get access to resources

# Competition



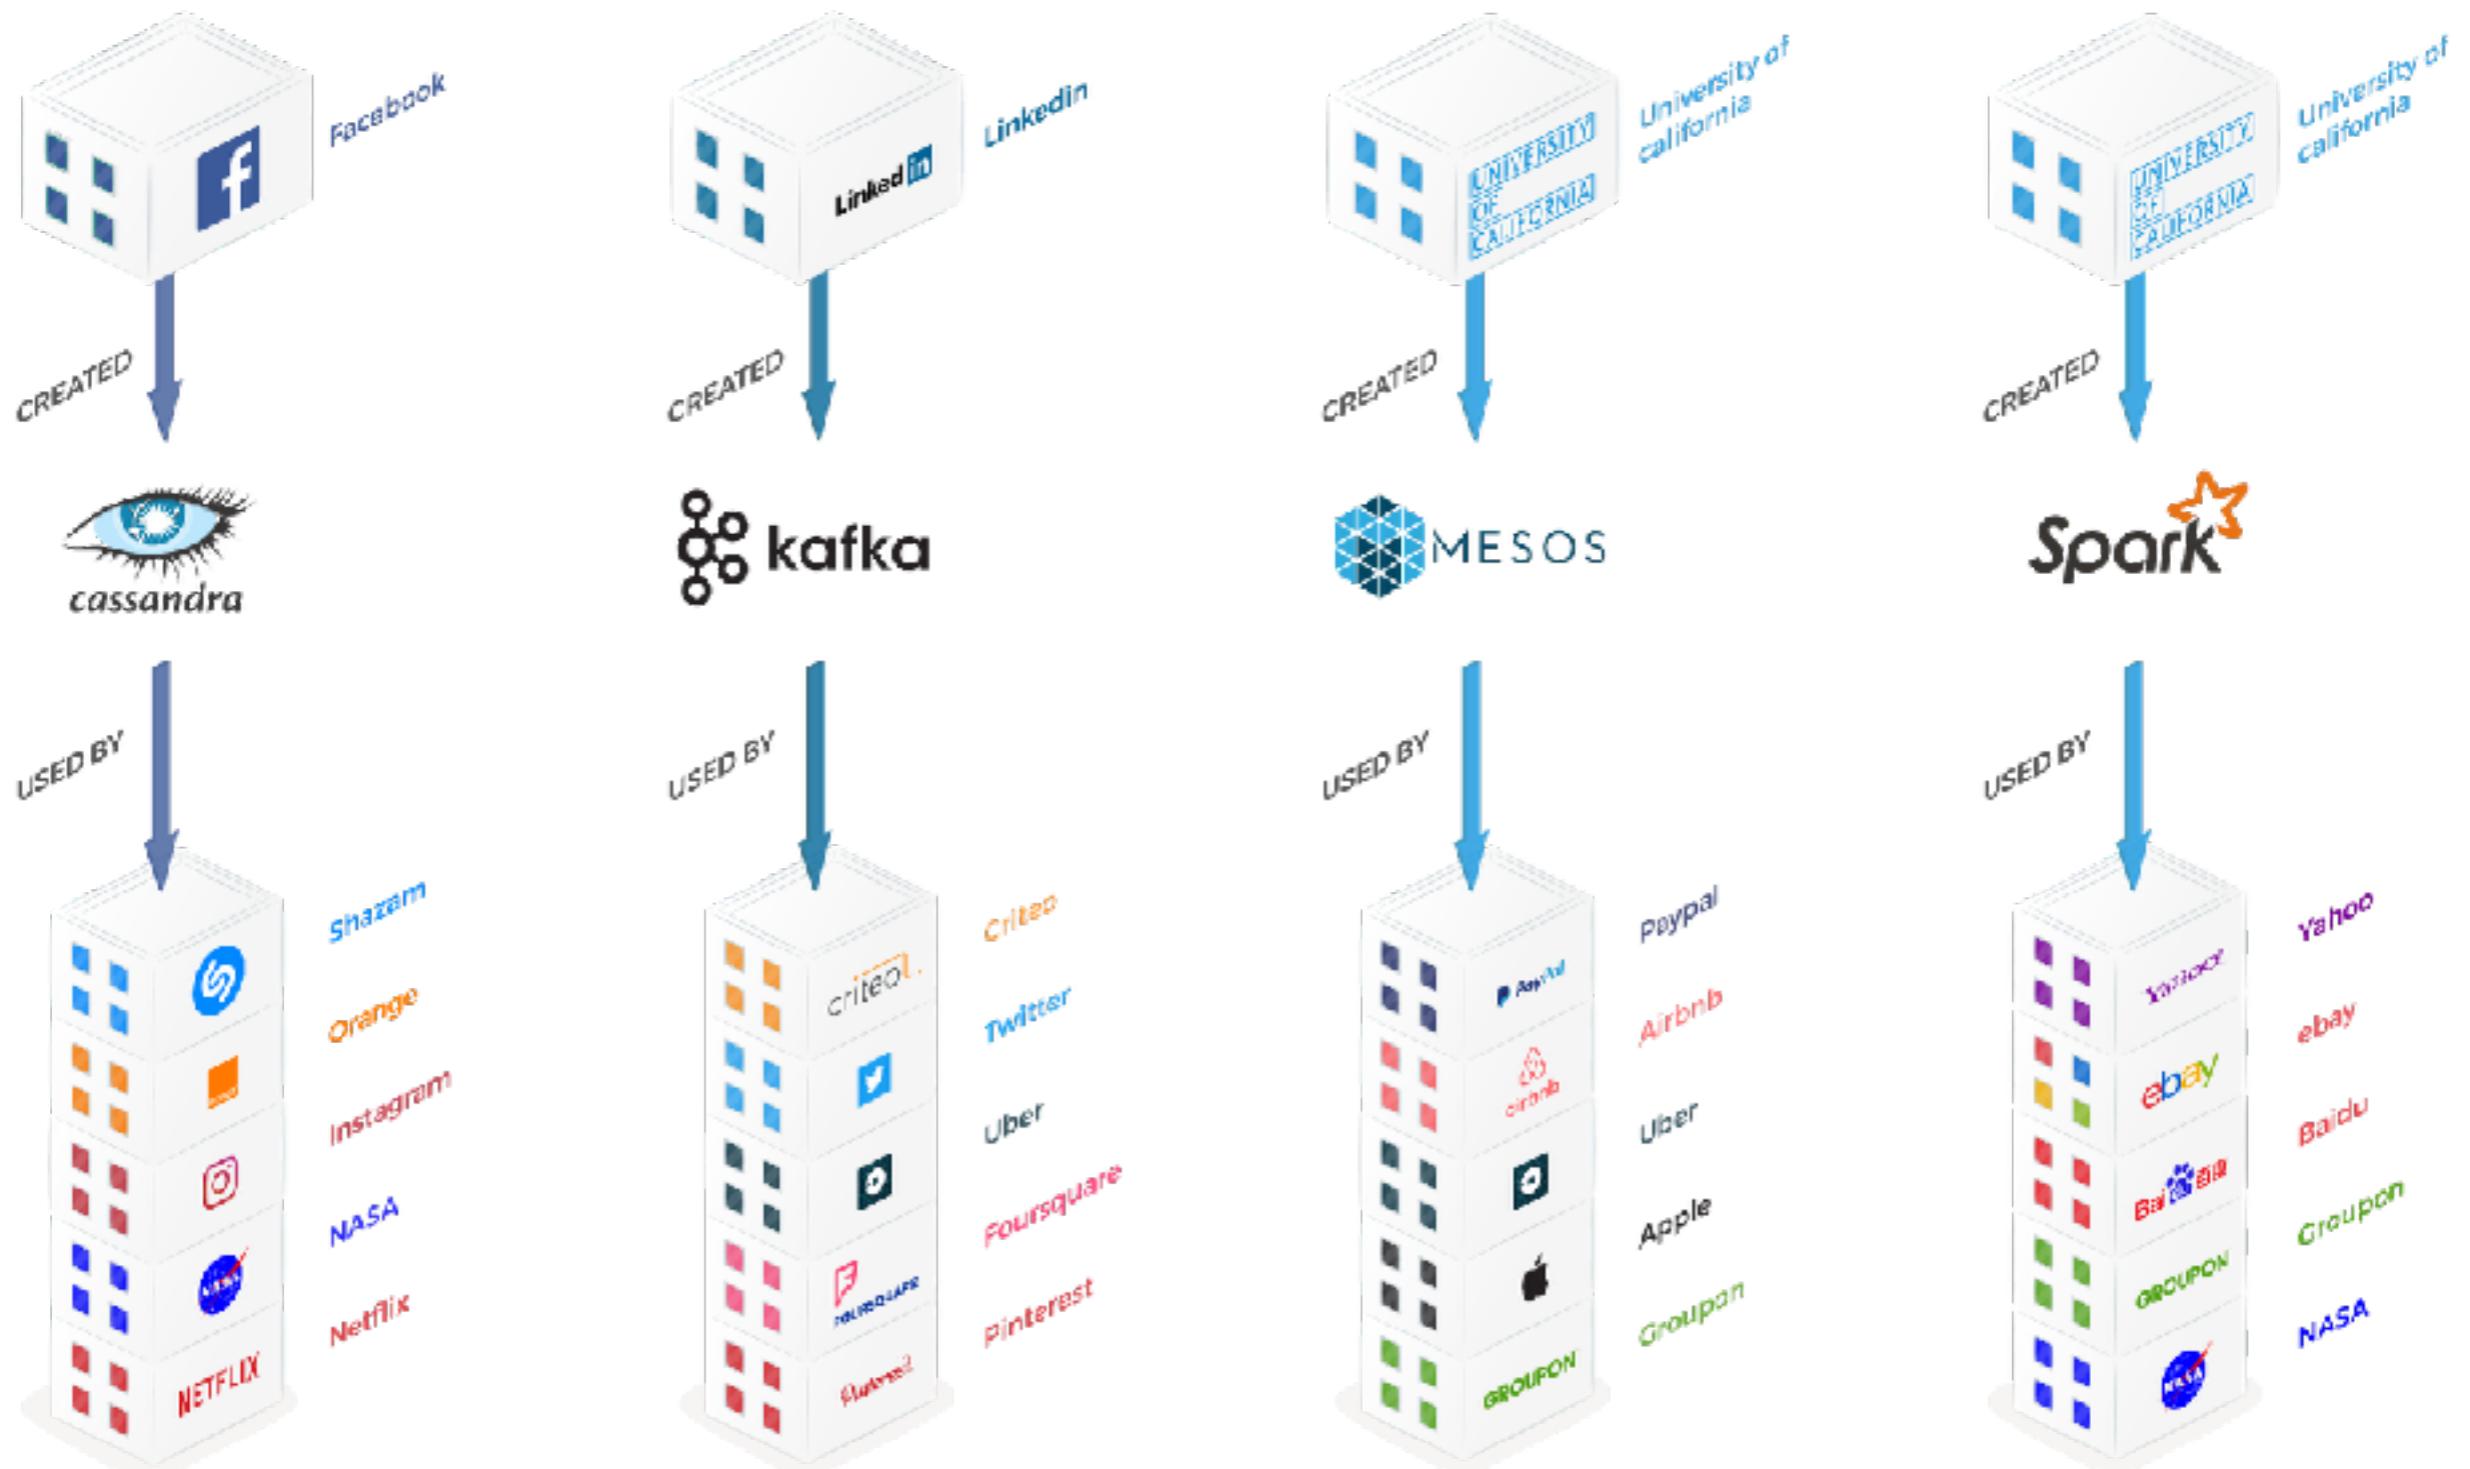
**nextGen Platform**  
aka  
**EIBP Platform V3**

# EIBP Platform V3



**GREENCOM**  
NETWORKS

# Our stack



# Few of the main components

- **Cassandra**: open source NoSQL distributed database management system (brings high availability with no single point of failure; scales horizontally). *Originally developed by Facebook.*
- **Kafka**: open source message broker (brings a unified, high-throughput, low-latency platform for handling real-time data feeds). *Originally developed by LinkedIn*
- **Mesos**: open source cluster manager (brings resource sharing capabilities in a highly distributed environment) .*Originally developed at the University of California, Berkeley. Used by Twitter, Airbnb, Apple*
- **Spark**: open source cluster computing framework (brings high performance in memory data processing capabilities). *Originally developed at UC Berkeley*

# Some figures:

## Production:

### Cassandra:

- 16 nodes
- 50 millions calls/day on low level API
- ~ 40 millions datapoints/day

### Kafka:

- 8 brokers
- 1700 topics / 7000 partitions
- 3500 messages/second (300 millions / day)

350+ services

1000+ instances

20 customers

100+ daily deployments

120 servers

2-5 minutes from commit to prod

3.5 cloud engineers

# Open points/questions

- **Monitoring:** Hard to monitor so many micro services in an highly distributed environment
- **Deployment:** Can't afford manual operations. Everything as to be automated
- **Testing/diagnostic:** micro services + loose coupling creates complexity
- **Number of clusters/split brain:** Minimum number of server is already high. Many different schedulers/cluster managers. High complexity.
- **Skillset:** Not so easy to recruit people with the right skill set



**GREENCOM**  
NETWORKS

# Roadmap

## **EIBP v3+ : What we are working on:**

- Multi Cloud
- Mesos to Kubernetes migration
- Better service mesh management



My grandma told me  
the best cloud provider  
is  
XXX

# Going multi-cloud



**GREENCOM**  
NETWORKS

# Going multi-cloud



**GREENCOM**  
NETWORKS

# Going multi-cloud



# Kubernetes



# Service mesh management

• Zero Downtime deployment	✓
• Dark Launch	✓
• Logging	✓
• A/B Testing	✓
• API Management	✓
• Tracing	✓
• Canary Release	✗
• Blue/Green Deployment	✗

# Container orchestrator

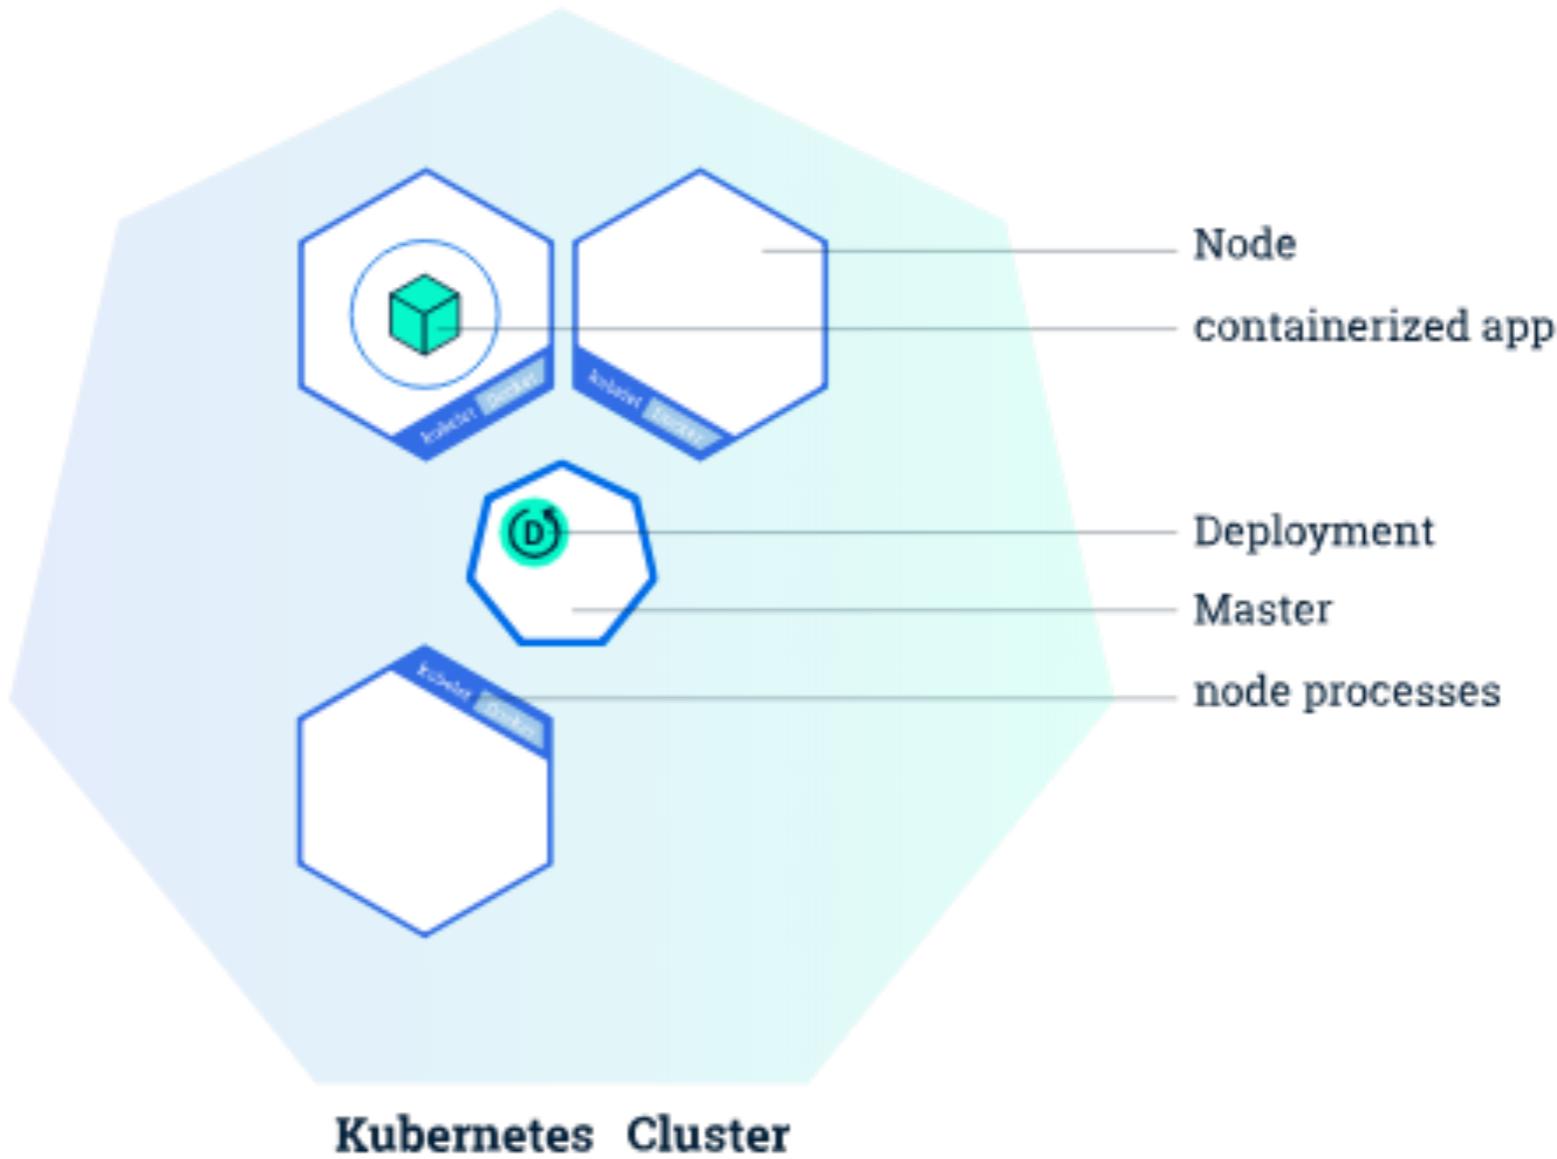
---

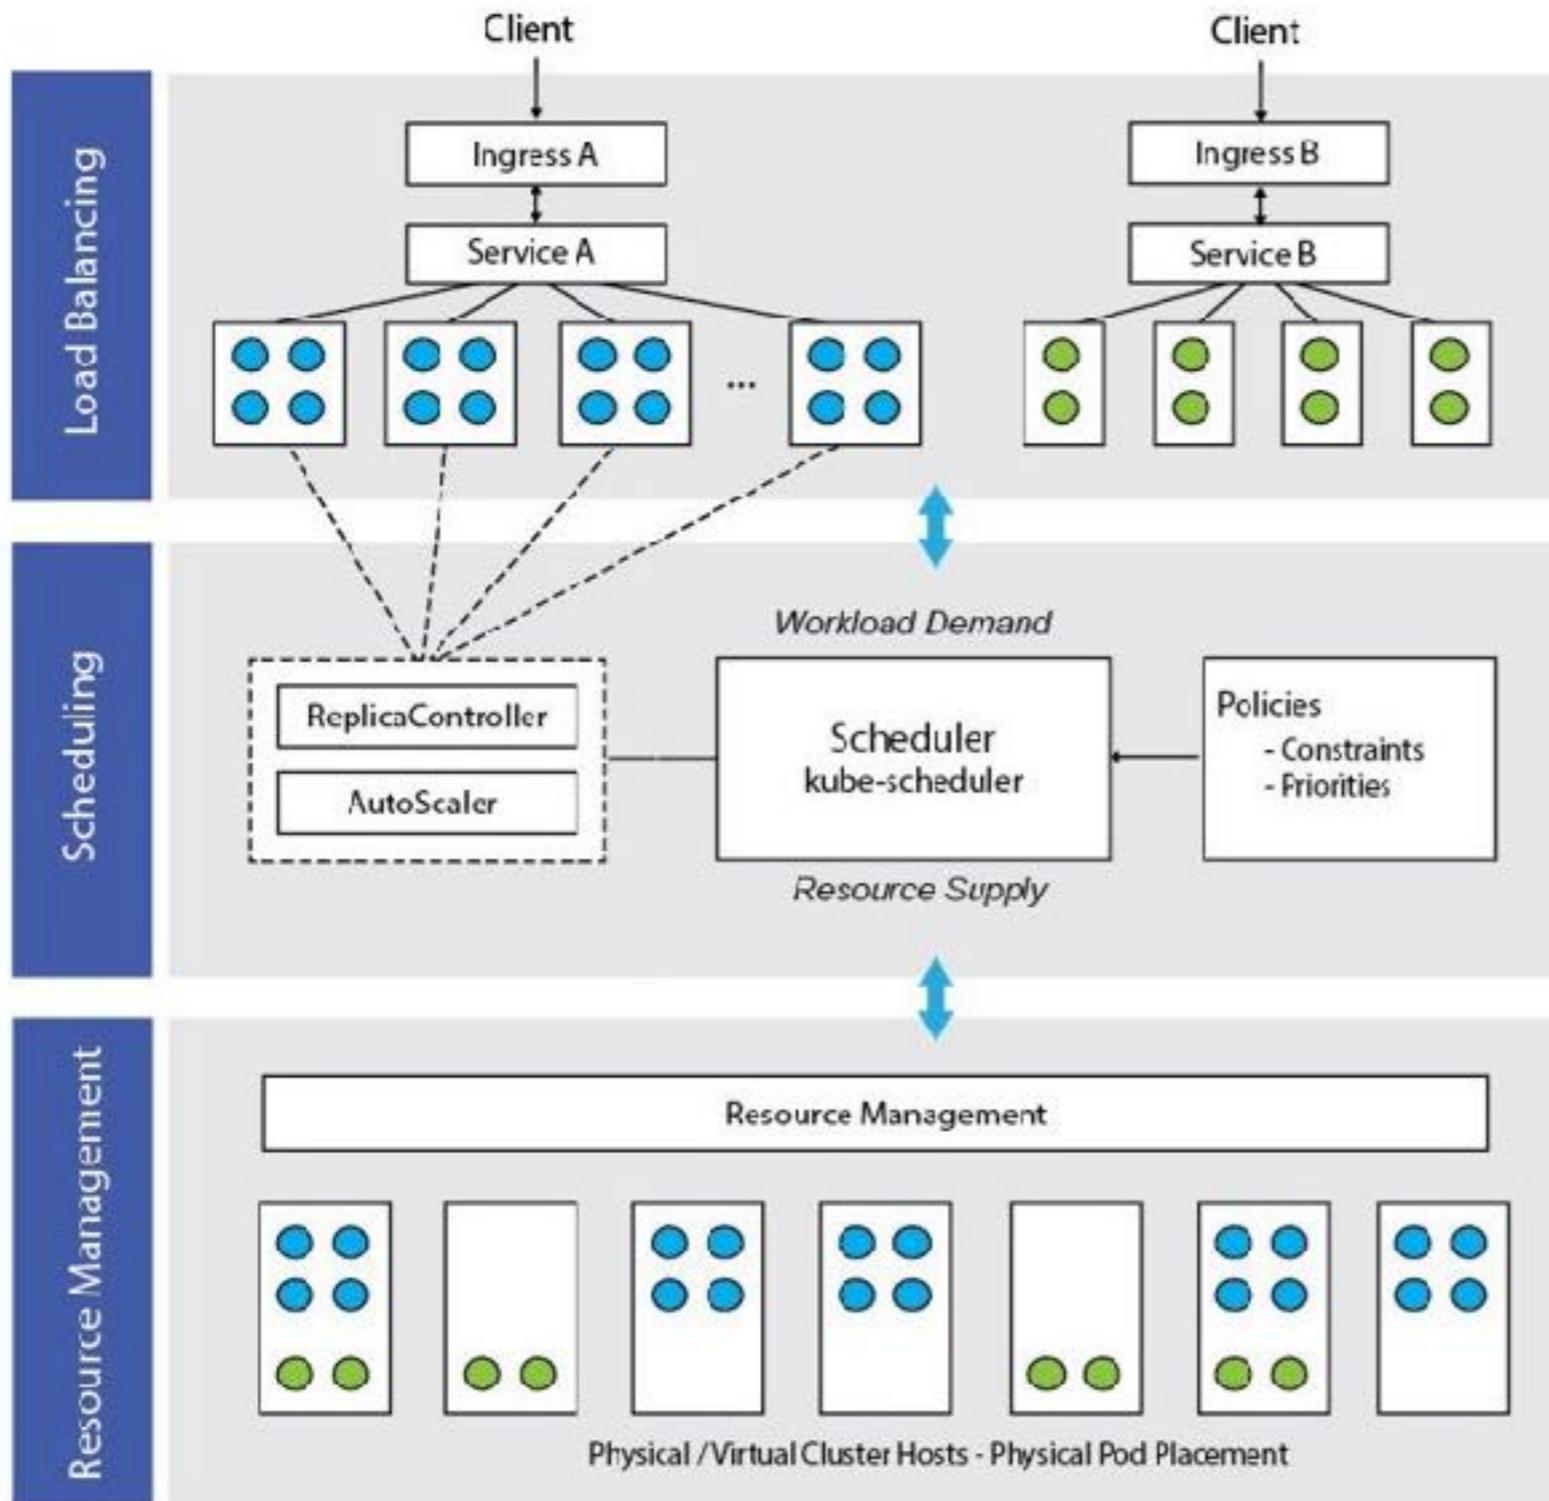


Open Source  
Industry choice  
Next big thing  
Container centric  
Inspired by Google's Borg  
Pretty damn cool



# KUBERNETES 101.





■ Availability - Networking

■ Lifecycles handling

■ Computation resources

# KUBERNETES 101.



```
1  apiVersion: apps/v1beta2
2  kind: Deployment
3  metadata:
4    name: python-basic-api
5    labels:
6      app: pba
7  spec:
8    replicas: 2
9    selector:
10      matchLabels:
11        app: pba
12    template:
13      metadata:
14        labels:
15          app: pba
16    spec:
17      containers:
18        - name: python-basic-container
19          image: murrau/pba:ready
20          ports:
21            - containerPort: 8066
22          imagePullPolicy: Always
23
24  ...
25
26  apiVersion: v1
27  kind: Service
28  metadata:
29    name: pba-service
30    labels:
31      app: pba
32  spec:
33    selector:
34      app: pba
35    type: ClusterIP
36    ports:
37      - protocol: TCP
38        name: http
39        port: 80
40        targetPort: 8000
41
```

```
1  apiVersion: extensions/v1beta1
2  kind: Ingress
3  metadata:
4    name: basic-apis-ingress
5  annotations:
6    kubernetes.io/ingress.class: "nginx"
7    nginx.ingress.kubernetes.io/rewrite-target: /
8  spec:
9    rules:
10      - host: rnd.gcn-eibp.com
11        http:
12          paths:
13            - path: /nba
14              backend:
15                serviceName: nba-service
16                servicePort: 80
17            - path: /pba
18              backend:
19                serviceName: pba-service
20                servicePort: 80
```

Am I writing  
**~300apps\*15customers\*4yaml**s ?  
**(PS: = 18000 files)**



Am I writing  
**~300apps\*15customers\*4yaml**s ?  
**(PS: = 18000 files)**

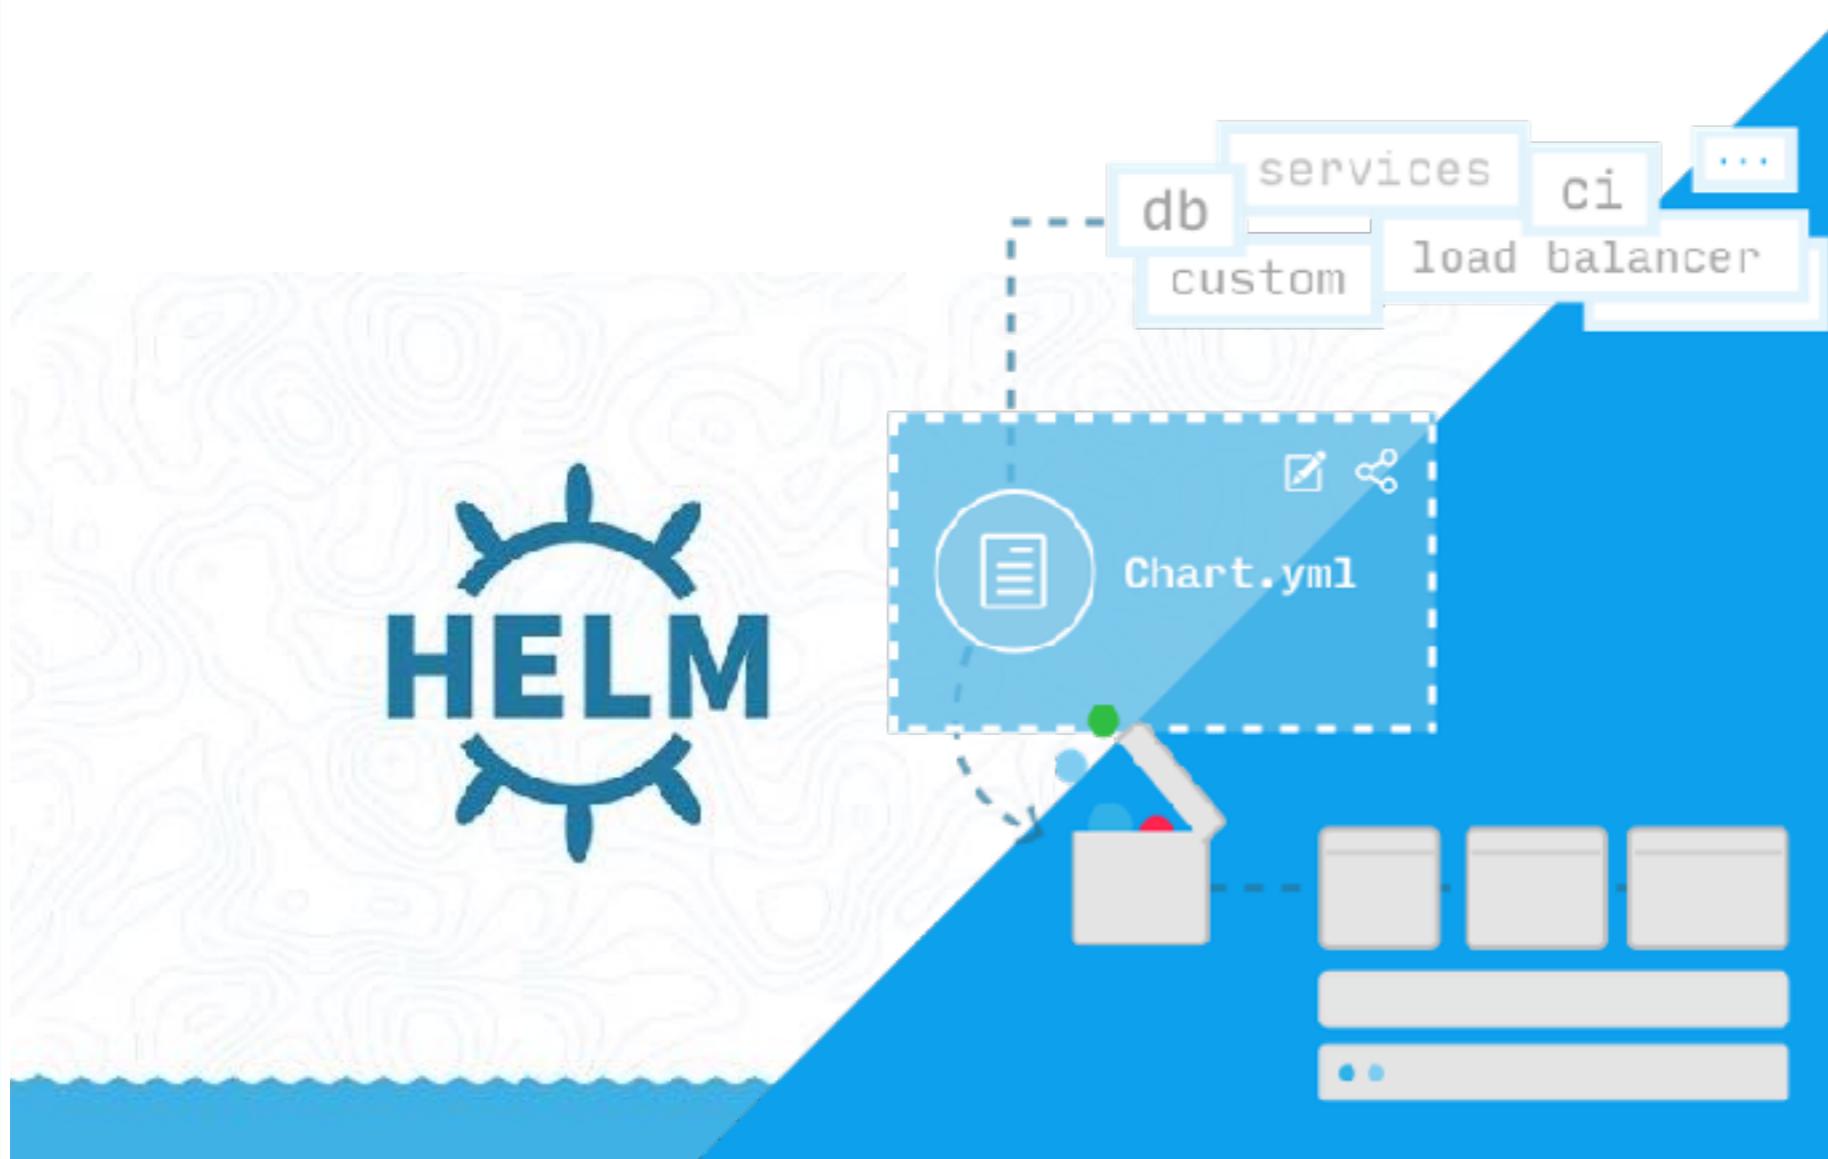
Also....

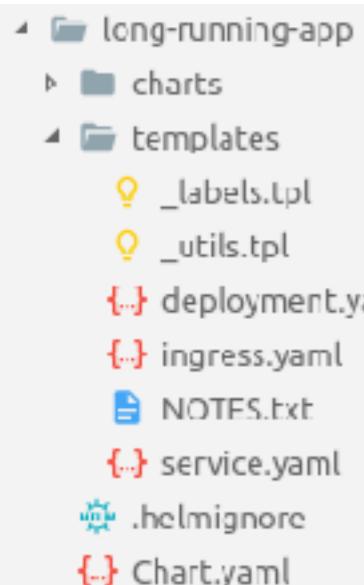
**Less code**



***Less mistakes, less problems, more fun***

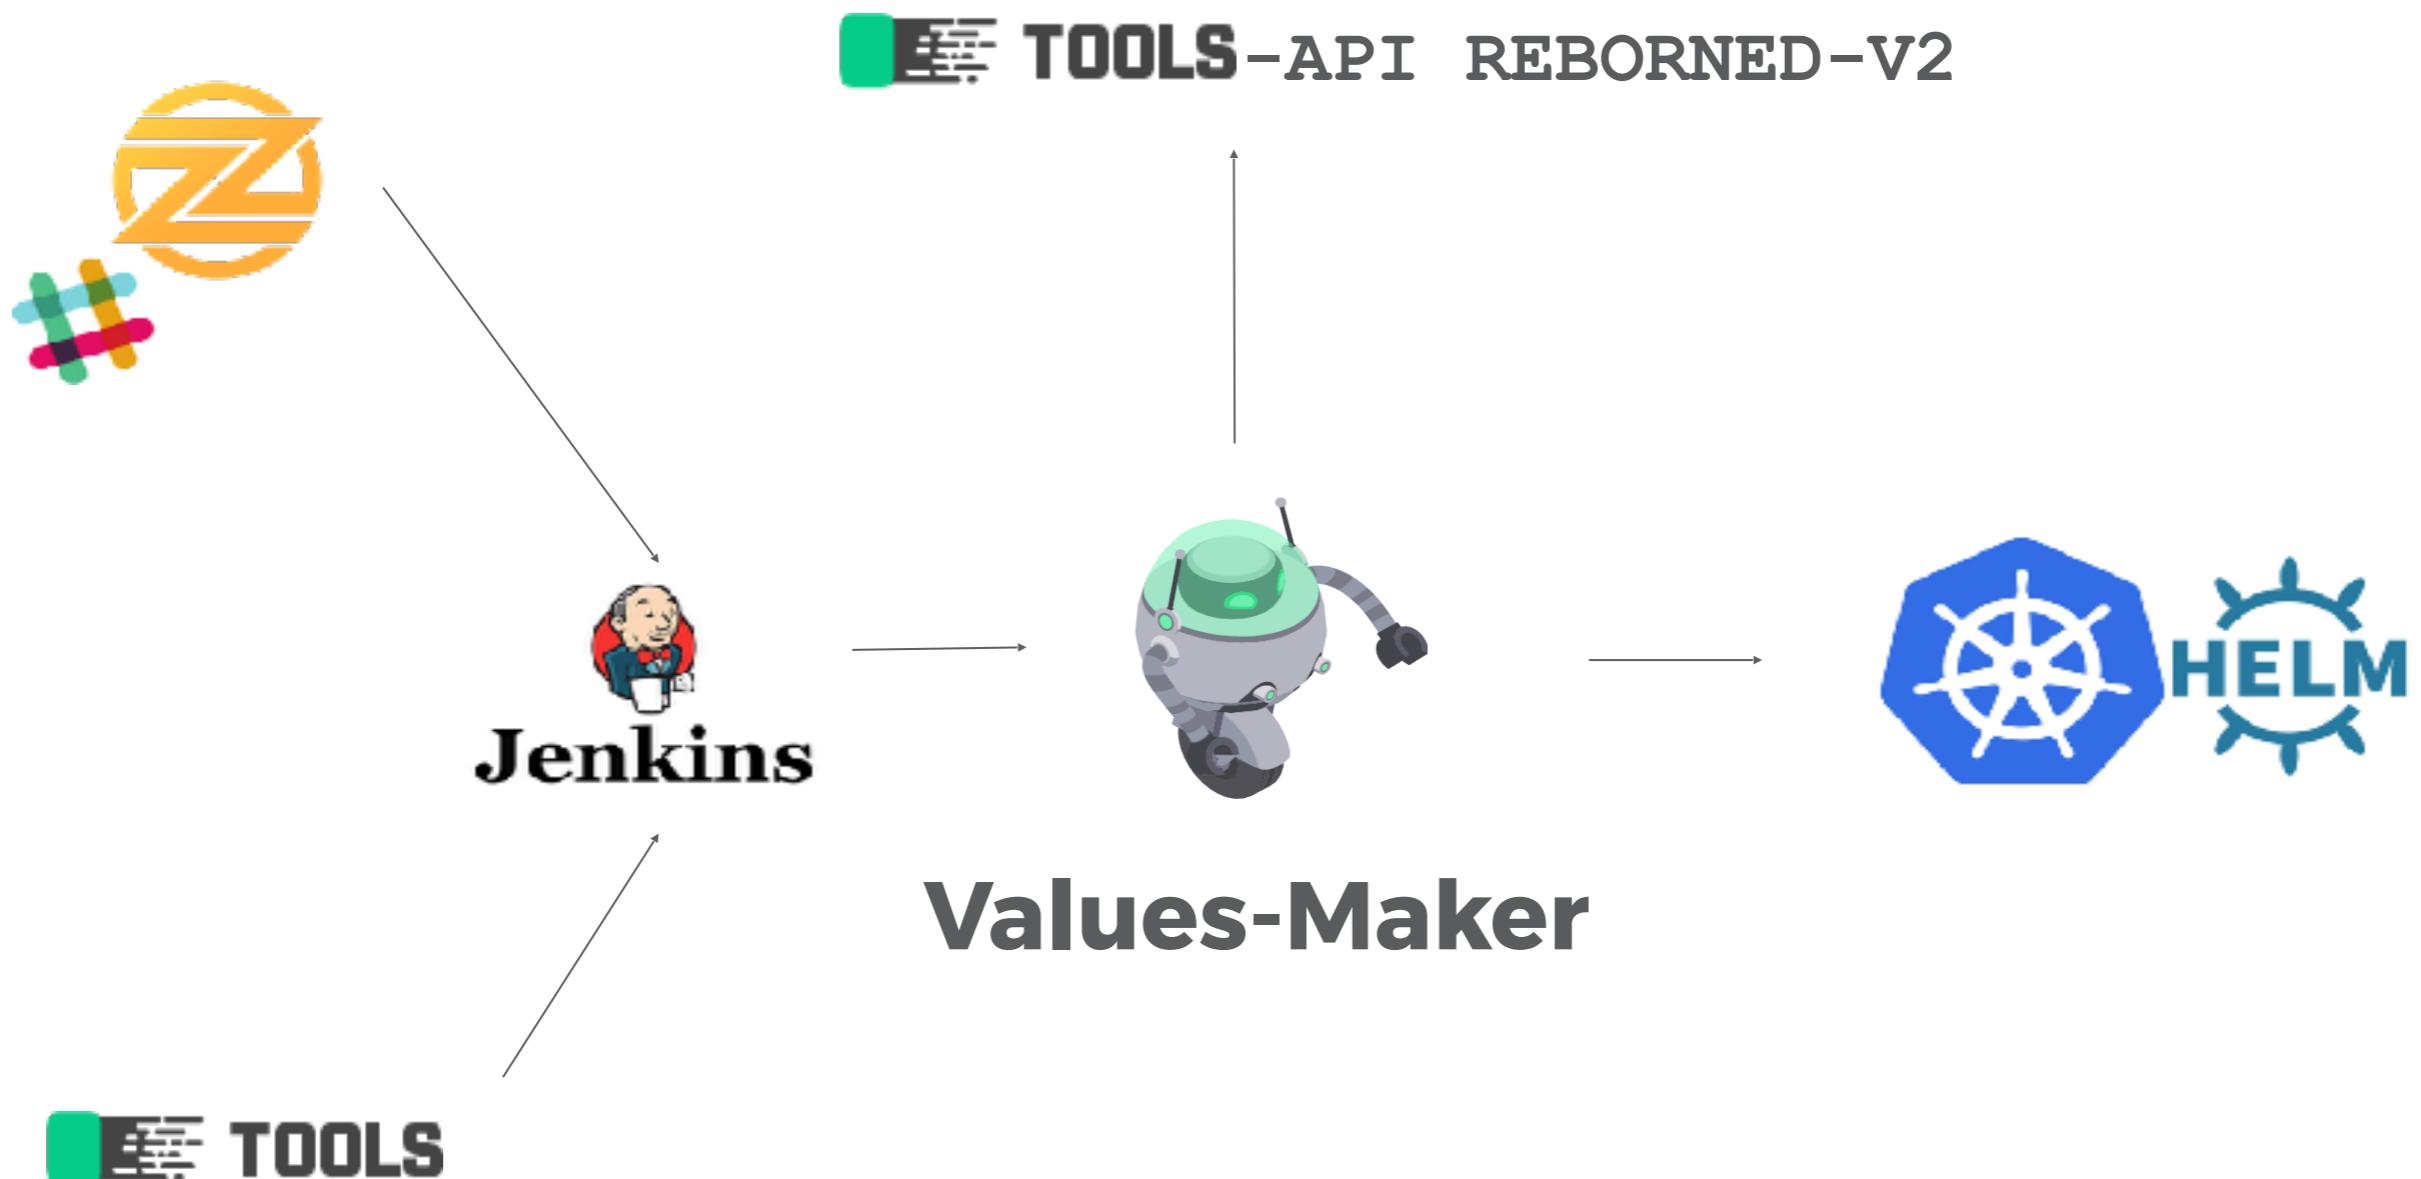
# KUBERNETES 101.





```
1  apiVersion: apps/v1beta2
2  kind: Deployment
3  metadata:
4    name: {{ .Values.app }}
5    labels: {{ include "labels.full" . | indent 2 }}
6  spec:
7    replicas: {{ .Values.replicaCount }}
8    selector:
9      matchLabels: {{ include "labels.short" . | indent 4 }}
10   template:
11     metadata:
12       labels: {{ include "labels.short" . | indent 6 }}
13     spec:
14       containers:
15         - name: {{ .Values.app }}
16           image: {{ .Values.registry.host }}/{{ .Values.registry.repository }}/{{ .Values.app }}:{{ .Values.tag }}
17
18         {{- if .Values.exposed -}}
19         ports:
20           - name: http
21             containerPort: {{ .Values.service.targetPort }}
22             protocol: TCP
23         {{- end -}}
24
25         {{- if .Values.resources -}}
26         resources:
27           {{- if .Values.resources.limits -}}
28             limits:
29               cpu: {{ .Values.resources.limits.cpu }}
30               memory: {{ .Values.resources.limits.memory }}
31           {{- end -}}
32         {{- end -}}
```

## ONE CHART TO RULE THEM ALL



## UPDATED (simplified) DEPLOYMENT PIPELINE

**GREENCOM**  
NETWORKS

**WE MADE IT.**



***PS: At that point, we haven't really moved forward though***

**TIME FOR BENEFITS.**

**Clean pipelines**

**Kubernetes is *almost* all-in-one**

**Load Balancing + Networking = *almost* free**

**On-the-shelf solutions**

**One place monitoring**

**Plenty of other cool stuff**



**AND FOR DOWNSIDES...**

***Too much cool stuff***

**Victims of industry choices**

**Learning curve, might be impressive**

**Many, many, many possibilities**



# WHERE IS THE CLOUD GONE ?

**Cloud providers were  
implementing their wrappers  
around Kubernetes.**



**GREENCOM**  
NETWORKS

**LET'S GET IN THE CLOUD.**

**Kubernetes integrations on-the-shelf**

**And let's deal with what's remaining**

---

**LET'S GET IN THE CLOUD.**

**Networking between components**

**Providing computing resources**

**On-demand installation and scalability**



**GREENCOM**  
NETWORKS

**LET'S GET IN THE CLOUD.**



**AVOID VENDOR LOCKING  
RE-USE AND MIRROR OUR ON-PREMS INSTALL  
MANAGE EVERYTHING THE SAME WAY**

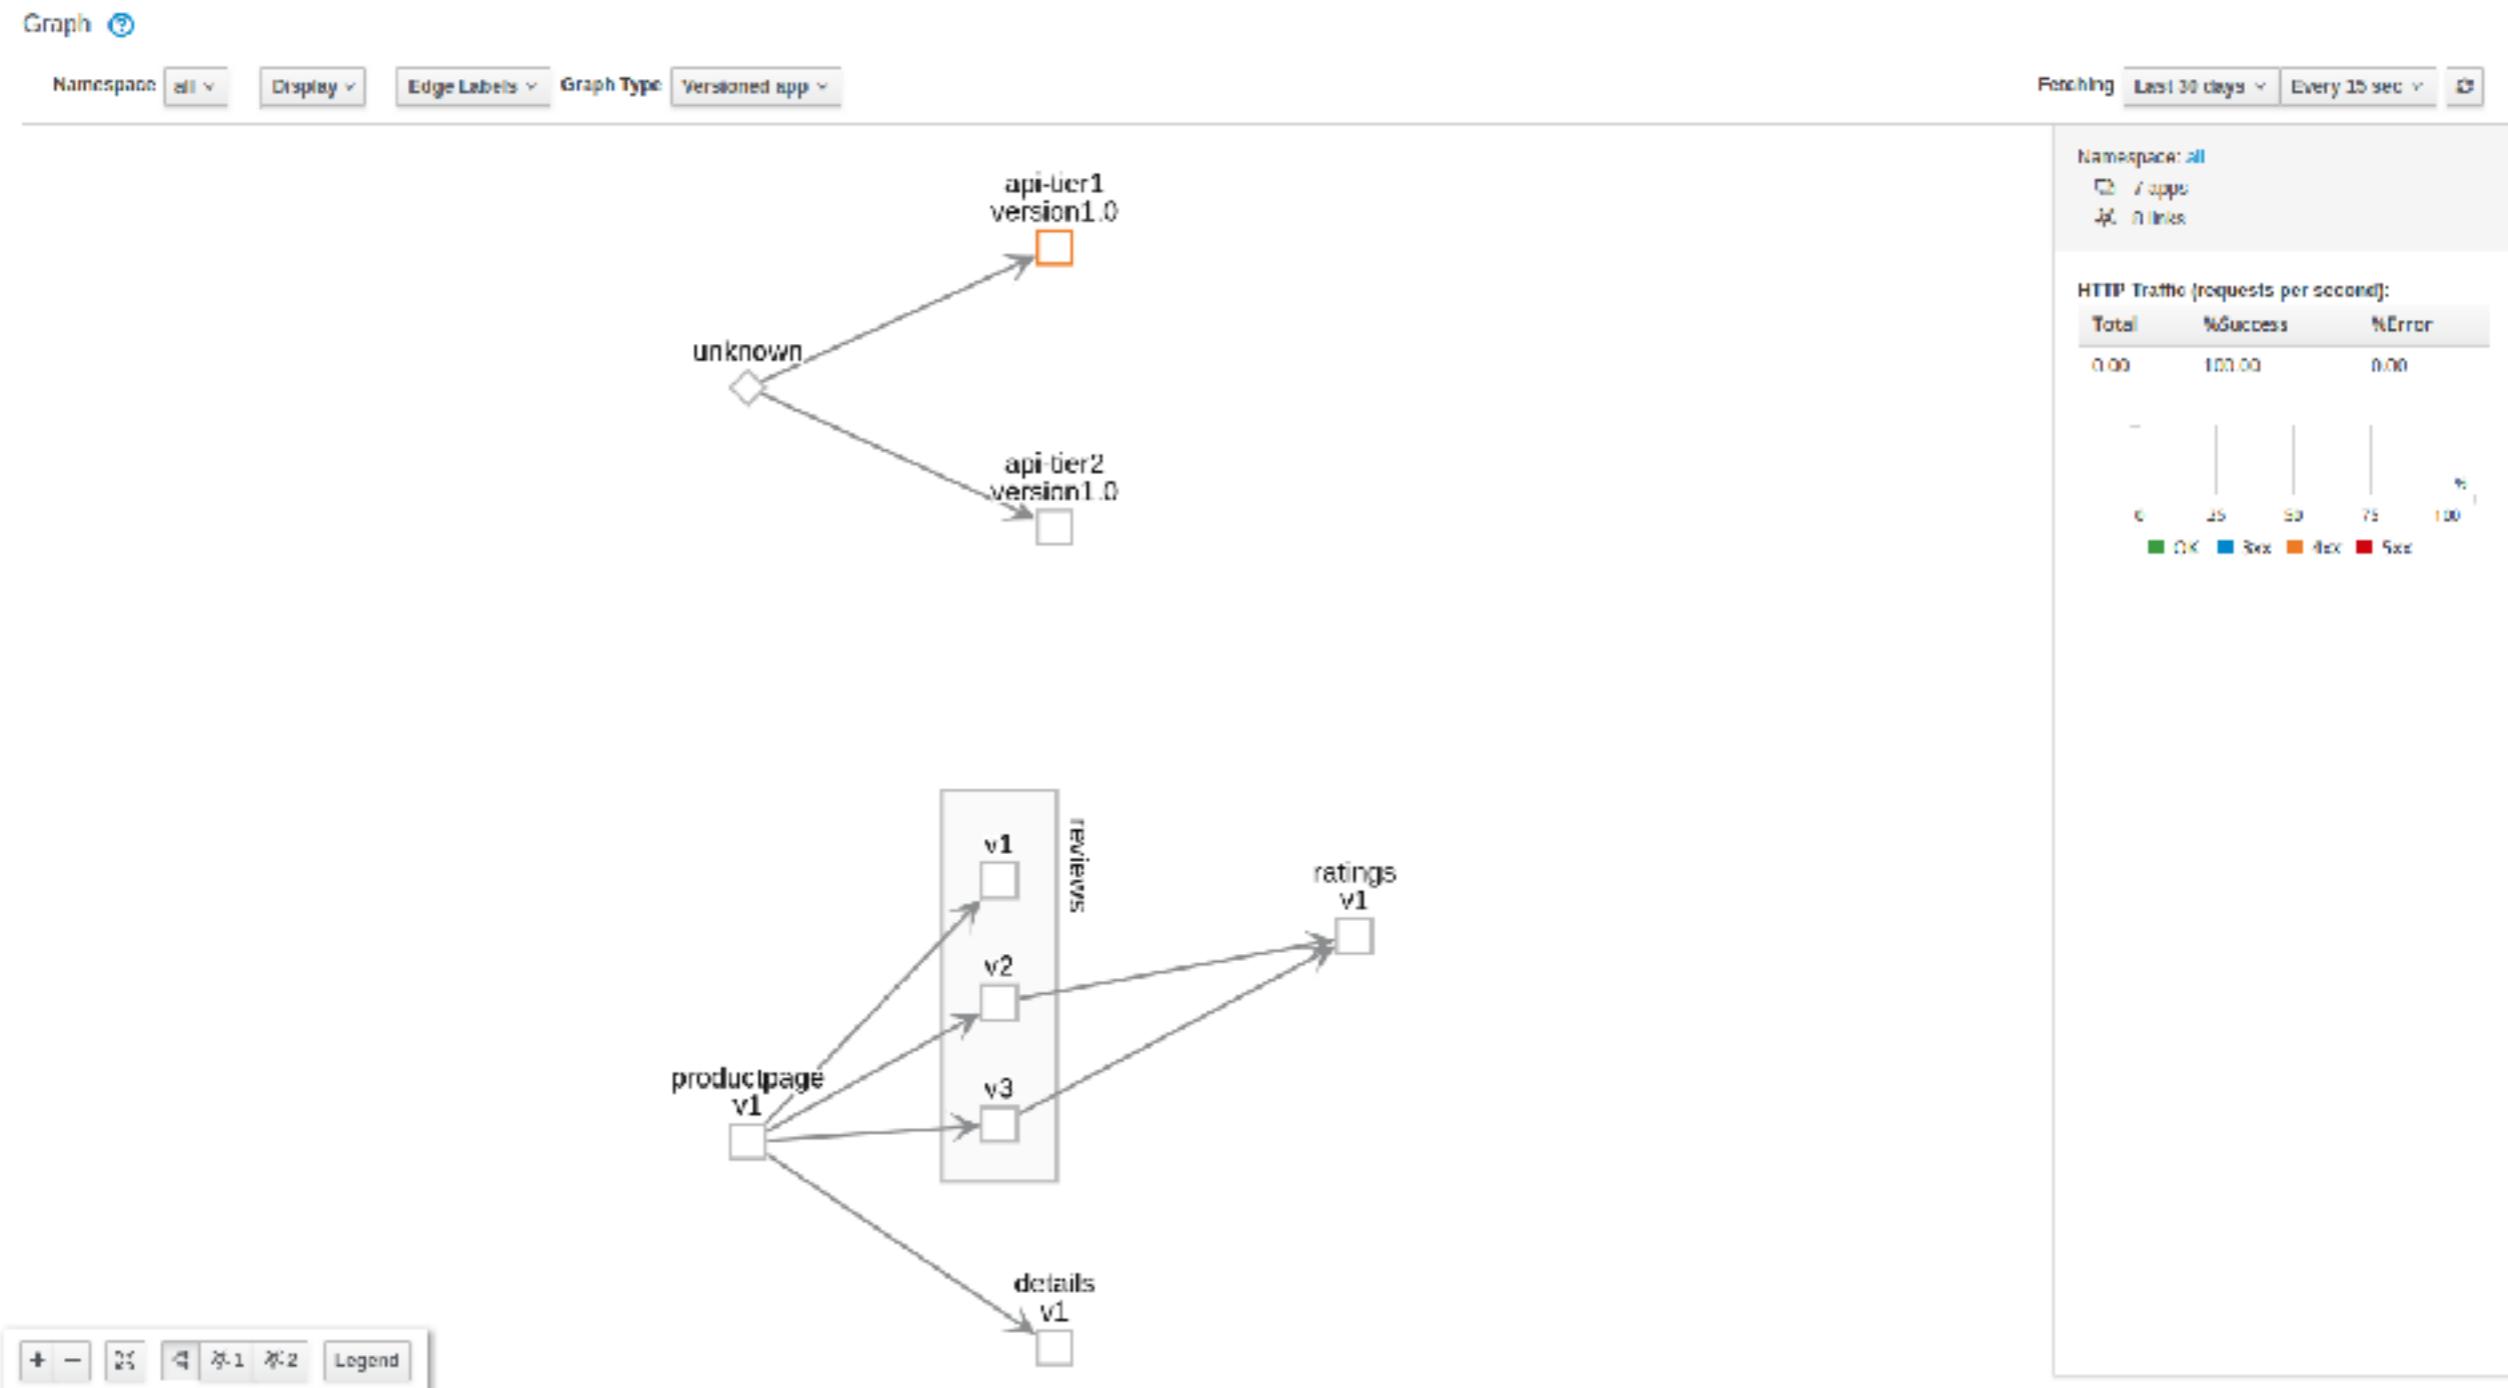
**GREENCOM**  
NETWORKS

**STILL A LONG WAY TO GO.**

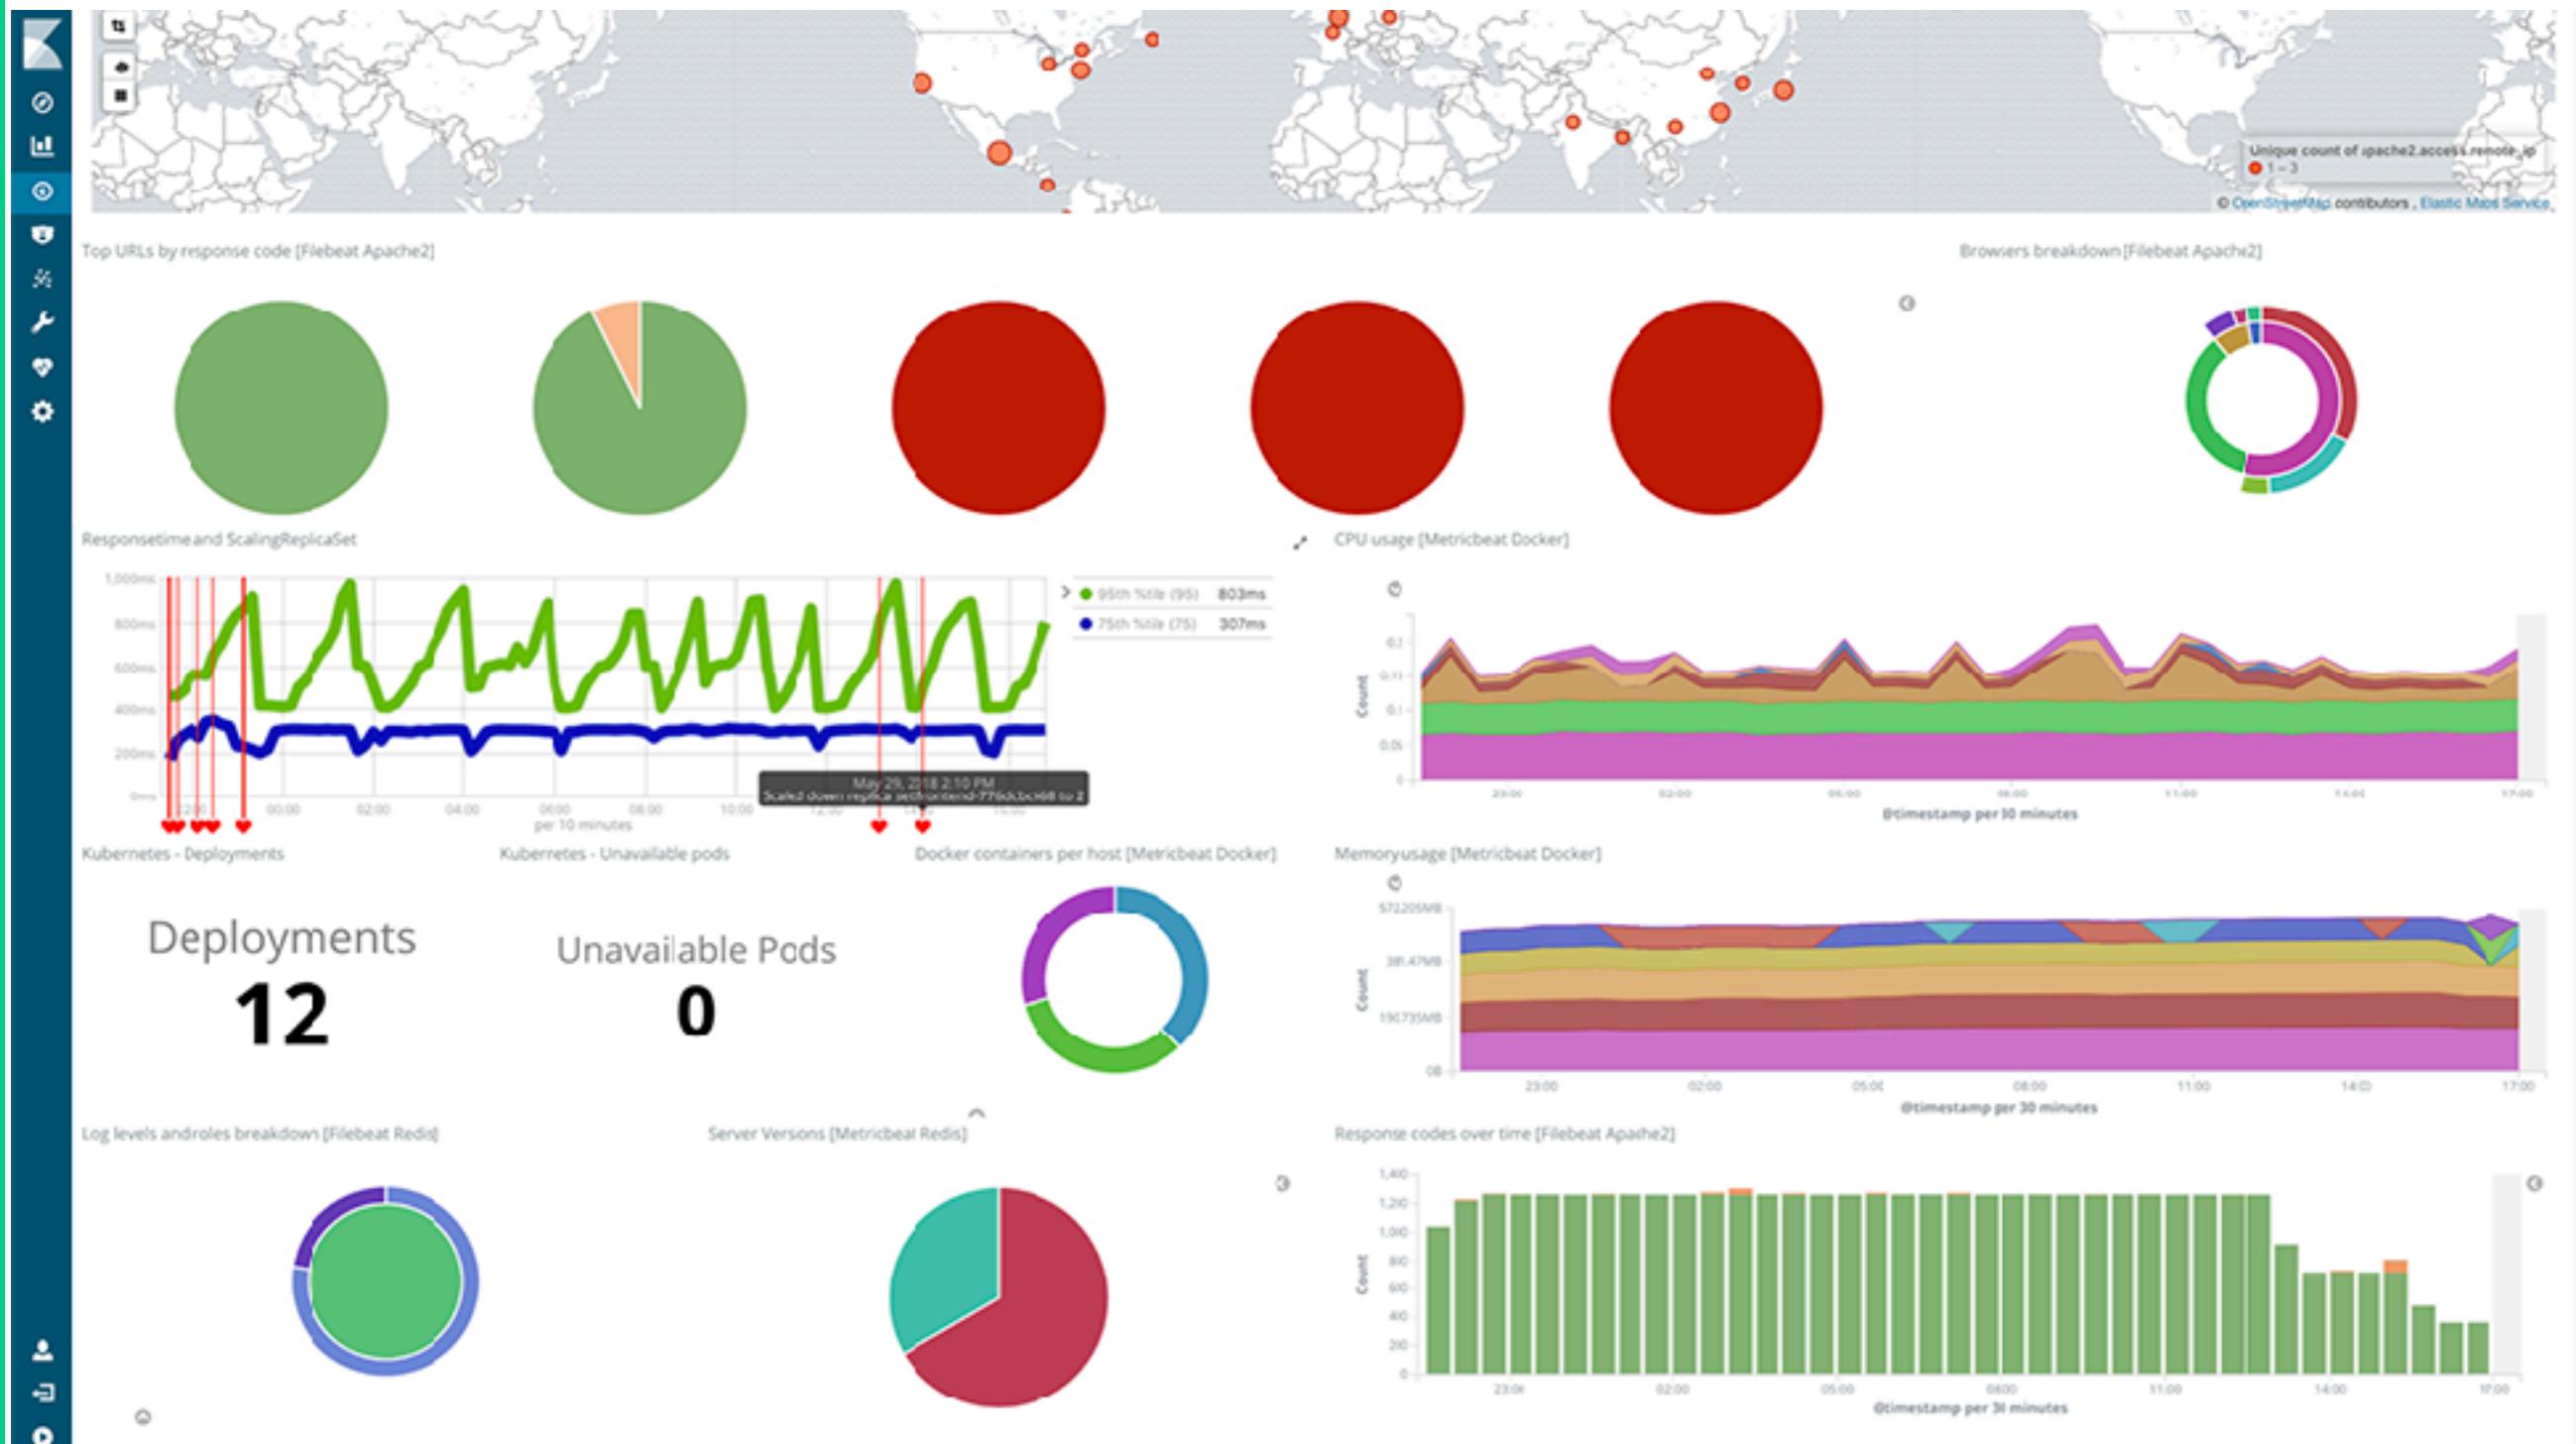


**But that's the easy part**

# SERVICE MESH LIVE STATUS.



# COST ANALYTICS.



**AND SO ON.**

**Automatic Blue-Green deployment**

**On-demand scalable storage**

**Circuits breakers**

**Self-recovering workload**

...

---

# Service mesh management

• Zero Downtime deployment	✓
• Dark Launch	✓
• Logging	✓
• A/B Testing	✓
• API Management	✓
• Tracing	✓
• Canary Release	✓
• Blue/Green Deployment	✓

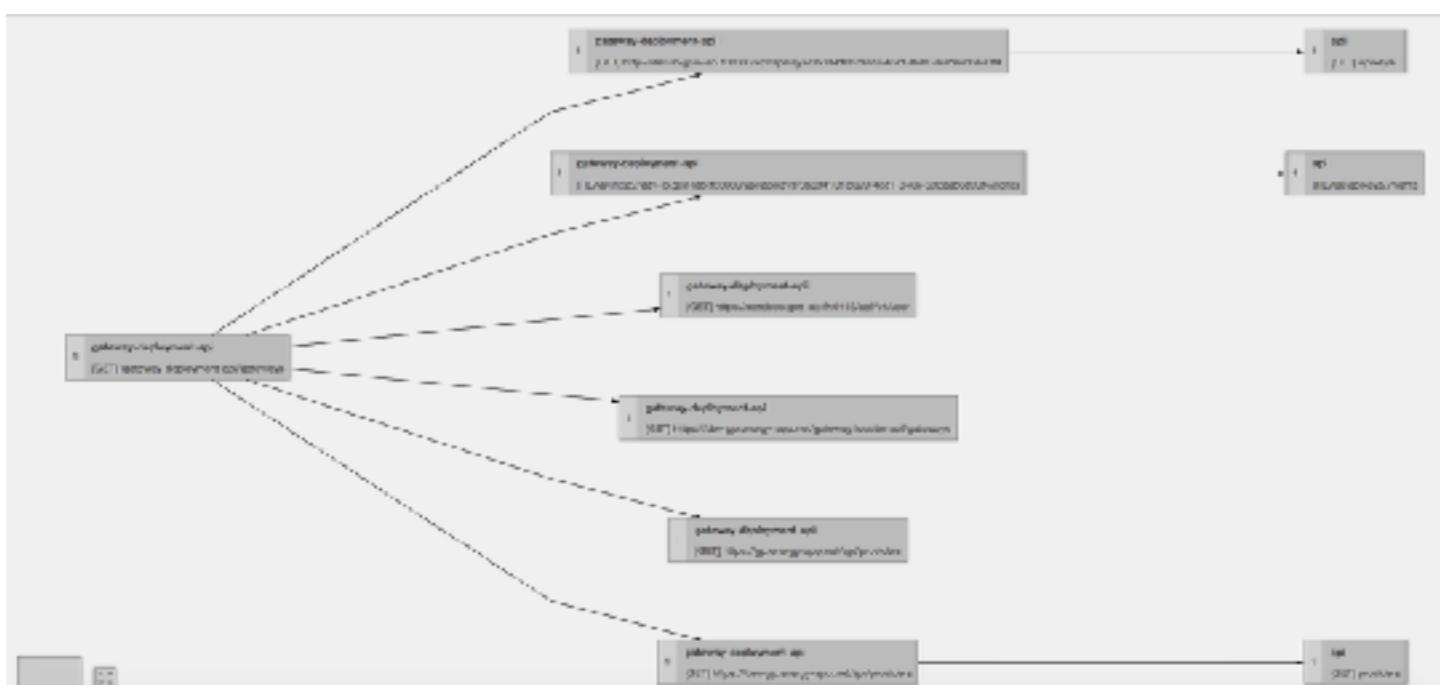
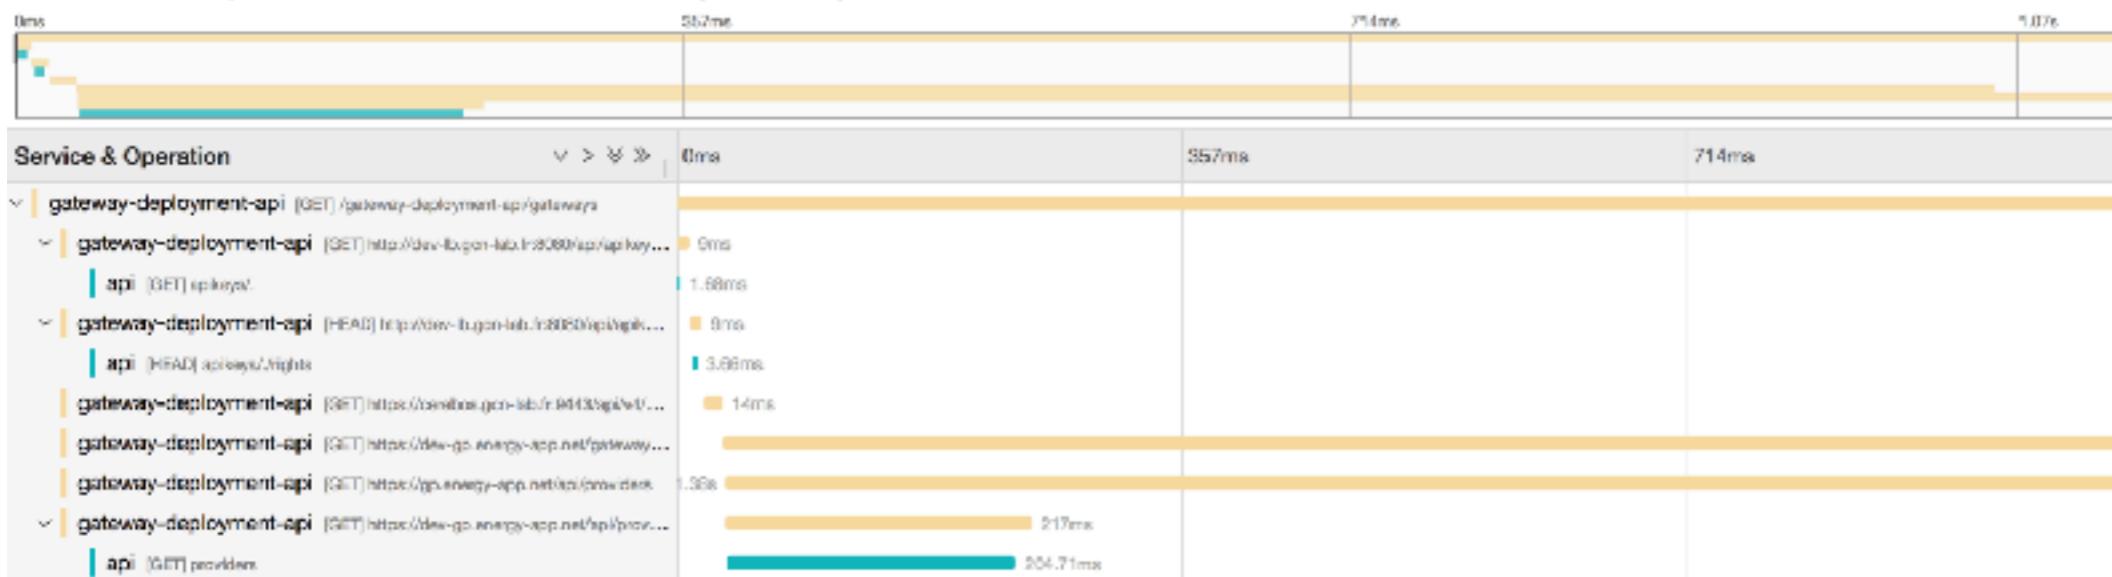
# Tracing



# JAEGER

## ✓ gateway-deployment-api: [GET] /gateway-deployment-api/gateways

Trace Start: October 24, 2018 8:45 PM | Duration: 1.43s | Services: 2 | Depth: 3 | Total Spans: 10



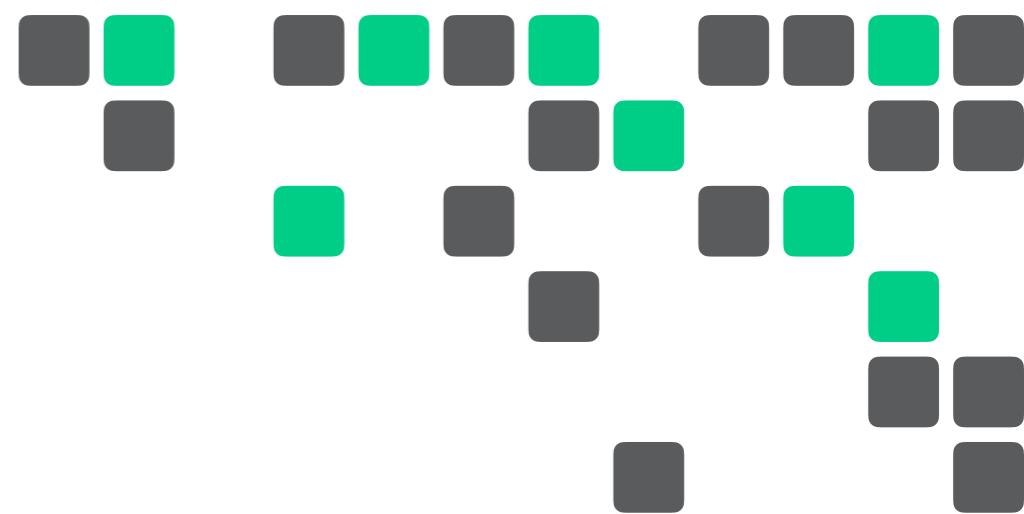


You have to stop now...  
They all are hungry, some of them  
are sleeping for the last 20  
minutes...

# Wrap Up

## Platform Design Principles:

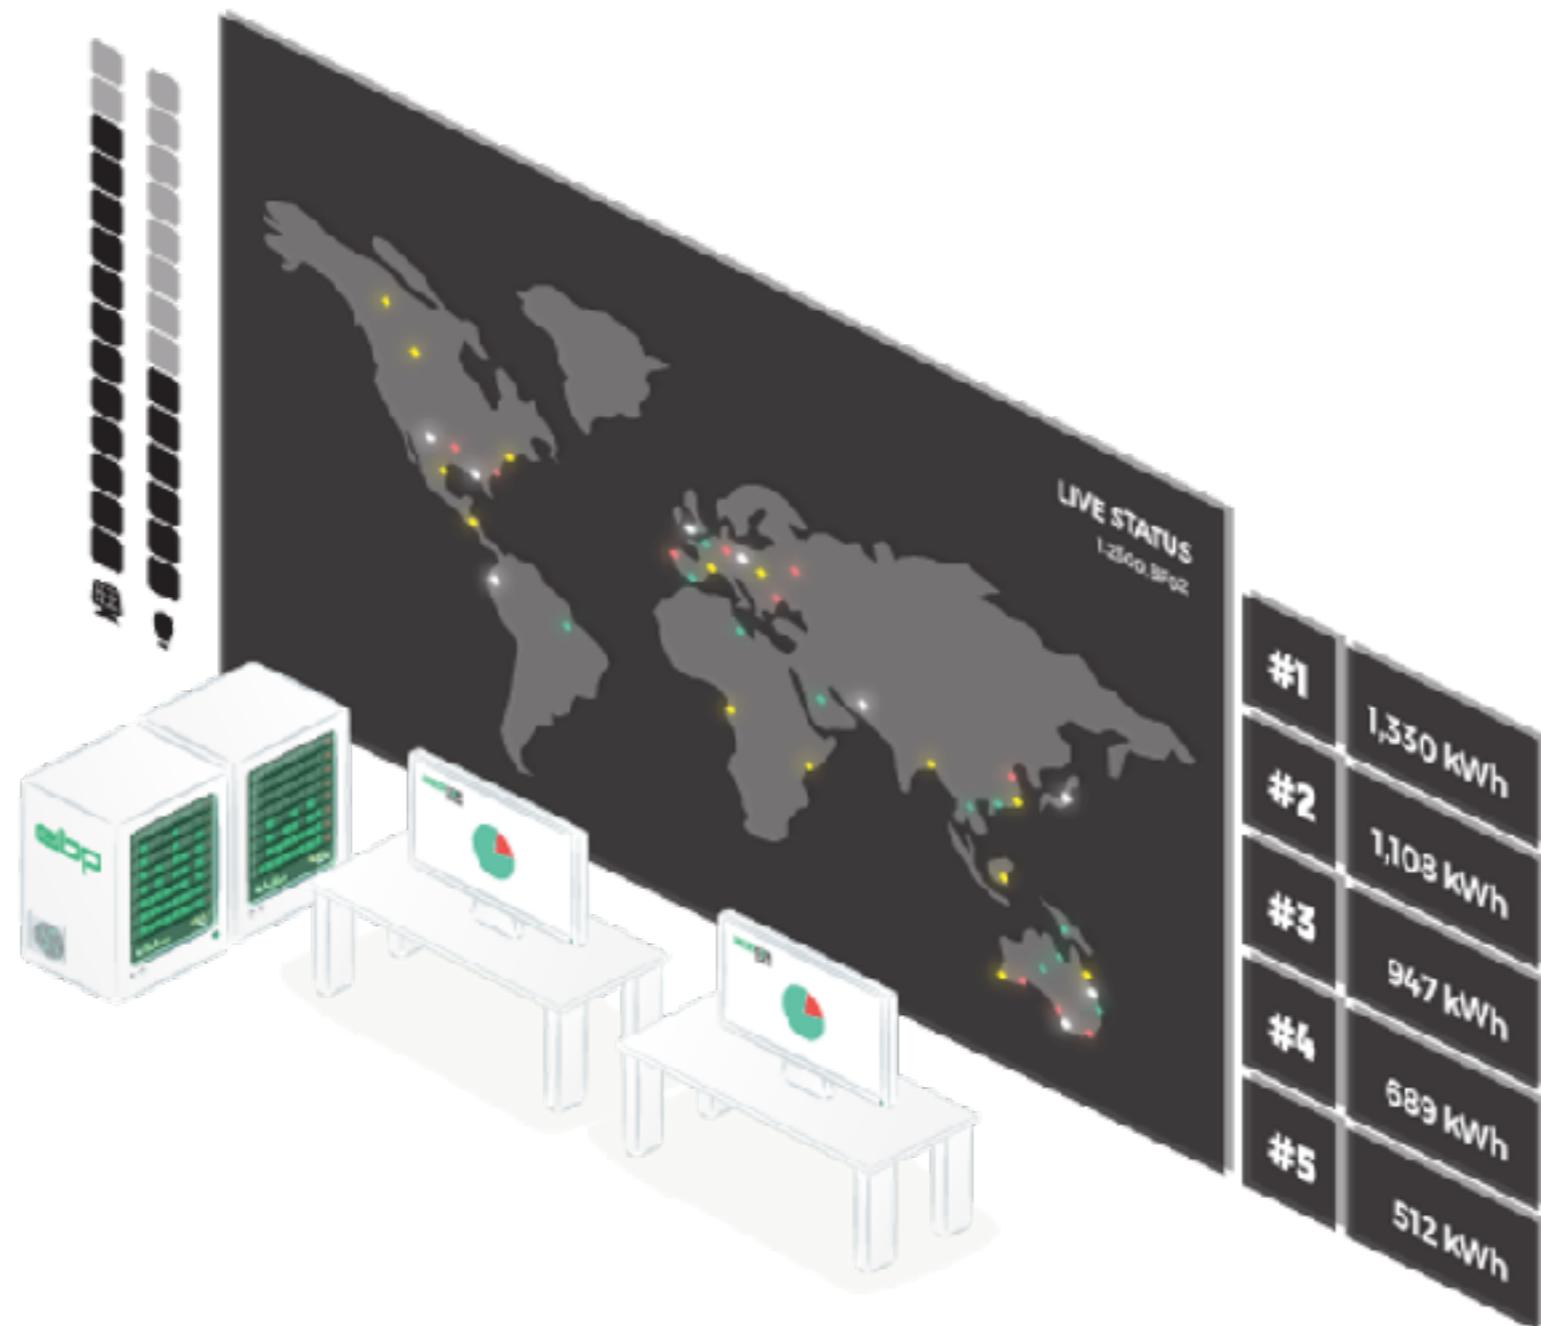
- Have a look to companies who are facing today volumes we will have to handle tomorrow
- Keep performance, scalability in mind
- No Single Point Of Failure
- Keep it simple / No over-engineering
- Agility is Key
- No choice if definitive



*Continuous improvement is better than delayed perfection.*

*Mark Twain*

**THAT'S IT.**



**GREENCOM**  
NETWORKS

